

Scaling Language Models

Weijia Shi



Feel free to interrupt with
questions / thoughts anytime!

Some slides are adapted from previous presentations by Sean Welleck, Yue Xiang, Niklas Muennighoff, Quentin Anthony, Liwei Jiang, Ilya Sutskever

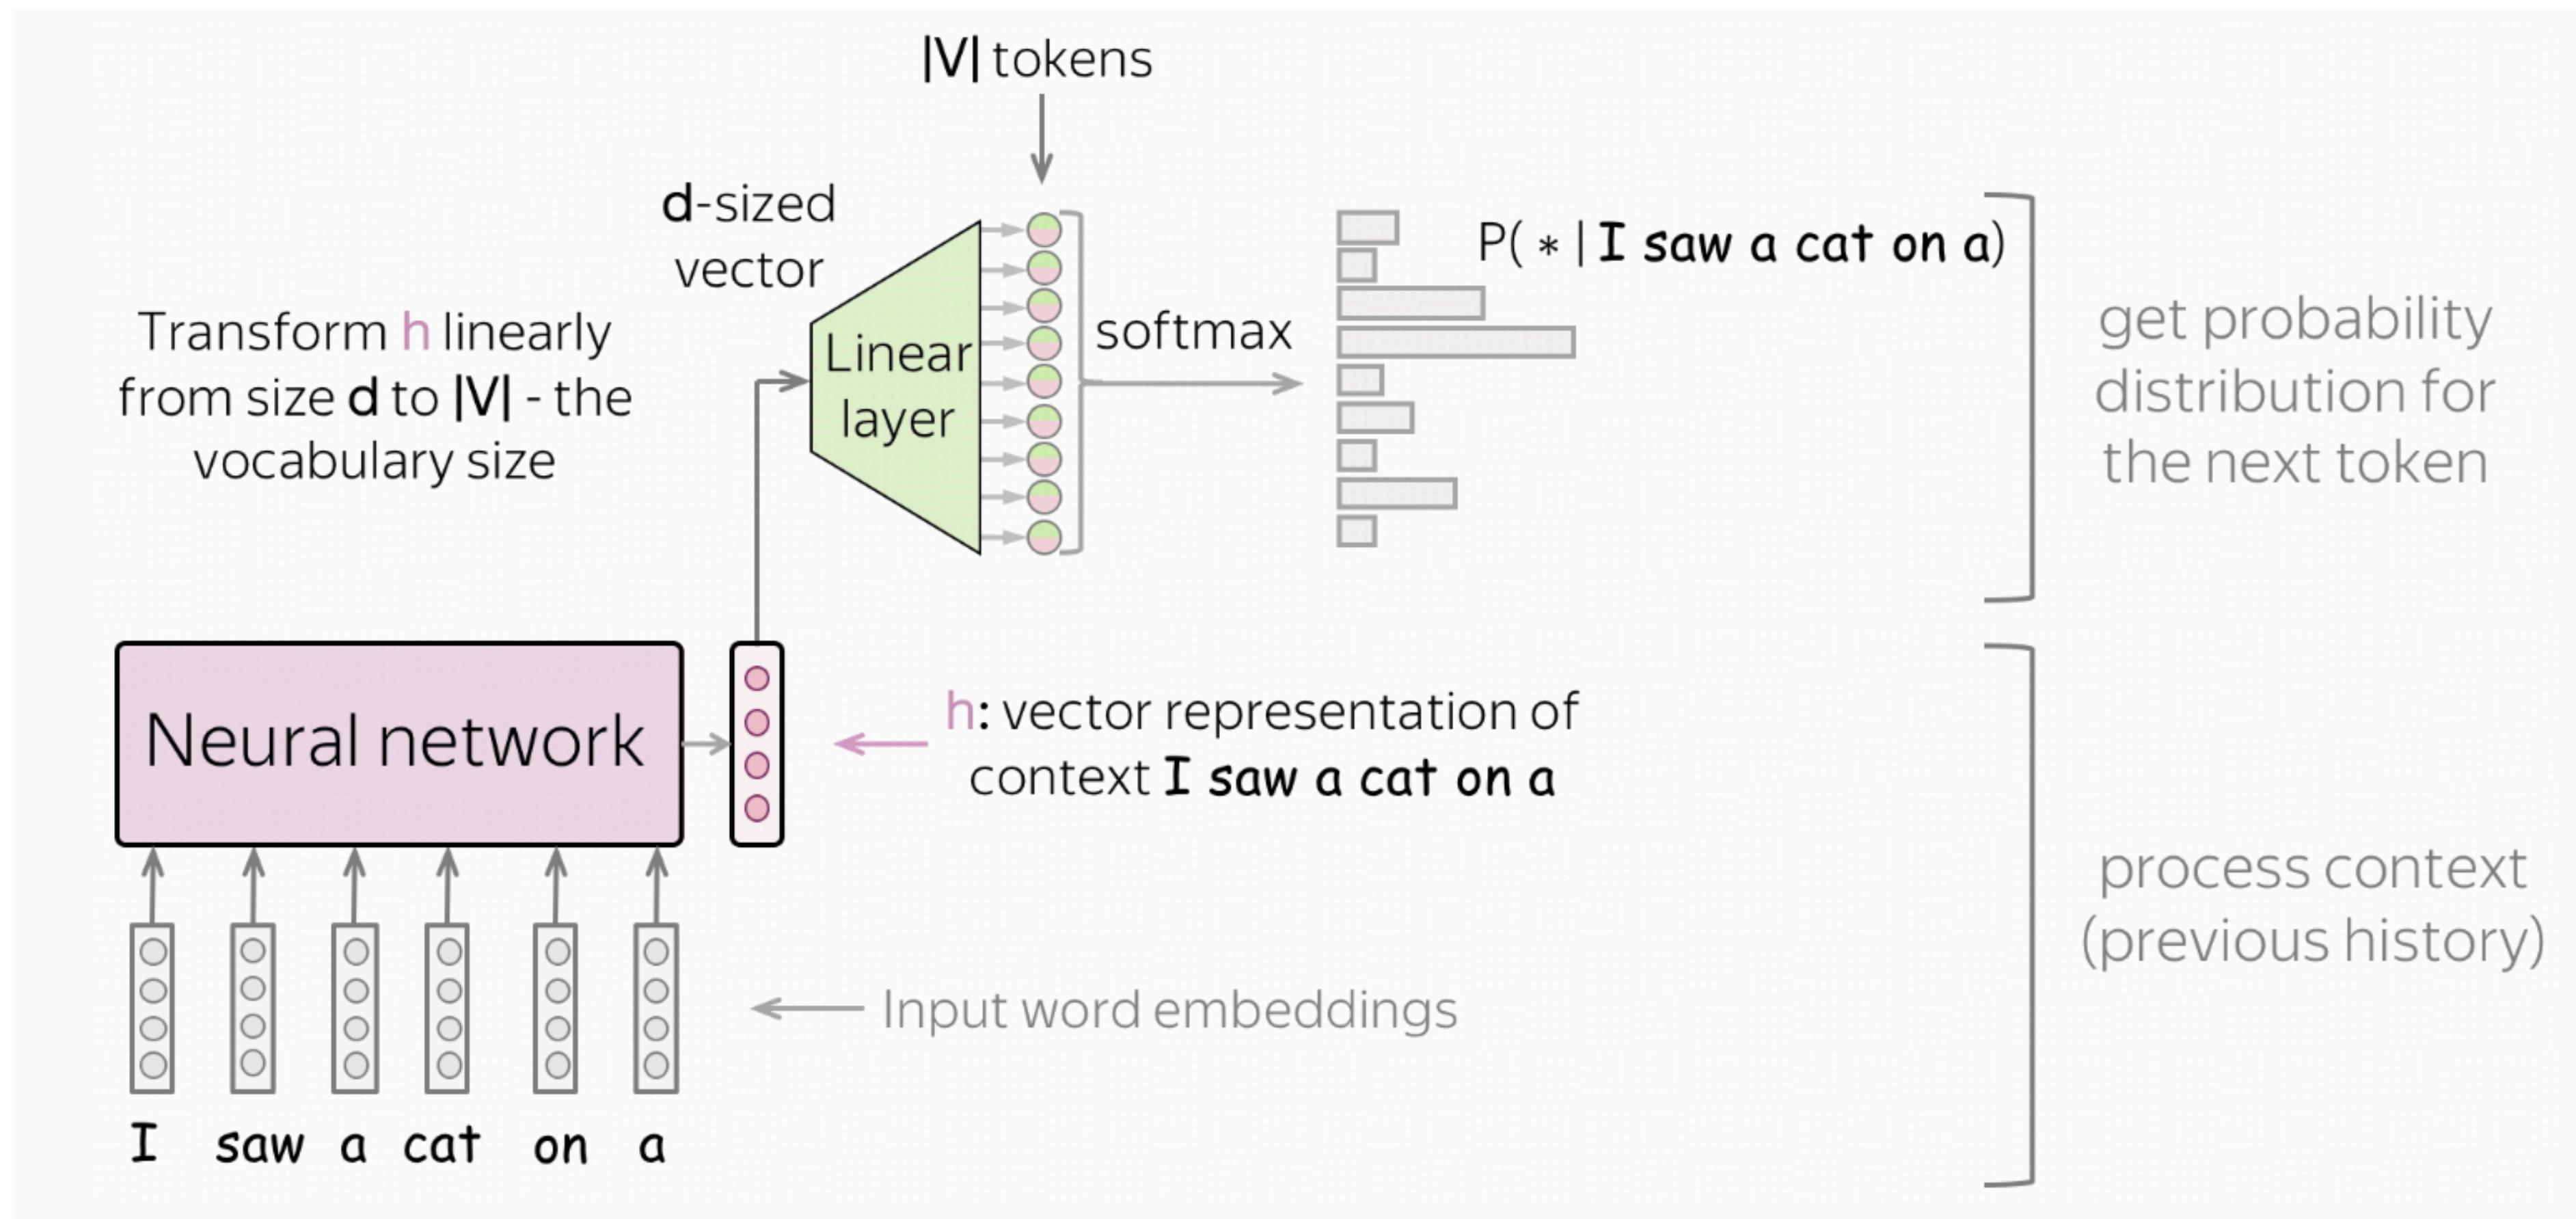
Recap of Language Modeling

Autoregressive Language Modeling

$$P(X) = \prod_{i=1}^I P(x_i \mid x_1, \dots, x_{i-1})$$

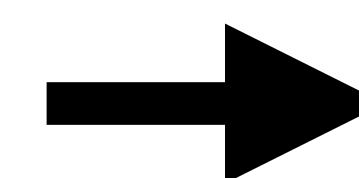
Next Token Context

Autoregressive Language Modeling

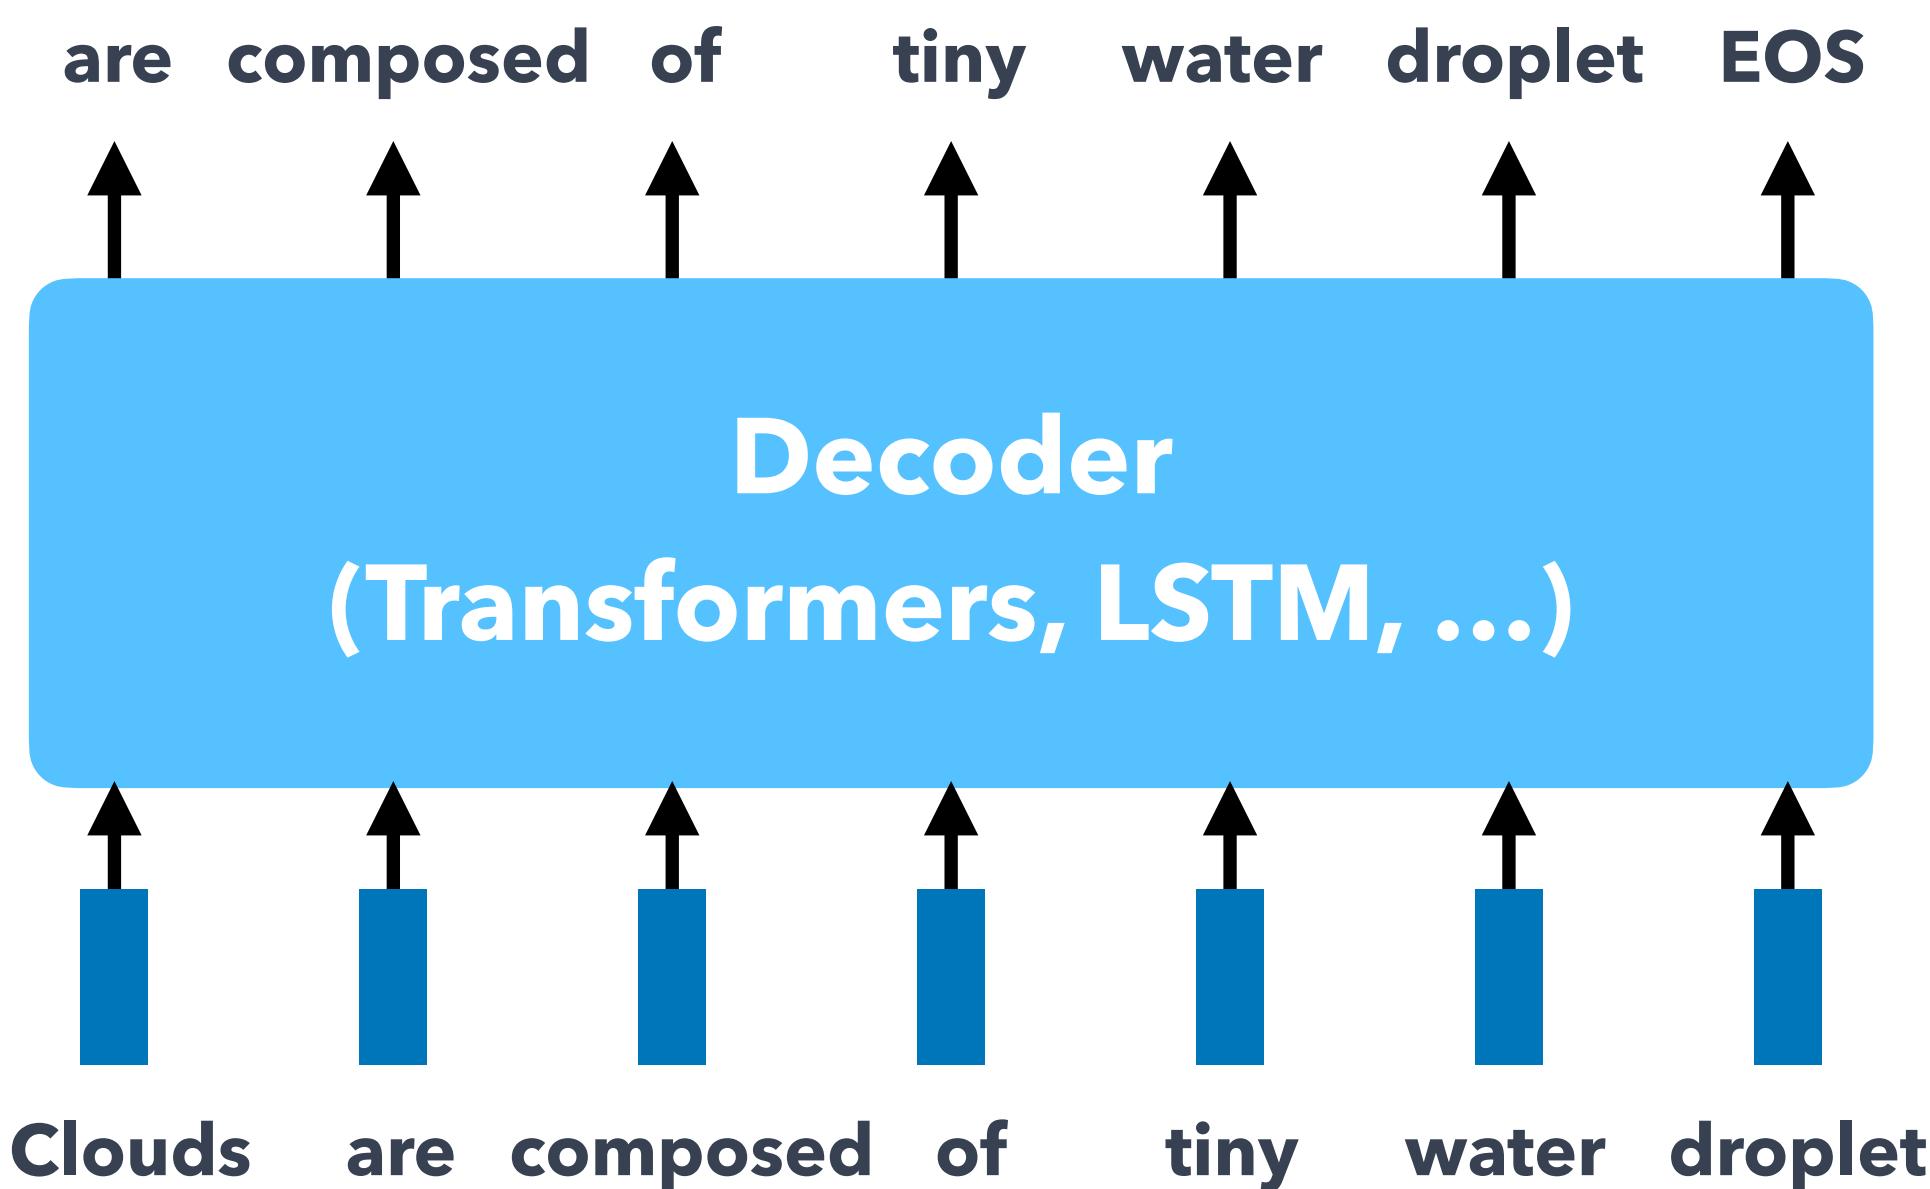


Why Is It Called Pretraining?

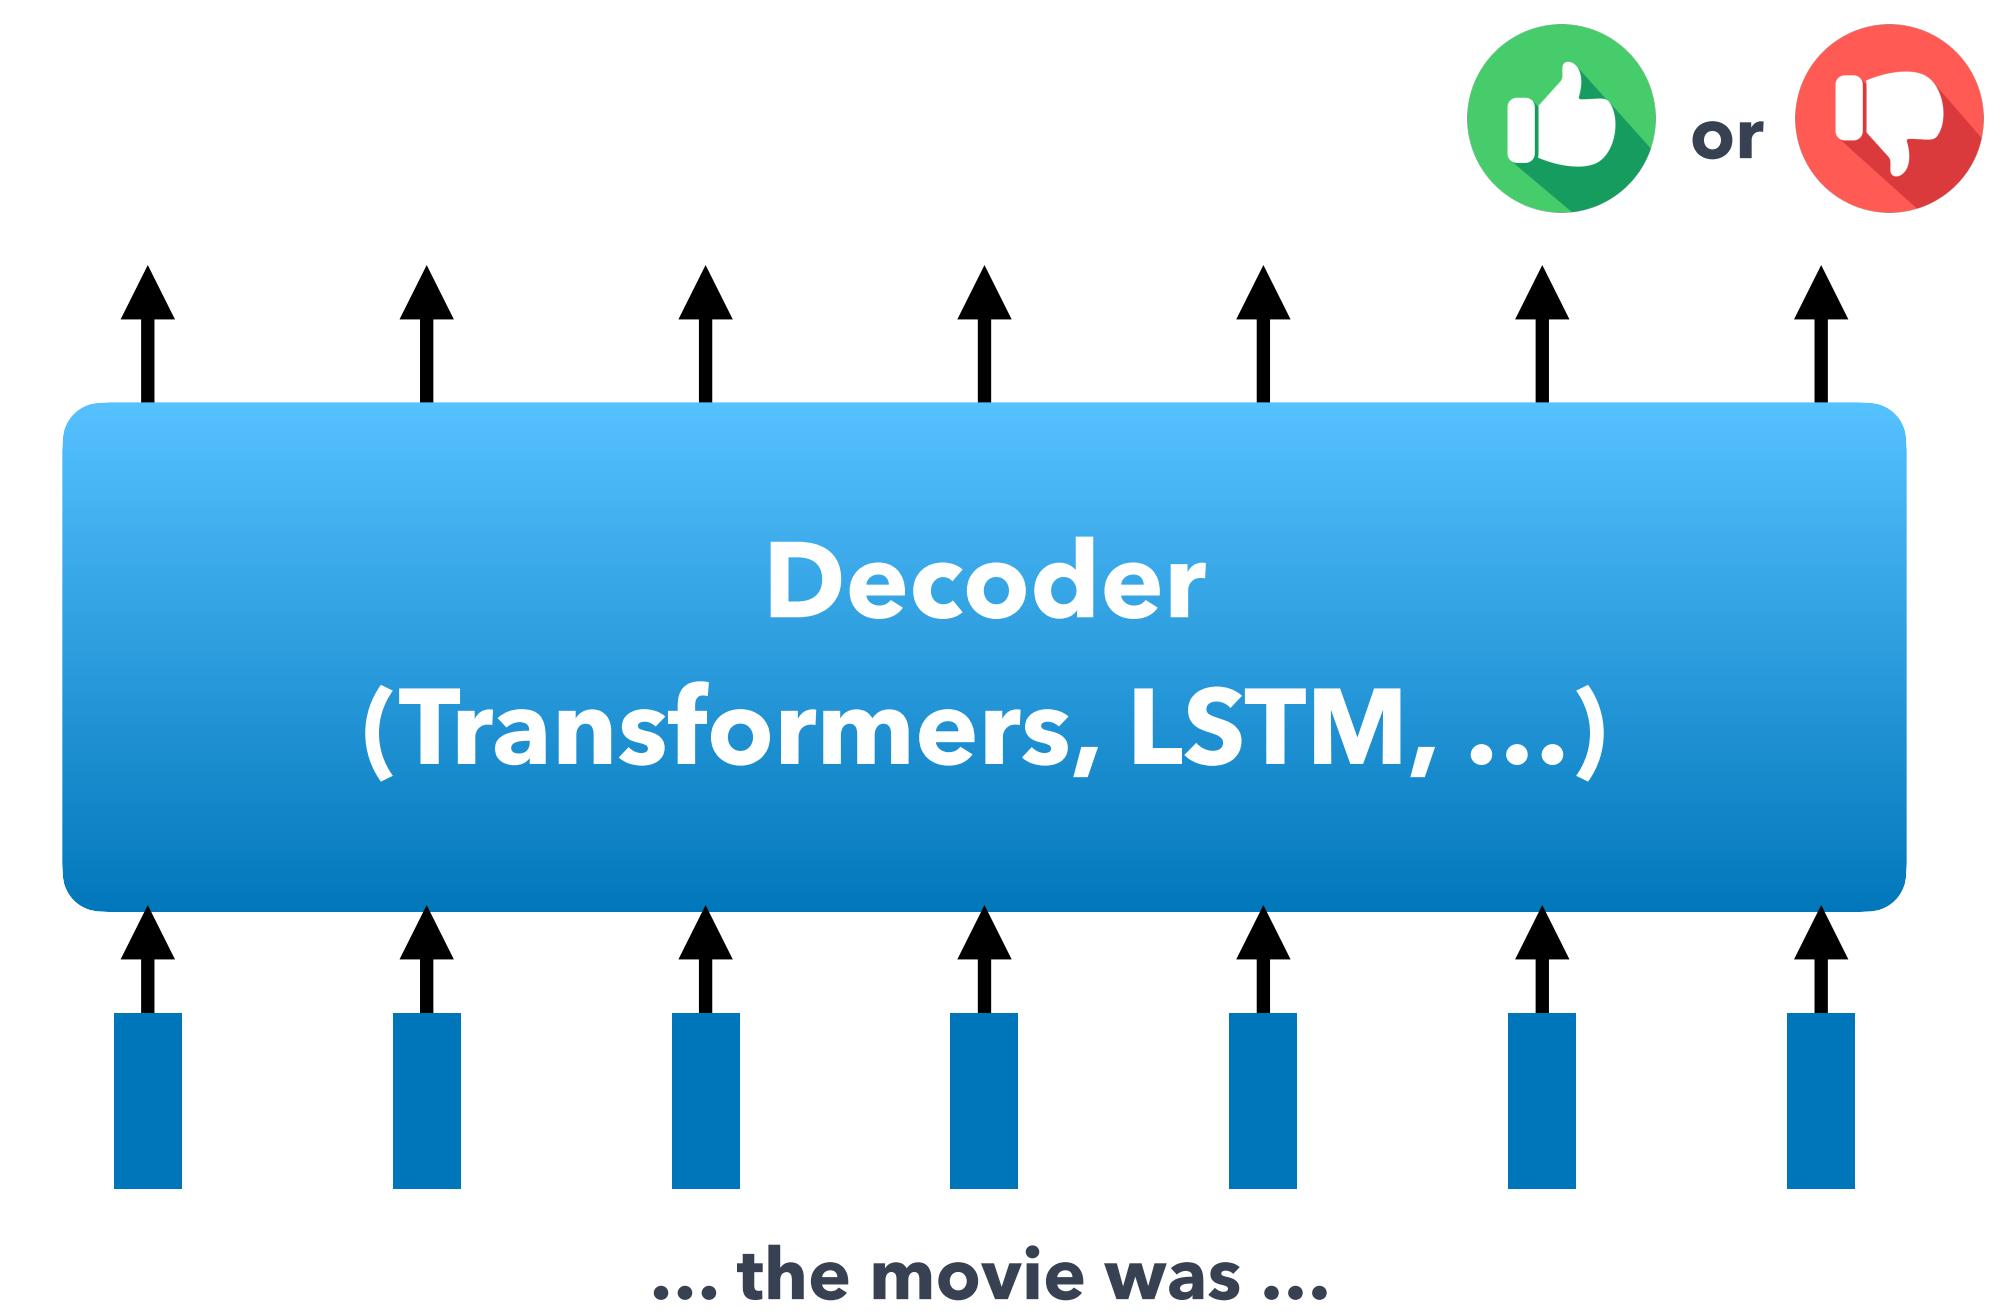
Step 1:
Pre-training



Step 2:
Fine-tuning



Abundant data; learn general language



Limited data; adapt to the task

“Pre”training happens before training (fine-tuning)!

Outline

- Overview of LM Pretraining
- Pretraining Data
- Pretraining Setups
- Scaling Law

Outline

- Overview of LM Pretraining
- Pretraining Data
- Pretraining Setups
- Scaling Law

Imagine you're developing Llama



Overview of Llama Training

Pretraining -> Instruction Fine-tuning -> RLHF

Why Pretraining?

Lots of Information in Raw Texts

Verb

I went to Hawaii for snorkeling, hiking, and whale _watching_.

Preposition

I walked across the street, checking for traffic _over_ my shoulders.

Commonsense

I use _knife_ and fork to eat steak.

Time

Ruth Bader Ginsburg was born in _1933_.

Location

University of Washington is located at _Seattle_, Washington.

Math

I was thinking about the sequence that goes 1, 1, 2, 3, 5, 8, 13, 21, _34_.

Chemistry

Sugar is composed of carbon, hydrogen, and _oxygen_.

Why Pretraining?

Pretraining Extracts Patterns, Structures, and Semantic Knowledge from Raw Texts

Outline

- Overview of LM Pretraining
- Pretraining Data
- Training Setups
- Scaling Law

What Matters for Pretraining Data?

- **Quantity:** How much data do I have?
- **Quality:** Is it beneficial for training?
- **Coverage:** Does the data cover enough domains for the end task?

Scale Up Data Quantity

	Tokens of training data
Llama 1	1.4 trillion
Llama 2	1.8 trillion
Llama 3	15 trillion
Deepseek 3	15 trillion

How Large are 1T Tokens?

Physical Size (if printed)

- Average words per page: A typical page contains about 300-500 words.
- Words from 1 trillion tokens: Assuming 750 billion words, and an average of 400 words per page:
 - Total Pages: Approximately **1.875 billion pages**.

Digital Storage

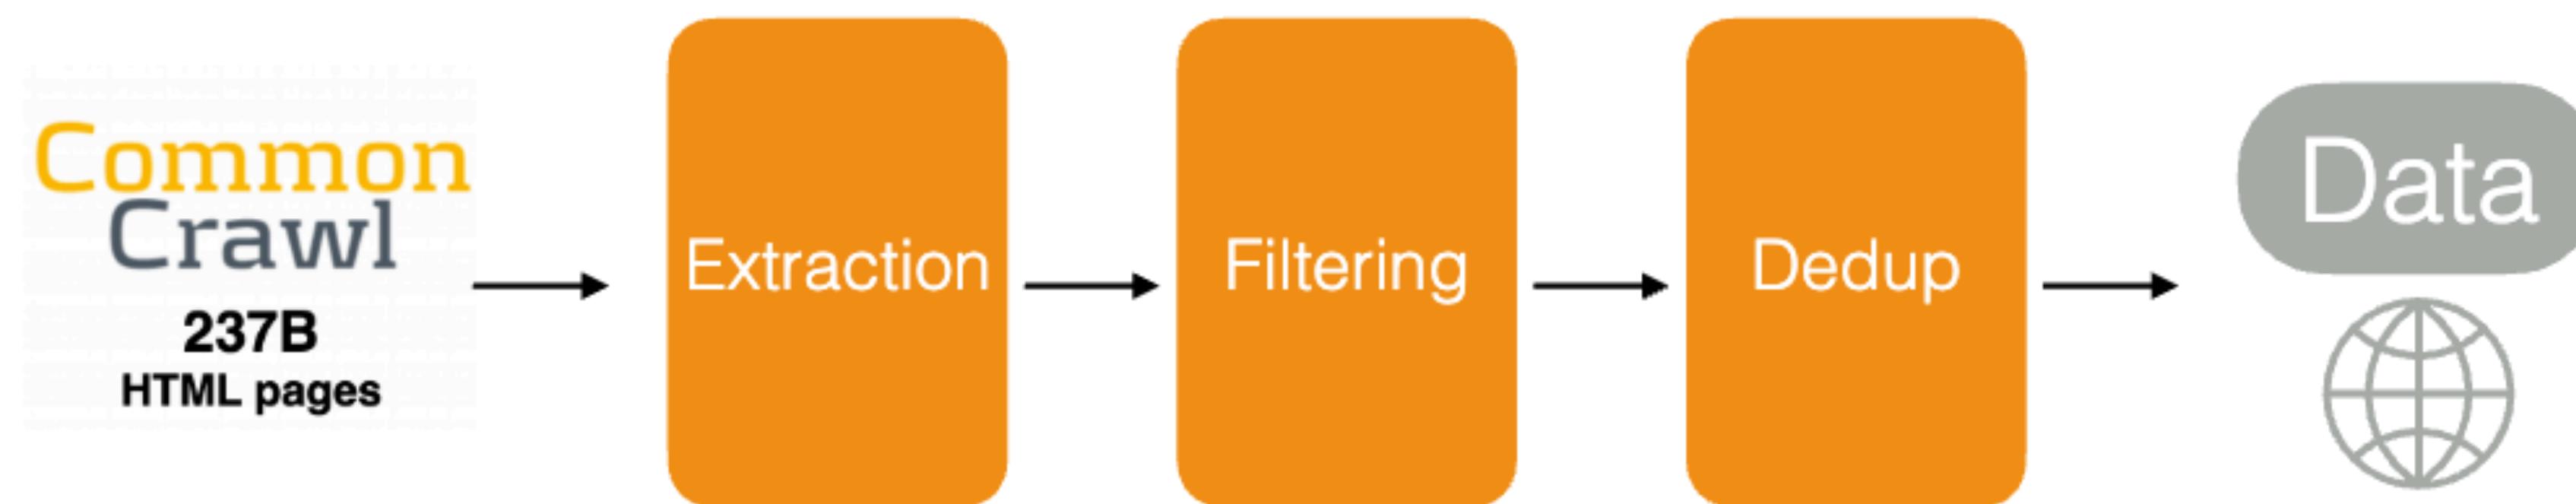
- Character Encoding: Assuming each character takes up 1 byte (in a simple encoding like ASCII), 1 trillion tokens (4 trillion characters) would require about **4 terabytes (TB)** of storage.

Reading Time

- Reading Speed: The average reading speed is about 200-250 words per minute.
- Time to Read 750 Billion Words: At 200 words per minute, it would take about 3.75 billion minutes, or approximately **7,125 years** of continuous reading.

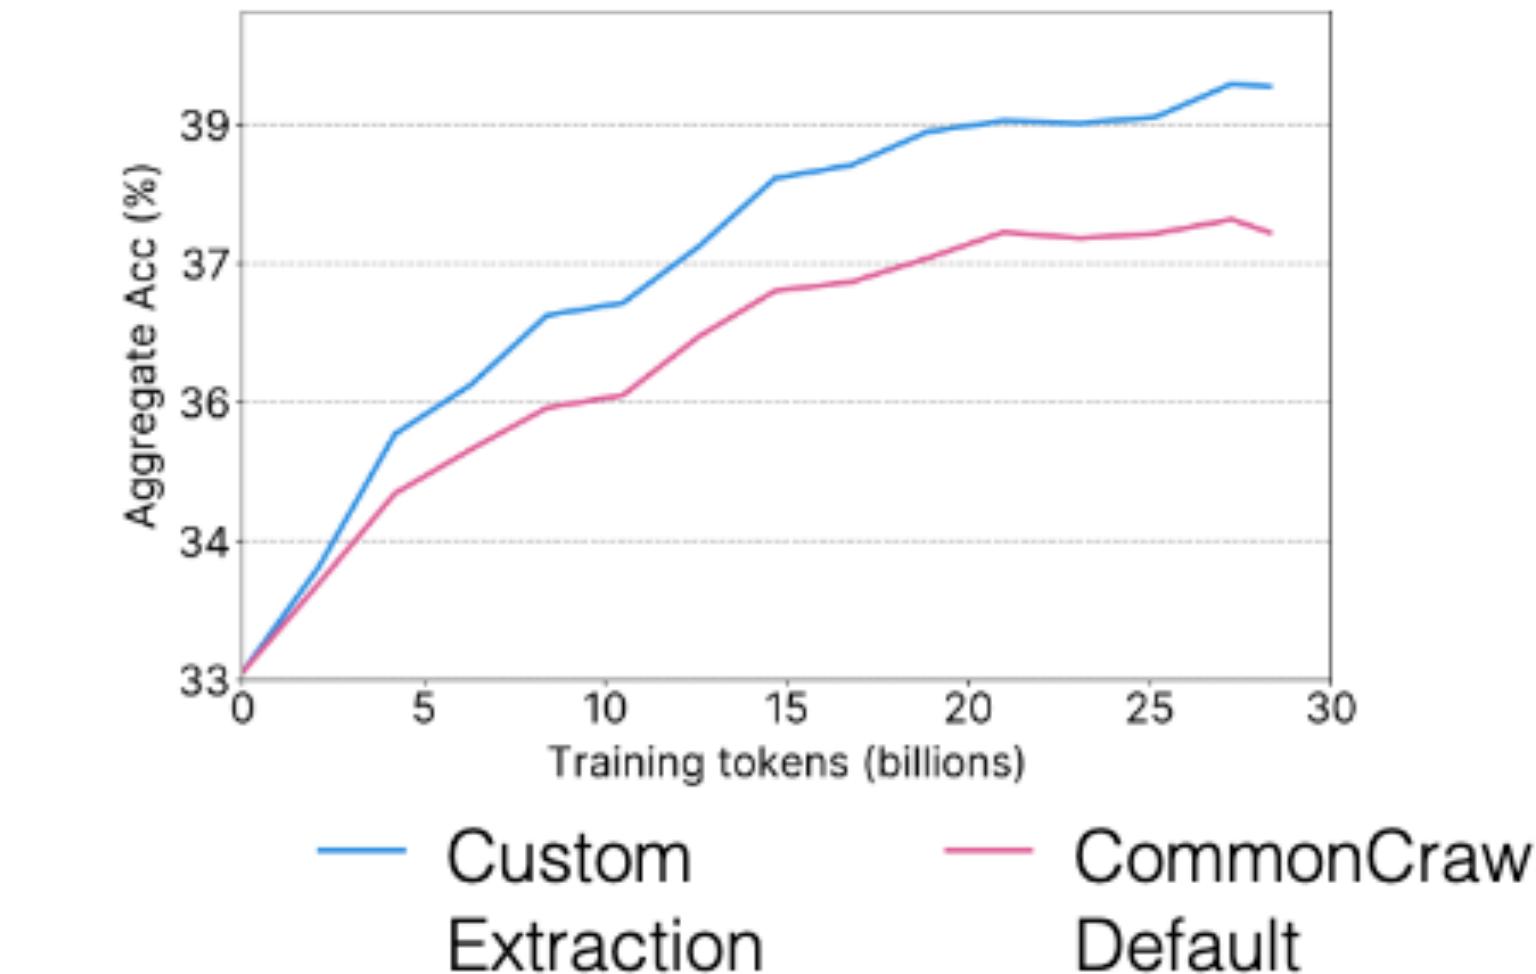
Pretraining Data Comes From Web

- Large snapshots of web pages.
 - Extraction: HTML to text
 - Filtering: filter out unwanted pages
 - Deduplication: many duplicate web pages



Extraction

- Extraction: HTML to text
 - Remove boilerplate
 - Retain Latex, code, etc.



Penedo et al 2024

This paper concerns the quantity
 defined as the length of the longest subsequence of the numbers from

Suppose I have a smooth map
$$f: \mathbb{R}^3 \rightarrow S^2$$
. If I identify \mathbb{R}^3 with $U_S = S^3 - \{(0,0,1)\}$ via stereographic projection

```
<math>
<semantics>
...
<annotation ...>
  {\displaystyle \mathrm{MA} \\
   =\frac{f_0}{f_E}}
</annotation>
</semantics>
</math>
```

Image Equations

Delimited Math

Special Tags

Filtering

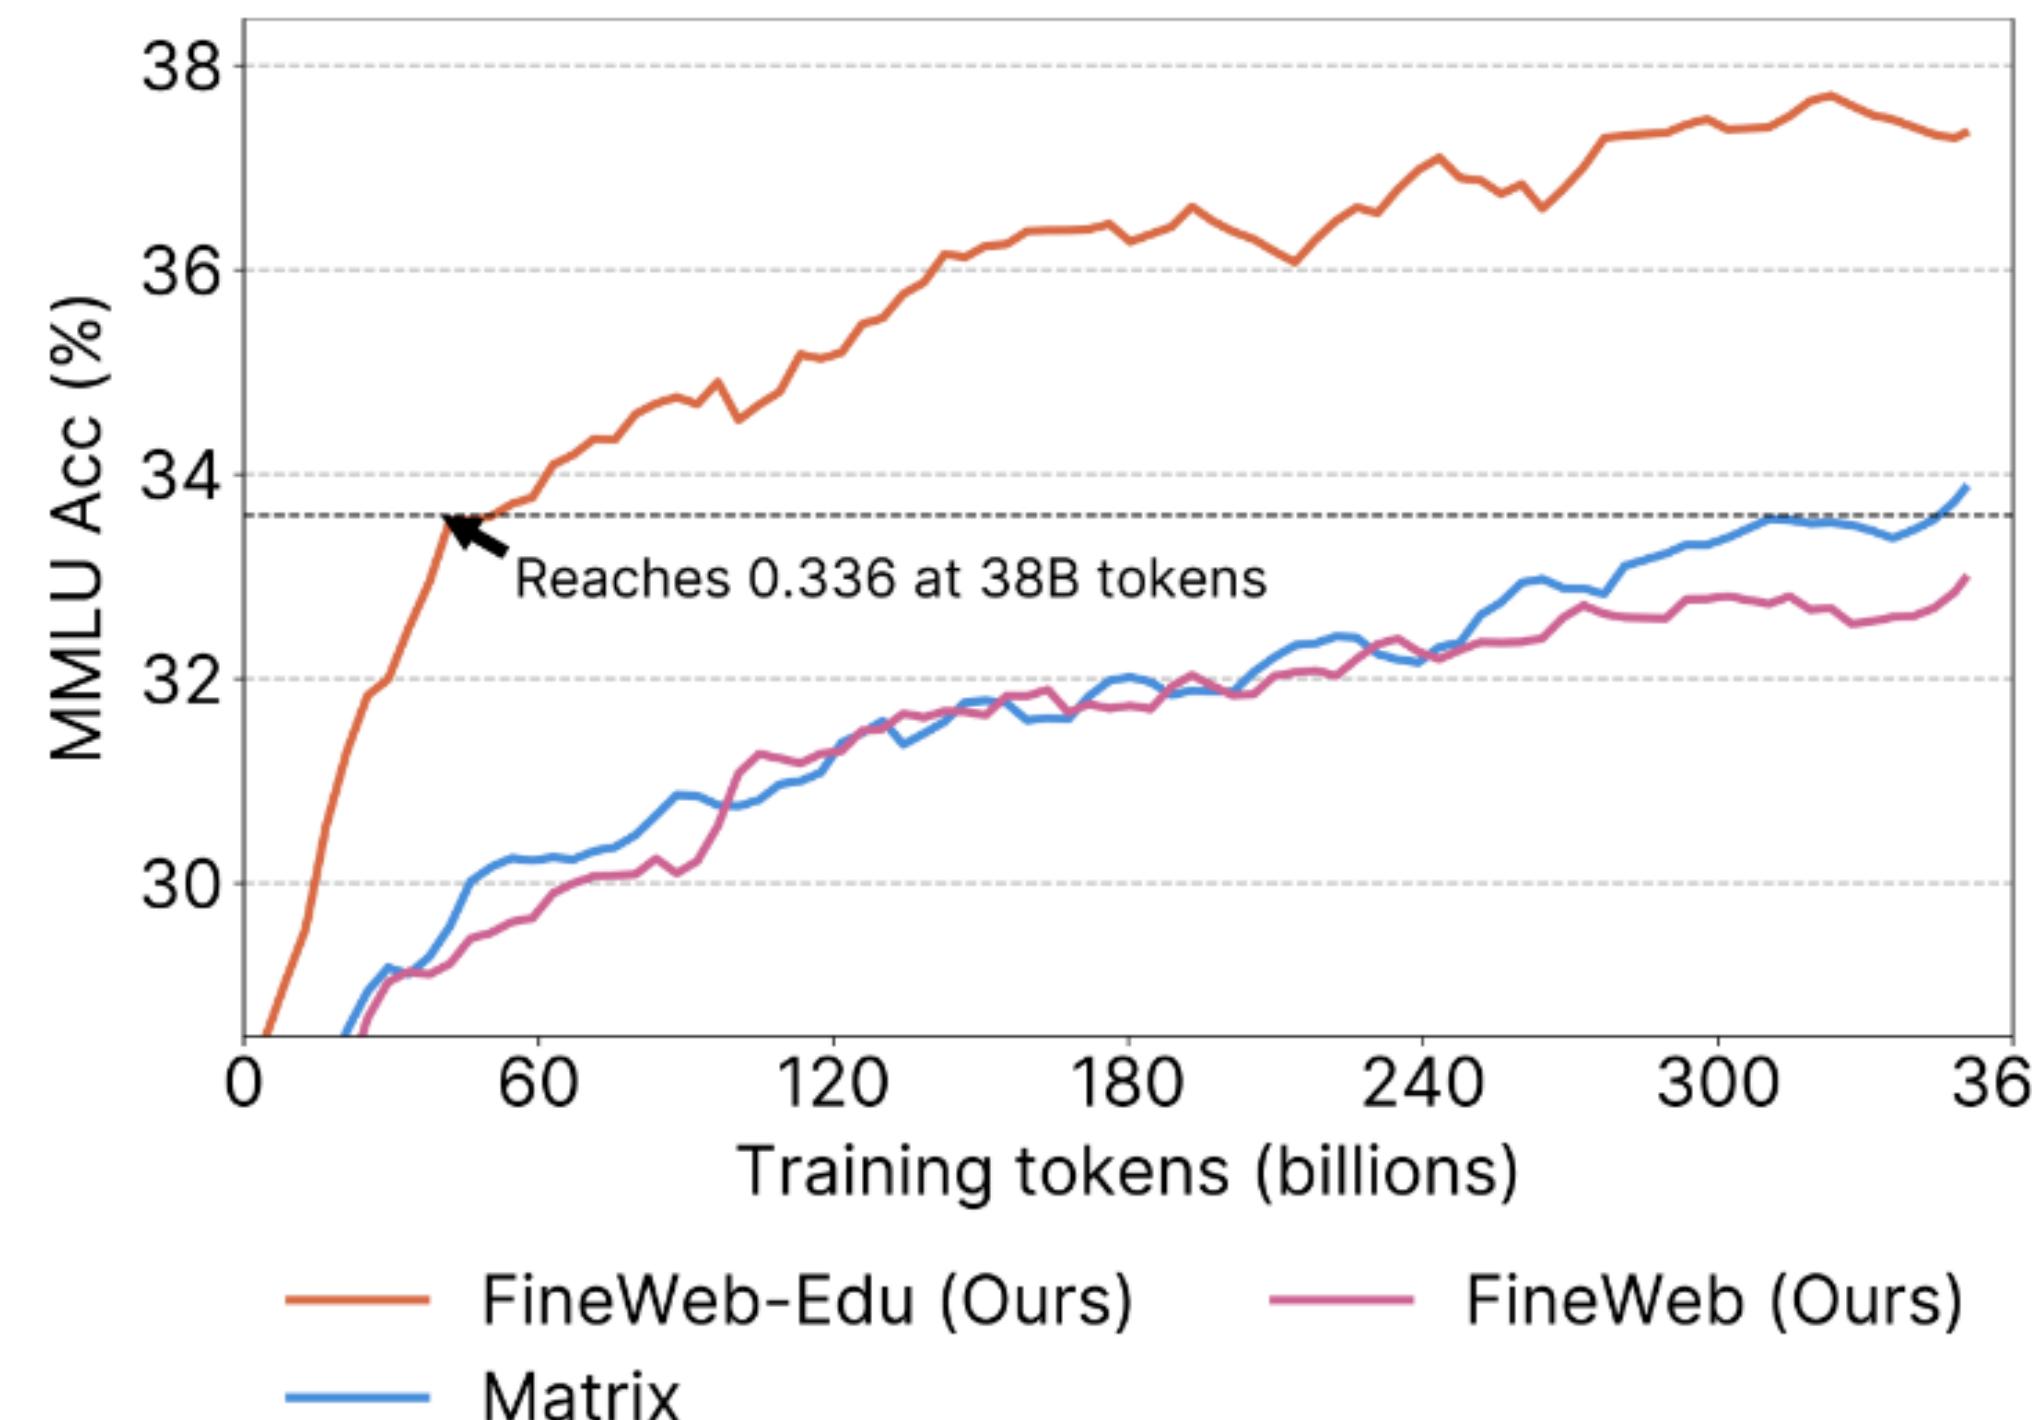
- Filter out unwanted text
 - Language filter
 - Repetitions
 - Too many short lines
 - ...

Deduplication

- Remove duplicate content
 - Fuzzy strategy: *minhash*

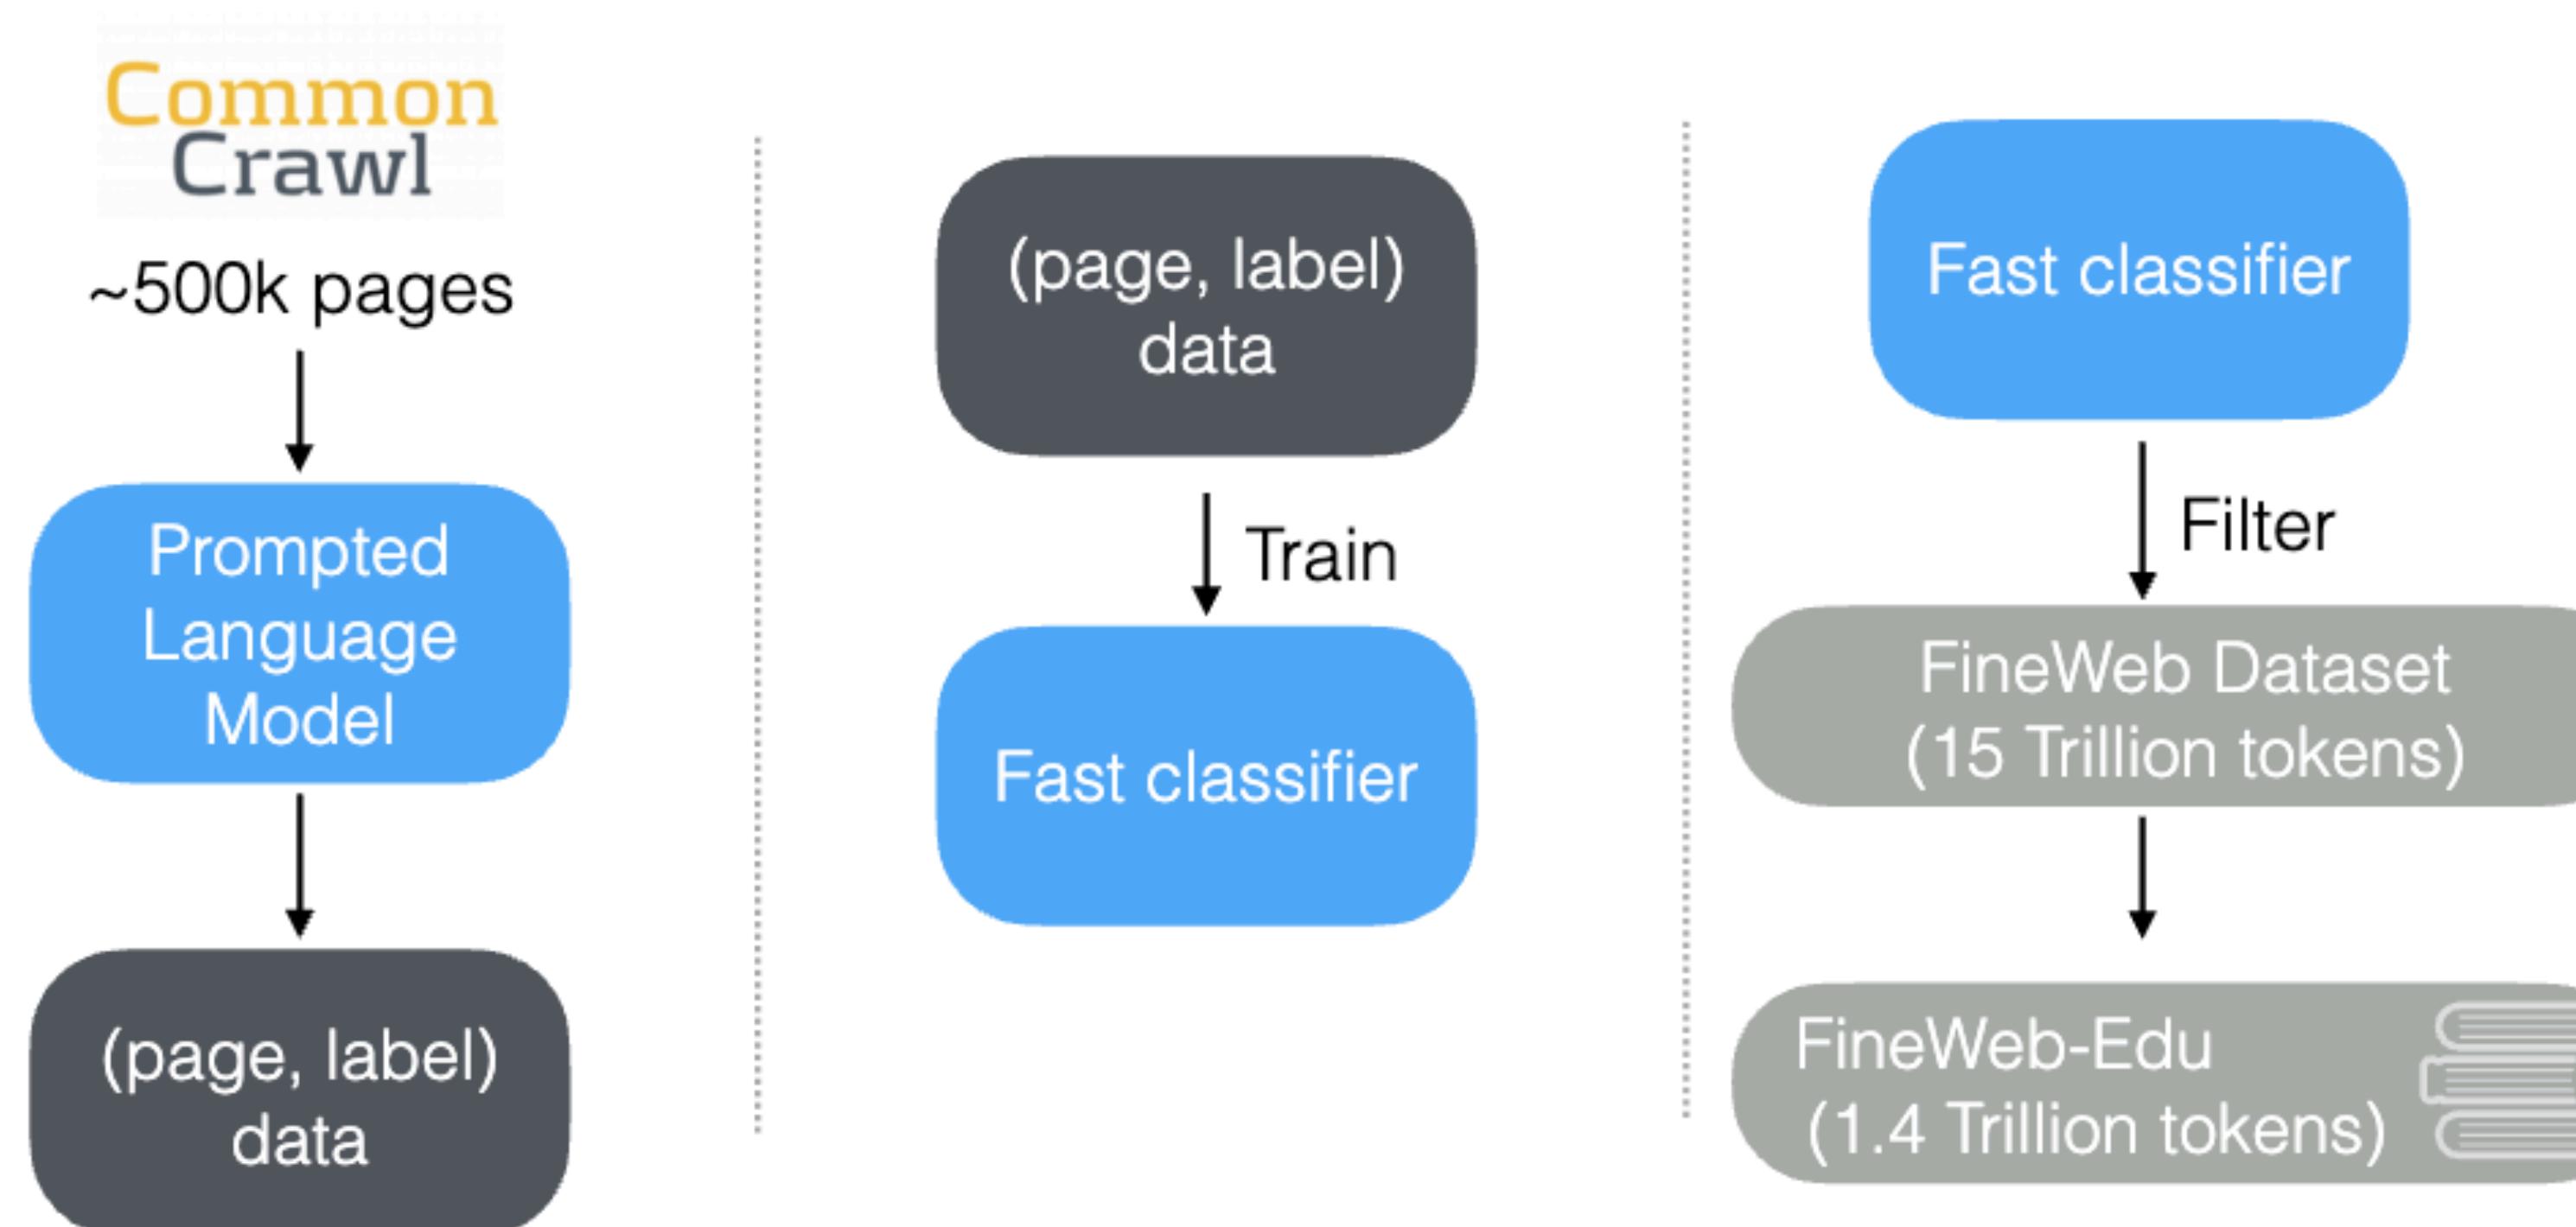
Data Quality: Model-based Selection

- Example: **FineWeb-Edu** [Penedo et al 2024]
 - Classifier to classify pages as “educational”



Data Quality: Model-based Filtering

- Example: **FineWeb-Edu** [Penedo et al 2024]
 - Classifier to classify pages as “educational”



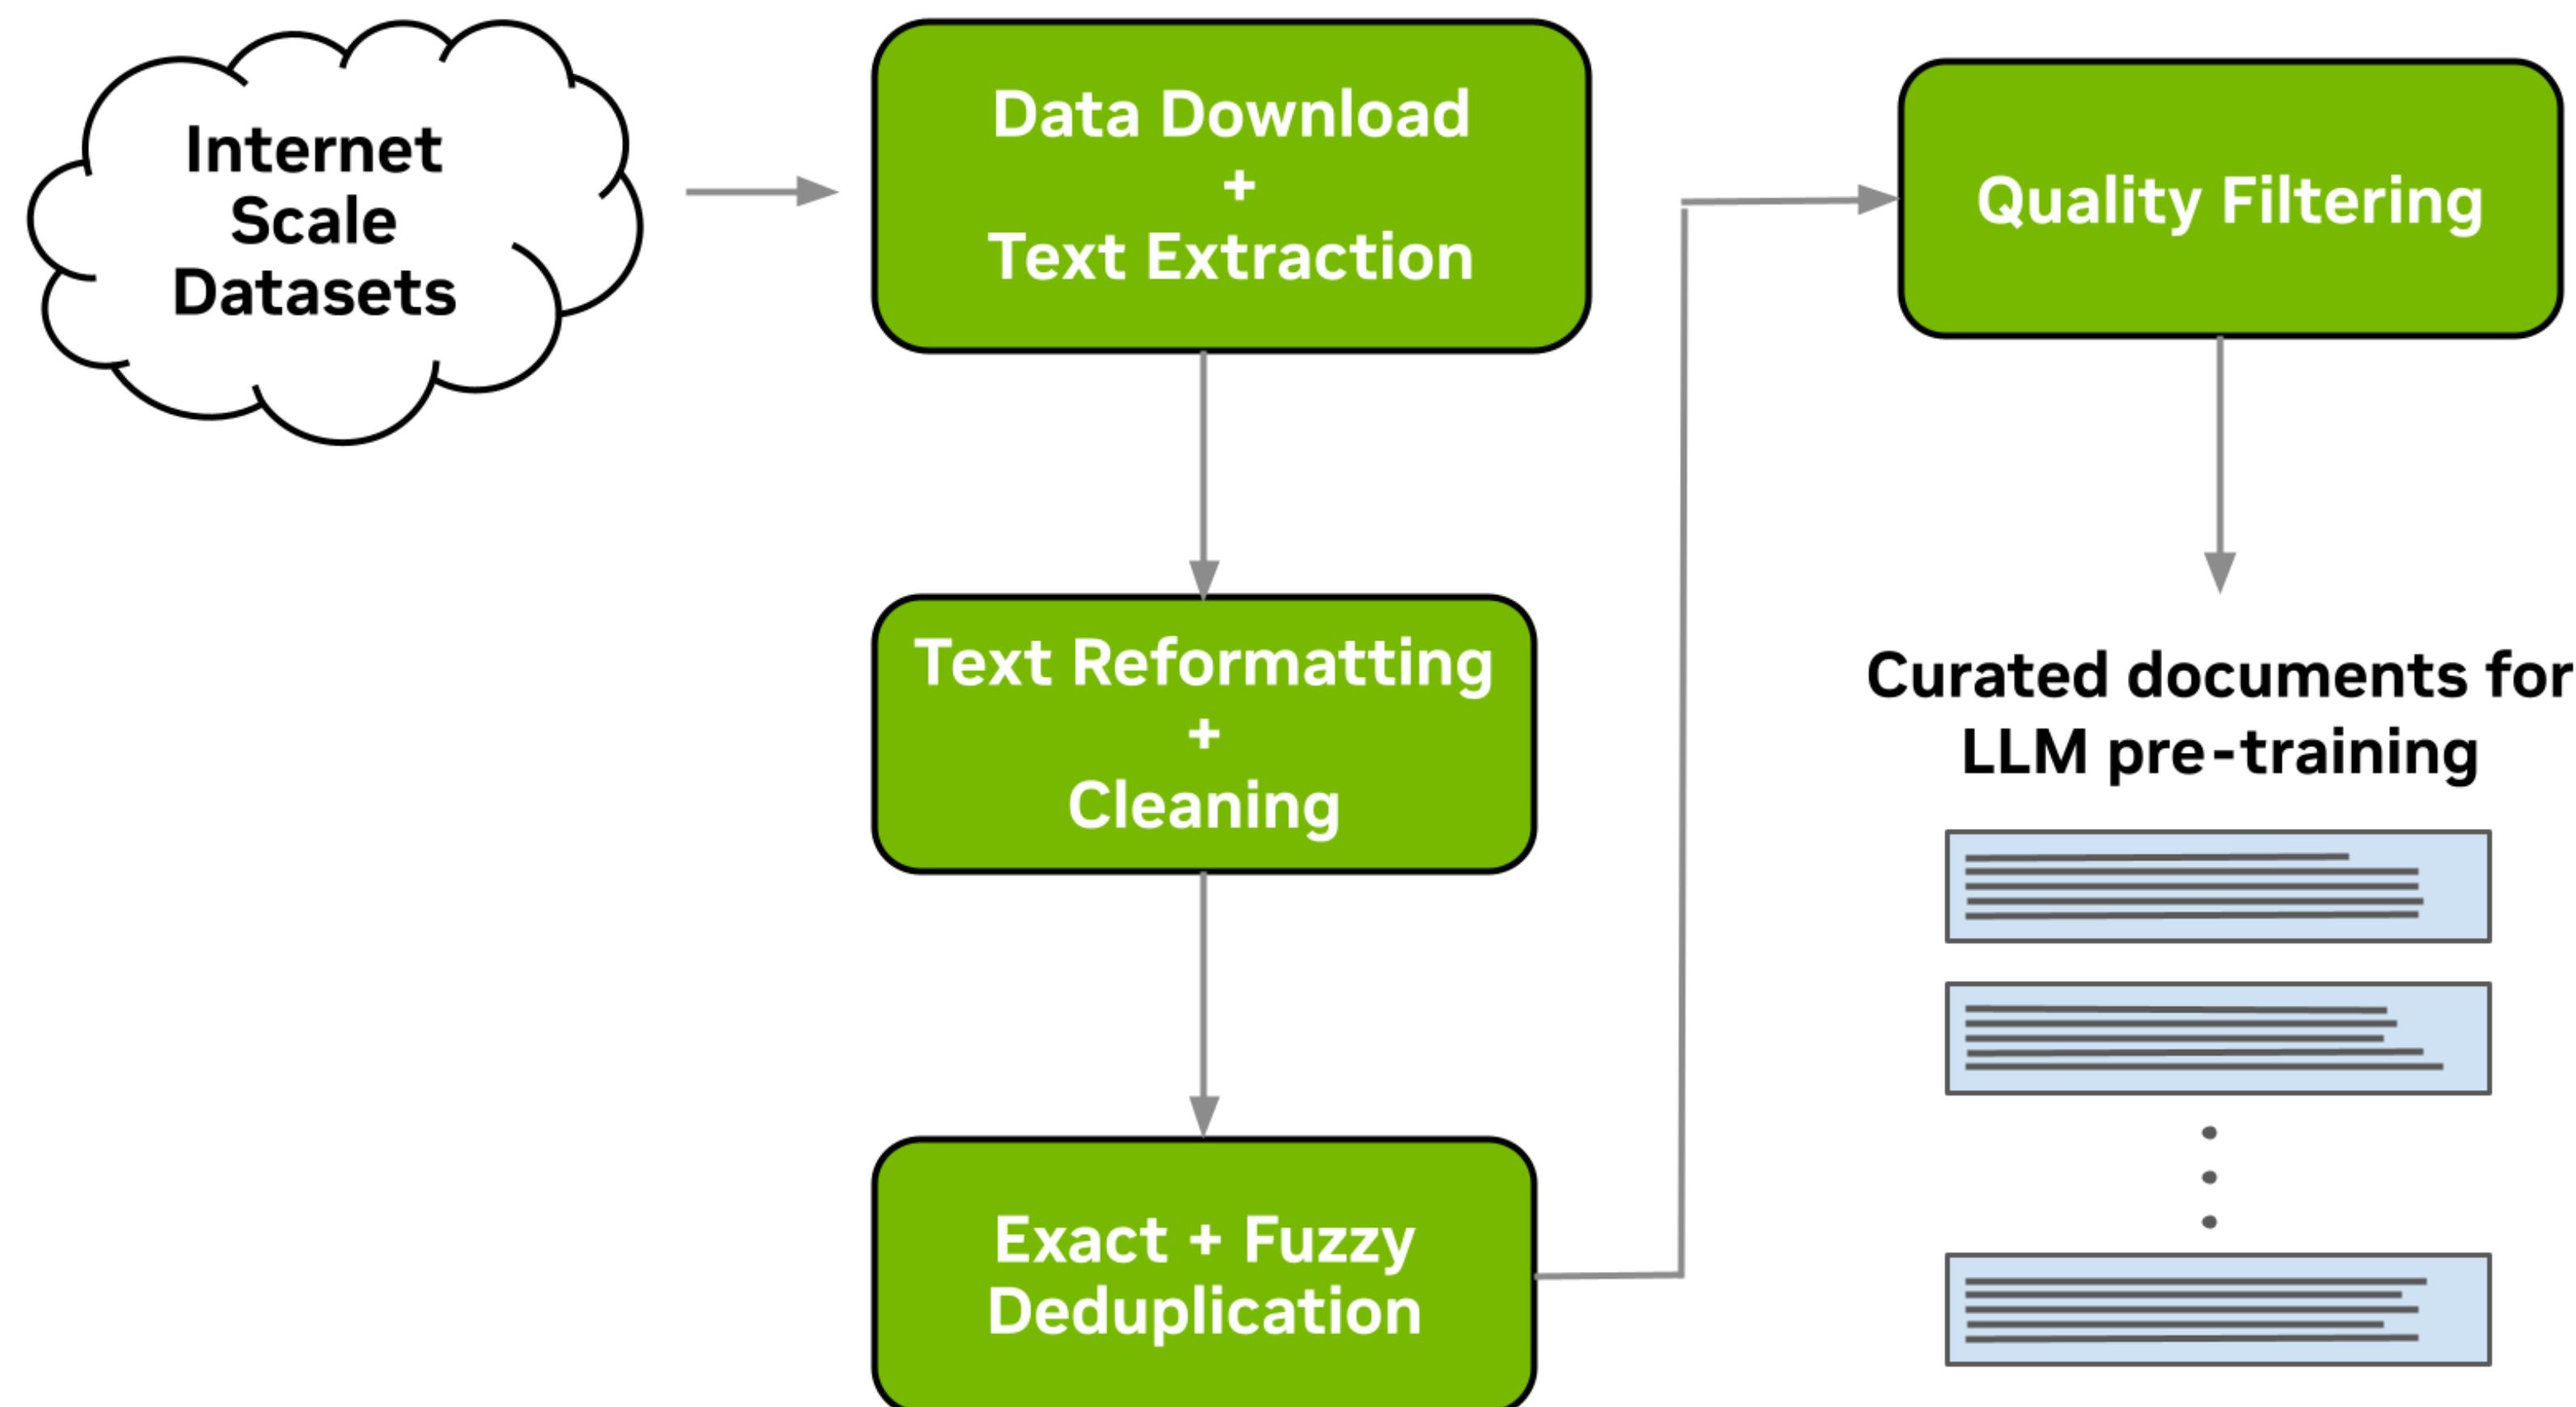
Data Coverage: Mixtures

- Training Data is a mixture of different sources

Dataset	Sampling prop.	Epochs	Disk size
CommonCrawl	67.0%	1.10	3.3 TB
C4	15.0%	1.06	783 GB
Github	4.5%	0.64	328 GB
Wikipedia	4.5%	2.45	83 GB
Books	4.5%	2.23	85 GB
ArXiv	2.5%	1.06	92 GB
StackExchange	2.0%	1.03	78 GB

Most model builders keep pretraining data private; Llama shared some details but not the data itself

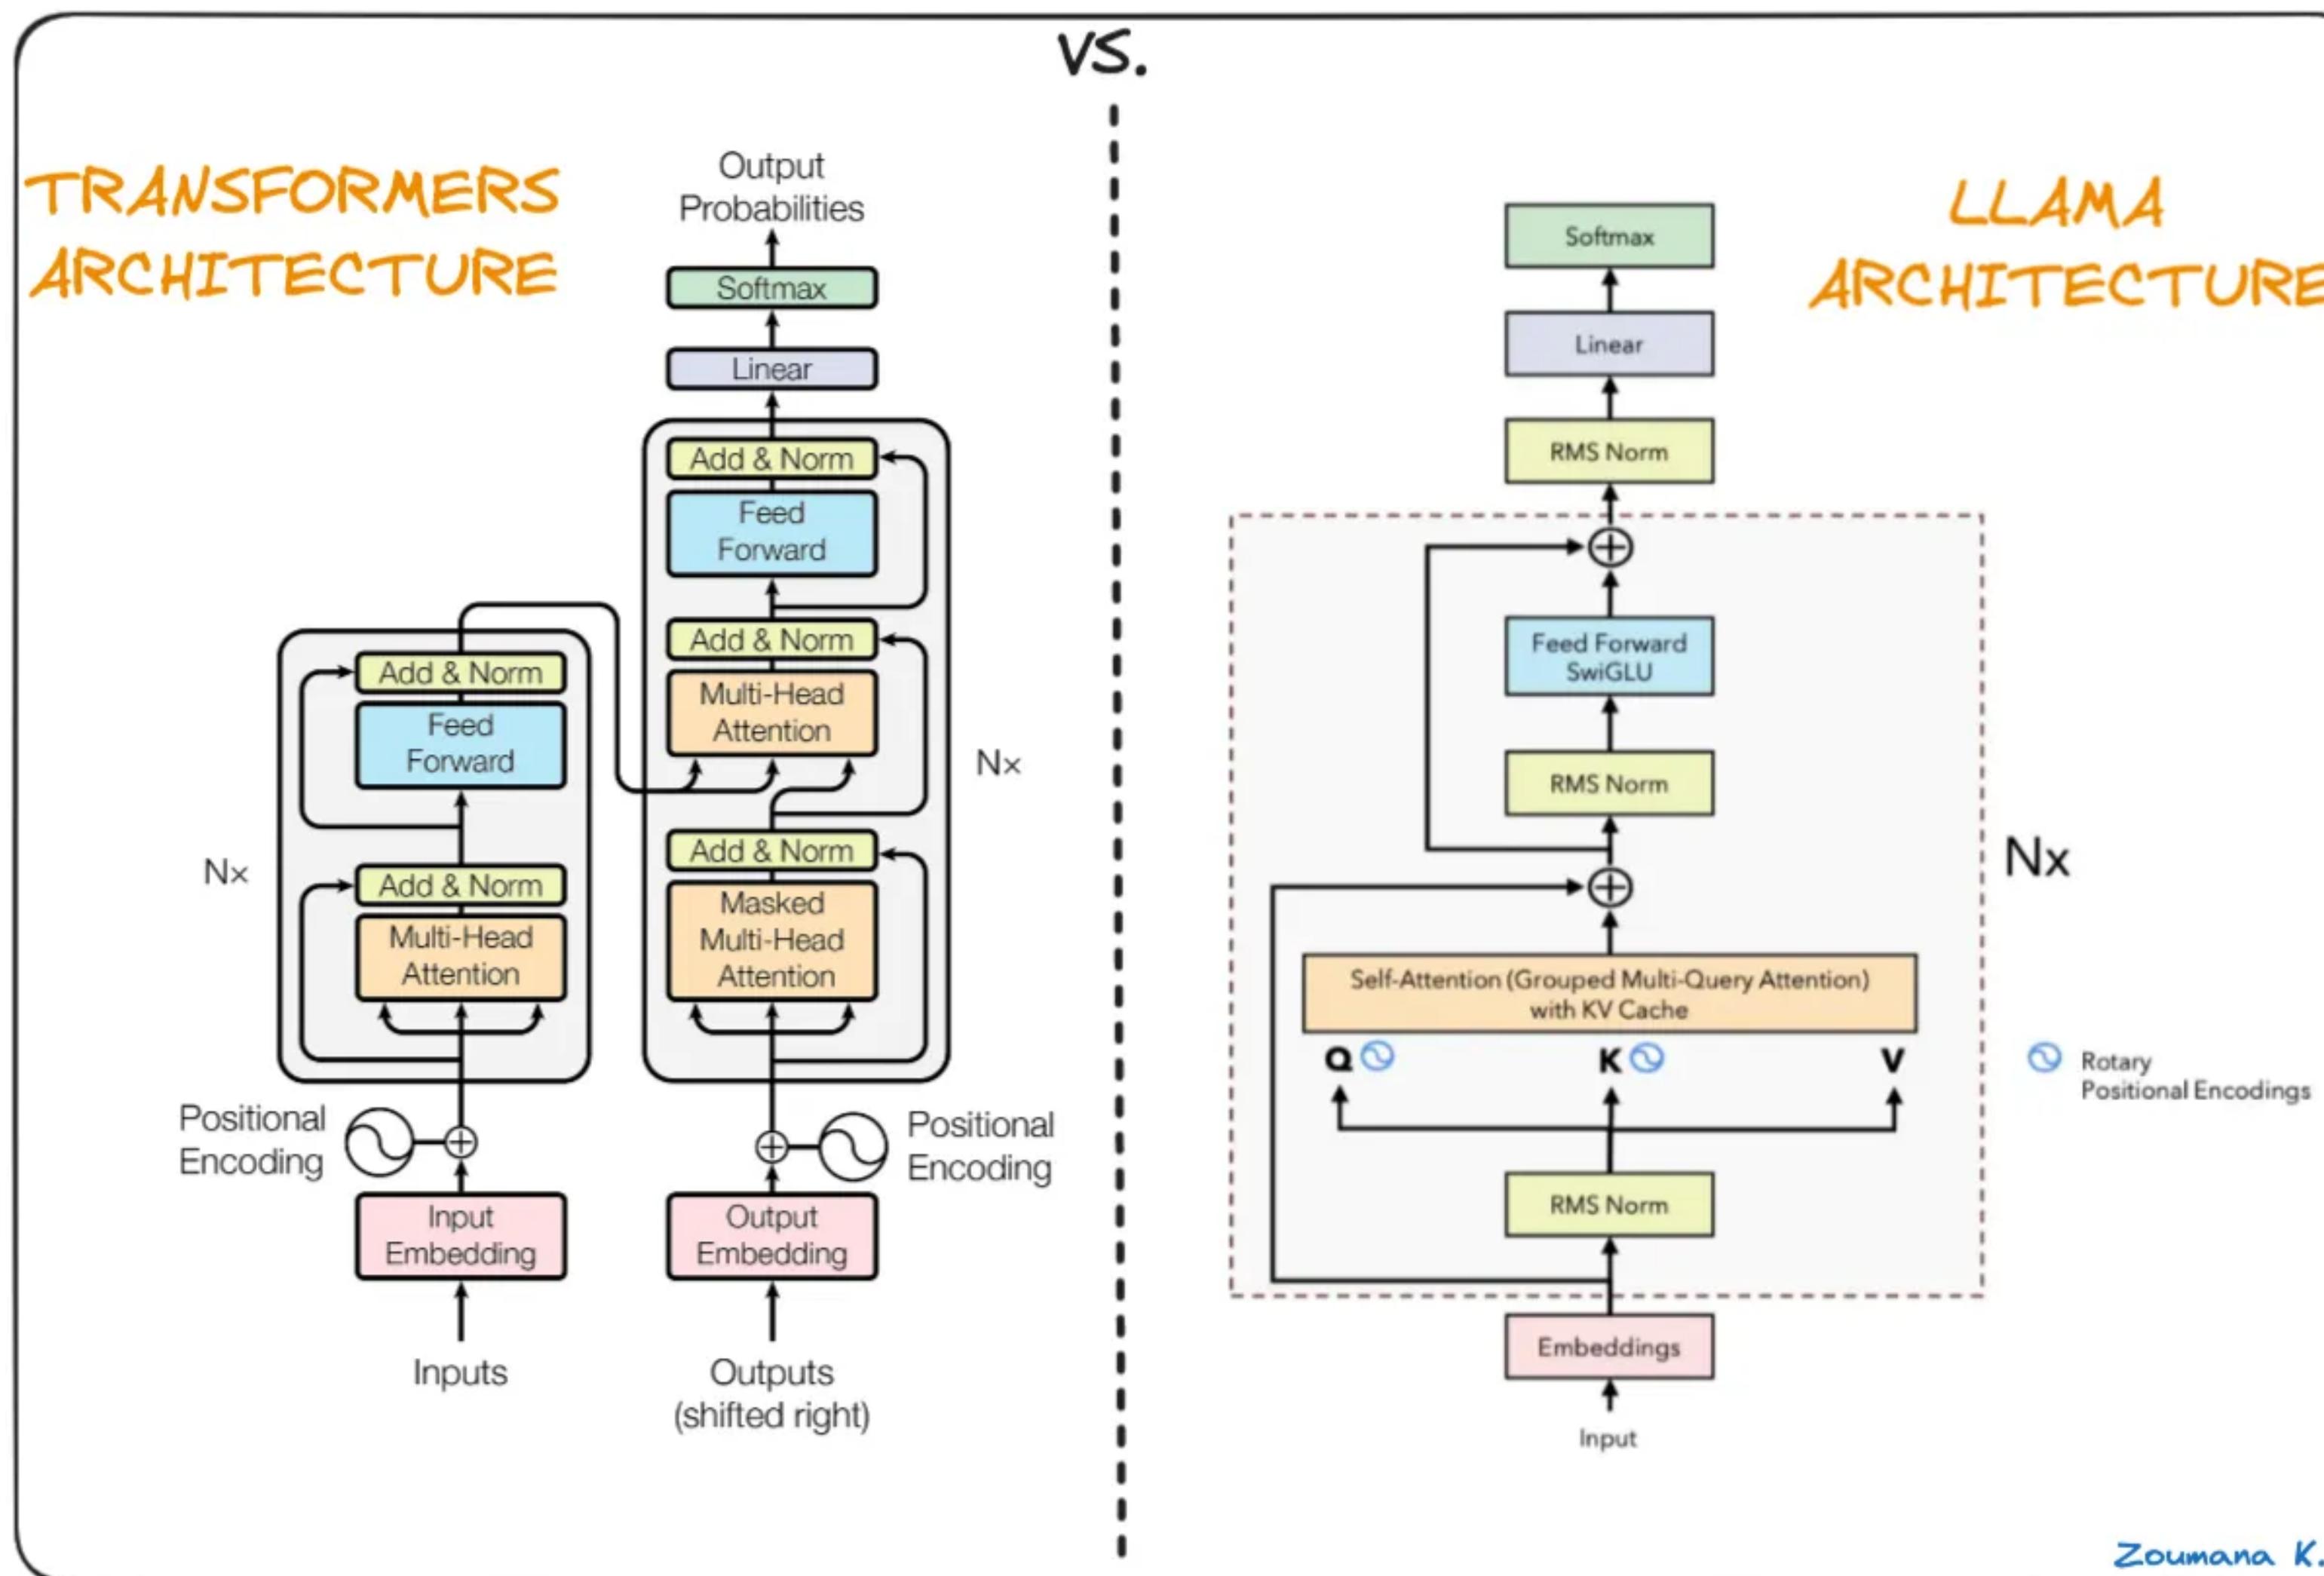
Pretraining Data Summary



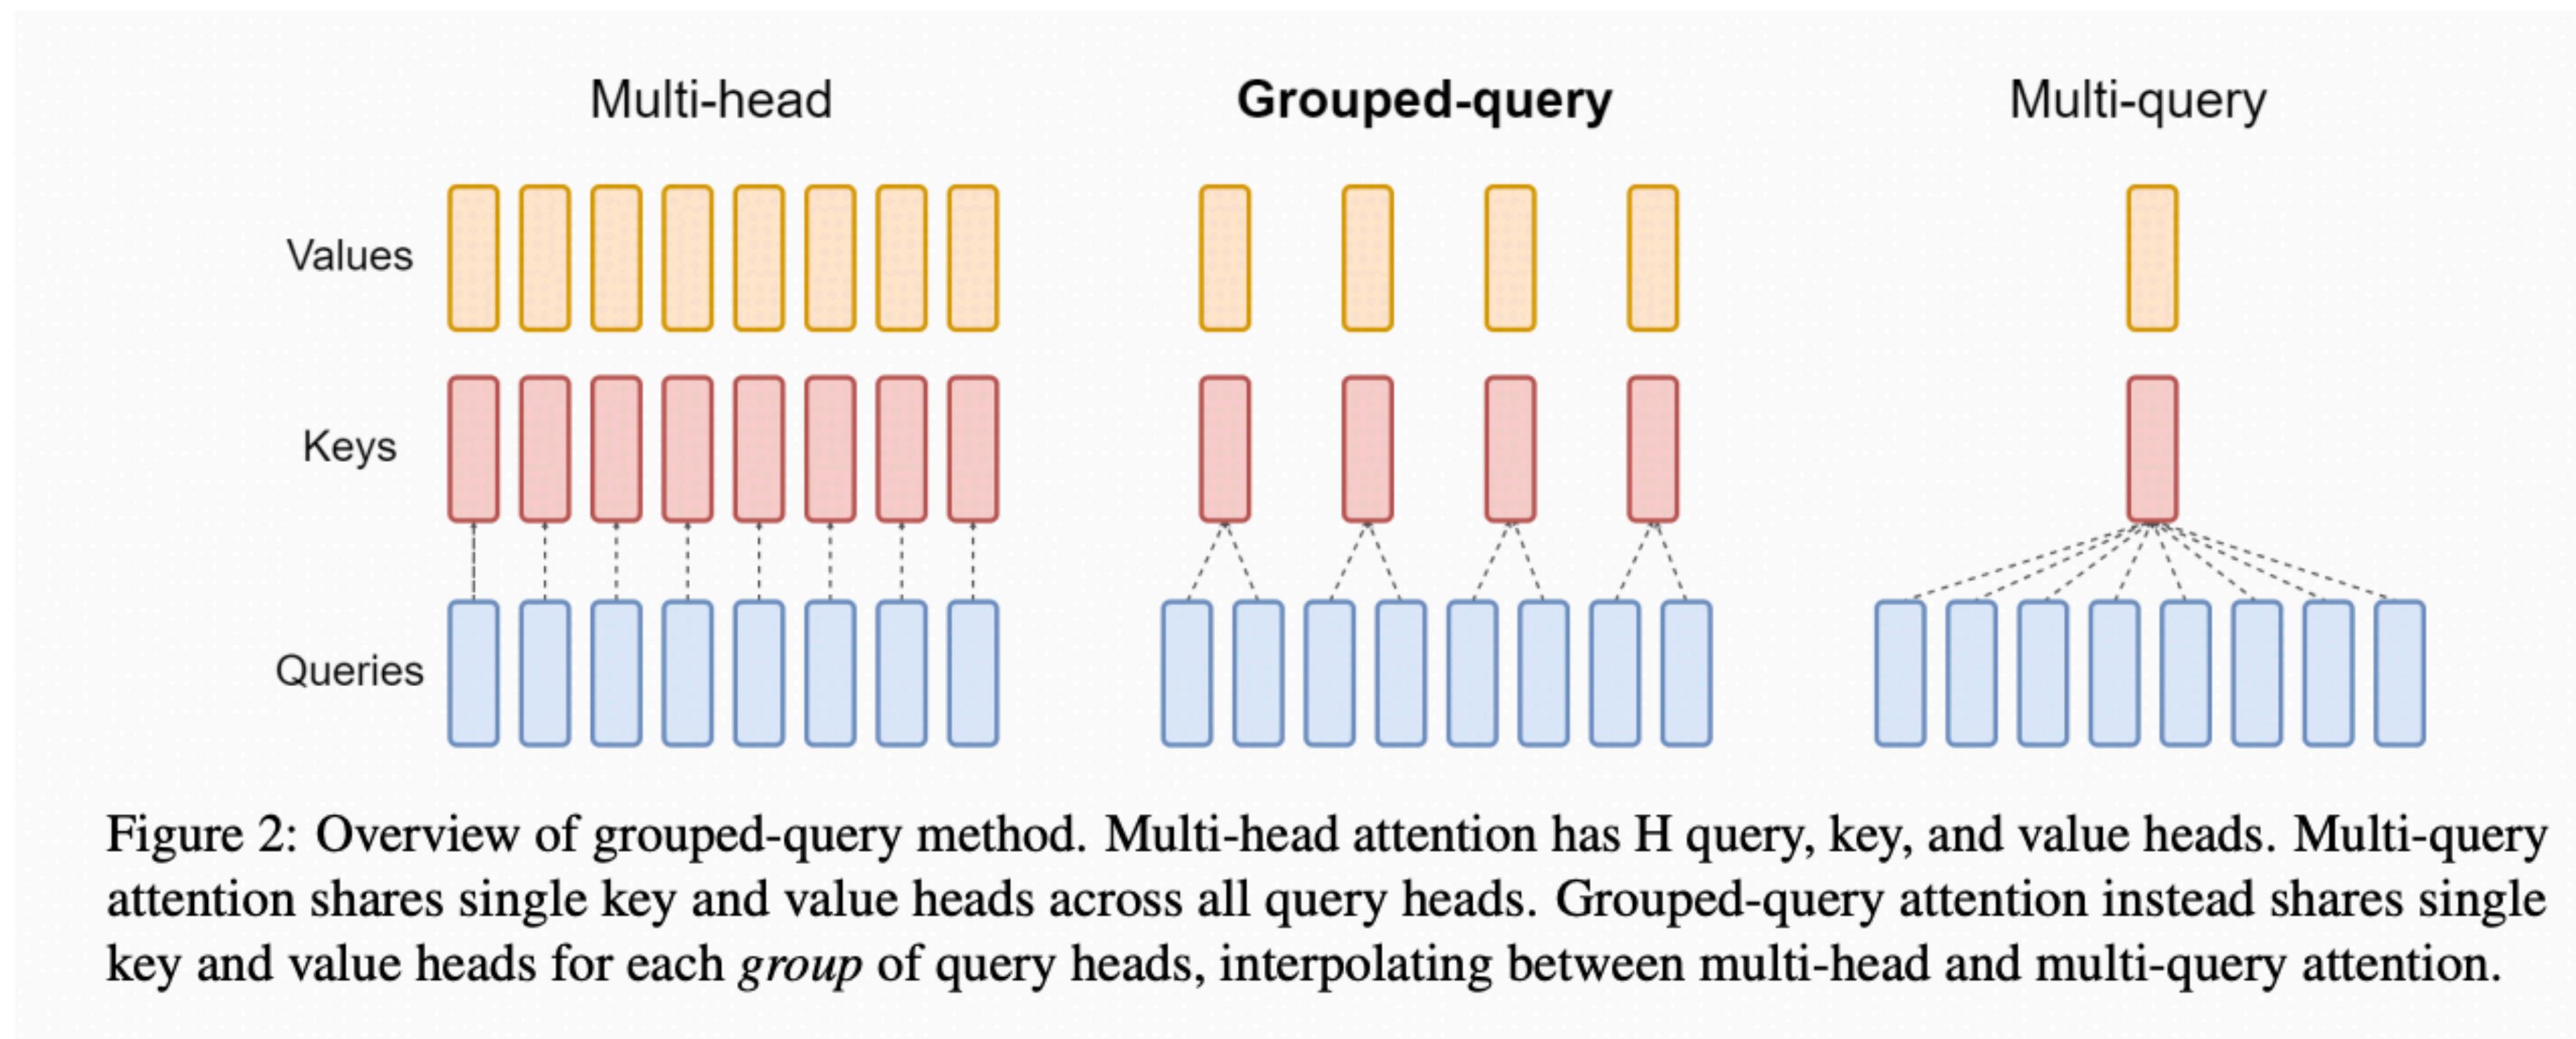
Outline

- Overview of LM Pretraining
- Pretraining Data
- Training Setups
- Scaling Law

Architecture (Recap)



Llama Architecture: Grouped Query Attention



Model	T _{infer}	Average	CNN	arXiv	PubMed	MediaSum	MultiNews	WMT	TriviaQA
	s		R ₁	BLEU	F1				
MHA-Large	0.37	46.0	42.9	44.6	46.2	35.5	46.6	27.7	78.2
MHA-XXL	1.51	47.2	43.8	45.6	47.5	36.4	46.9	28.4	81.9
MQA-XXL	0.24	46.6	43.0	45.0	46.9	36.1	46.5	28.5	81.3
GQA-8-XXL	0.28	47.1	43.5	45.4	47.7	36.3	47.2	28.4	81.6

Llama Architecture: Other Setups

params	dimension	n heads	n layers	learning rate	batch size	n tokens
6.7B	4096	32	32	$3.0e^{-4}$	4M	1.0T
13.0B	5120	40	40	$3.0e^{-4}$	4M	1.0T
32.5B	6656	52	60	$1.5e^{-4}$	4M	1.4T
65.2B	8192	64	80	$1.5e^{-4}$	4M	1.4T

Table 2: Model sizes, architectures, and optimization hyper-parameters.

Llama Architecture: Other Setups

params	dimension	n heads	n layers	learning rate	batch size	n tokens
6.7B	4096	32	32	$3.0e^{-4}$	4M	1.0T
13.0B	5120	40	40	$3.0e^{-4}$	4M	1.0T
32.5B	6656	52	60	$1.5e^{-4}$	4M	1.4T
65.2B	8192	64	80	$1.5e^{-4}$	4M	1.4T

Table 2: Model sizes, architectures, and optimization hyper-parameters.

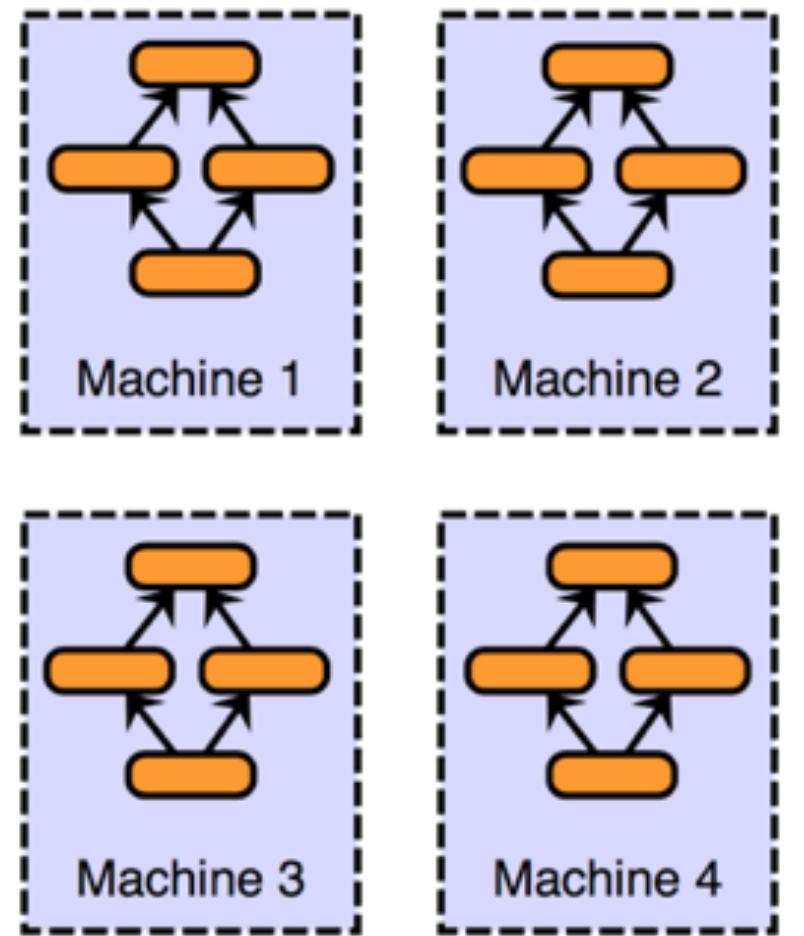
Cannot fit with 1 GPU's memory. How to train?

Bottlenecks of Training Big Models

- **Bottleneck 1:** Iteration time 
 - Each training sample takes longer to propagate through more parameters
- **Bottleneck 2:** Processor Memory 
 - Billions of parameters and 2 bytes each, GPUs have comparatively little memory!

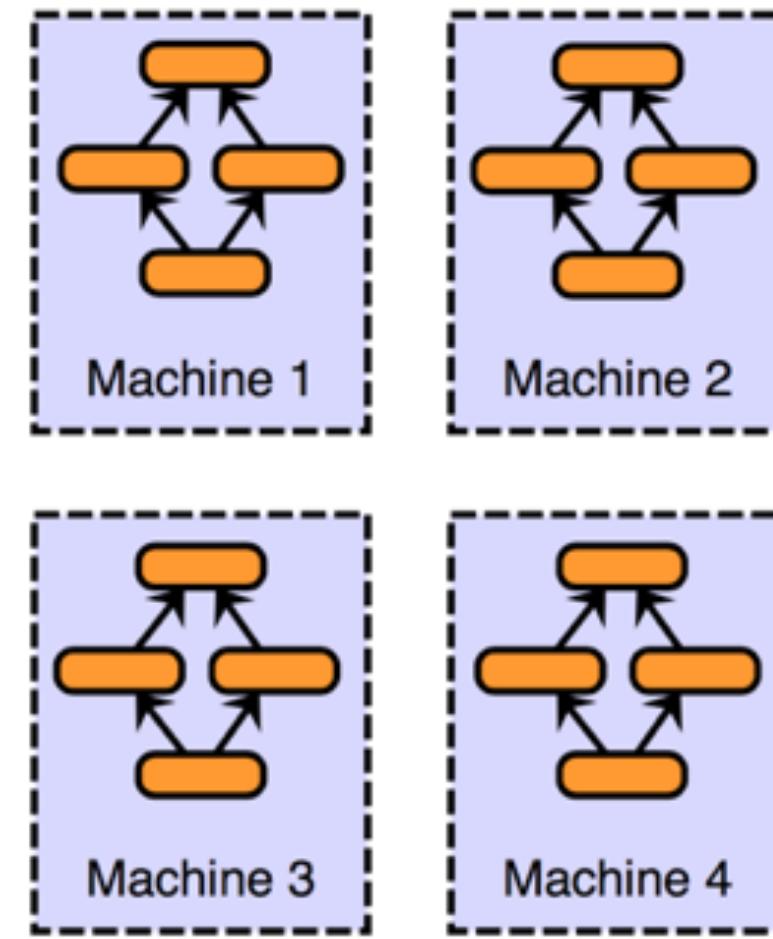
Parallelization Strategies

Data Parallelism

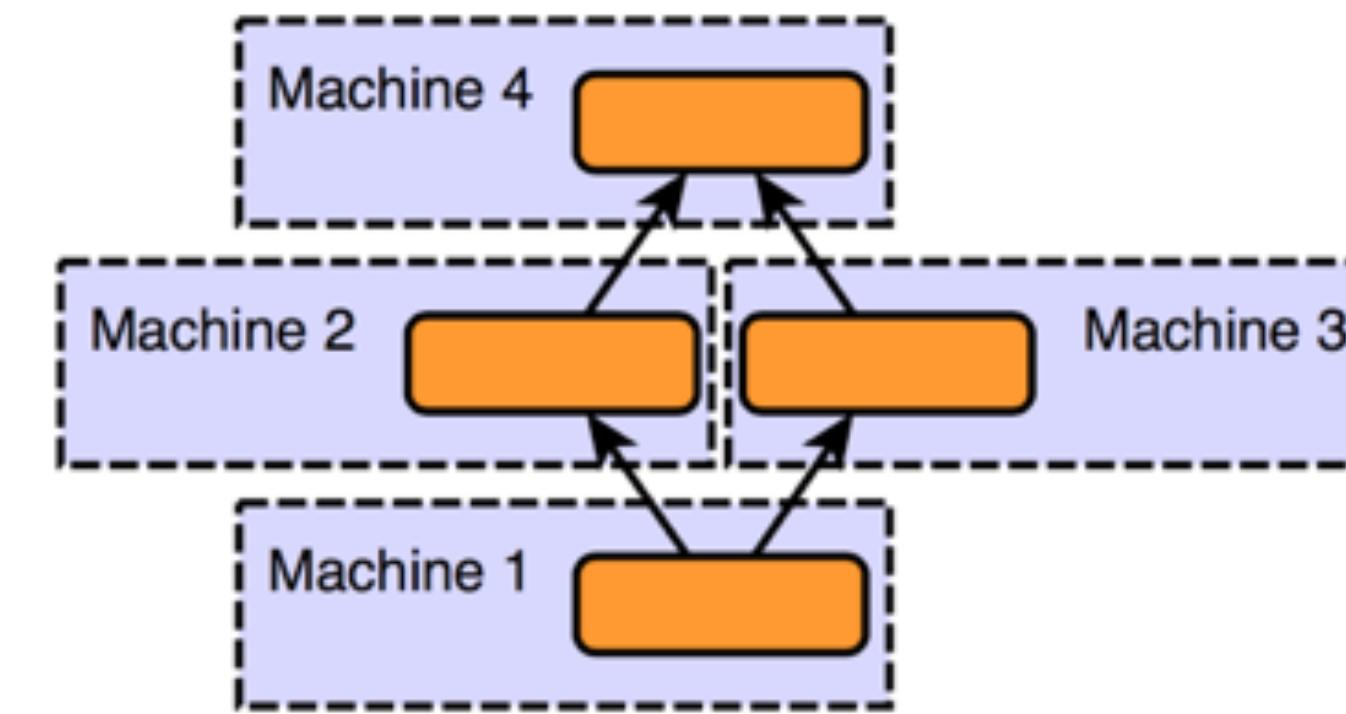


Parallelization Strategies

Data Parallelism



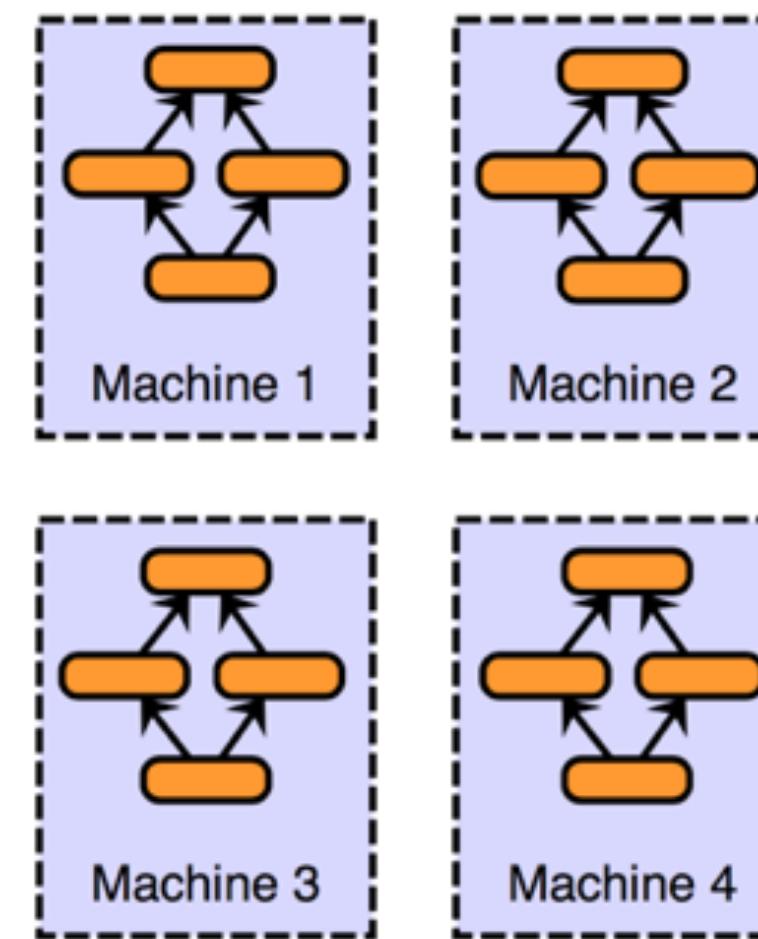
Model Parallelism



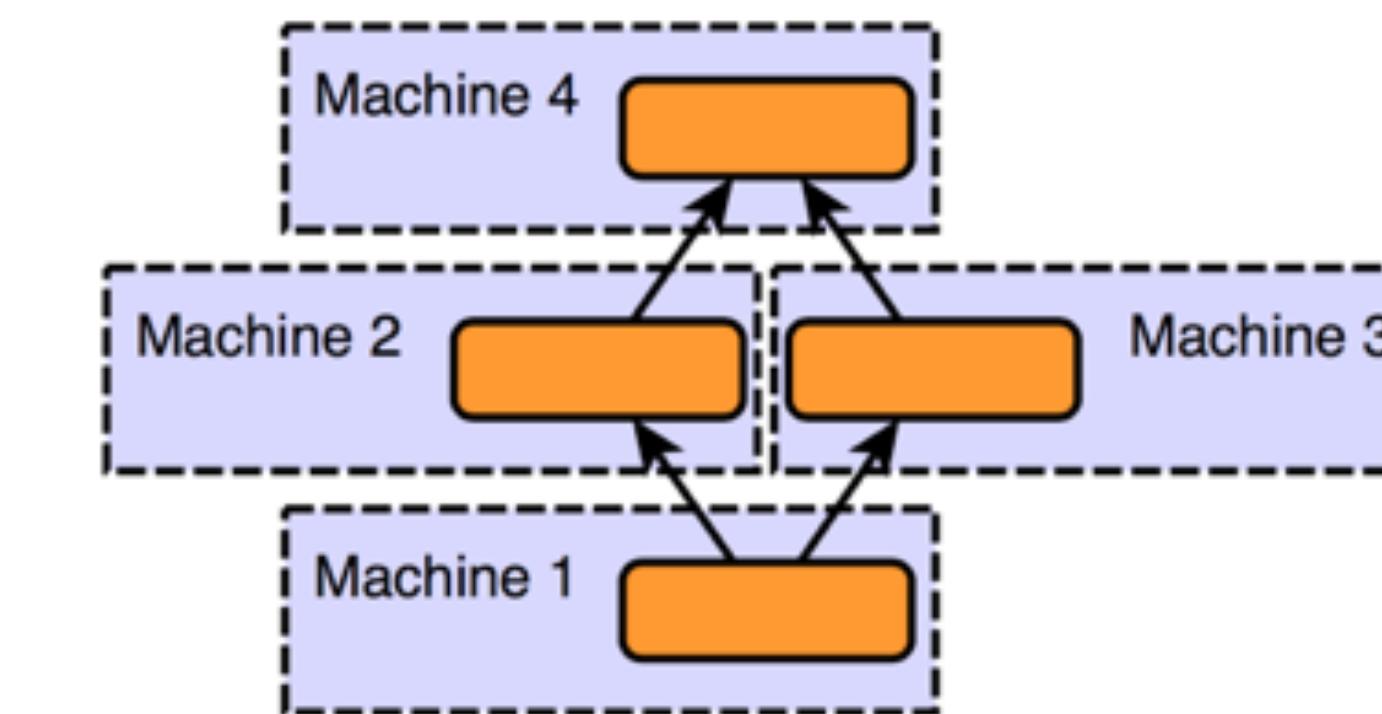
Hybrid (Model and Data) Parallelism

Parallelization Strategies

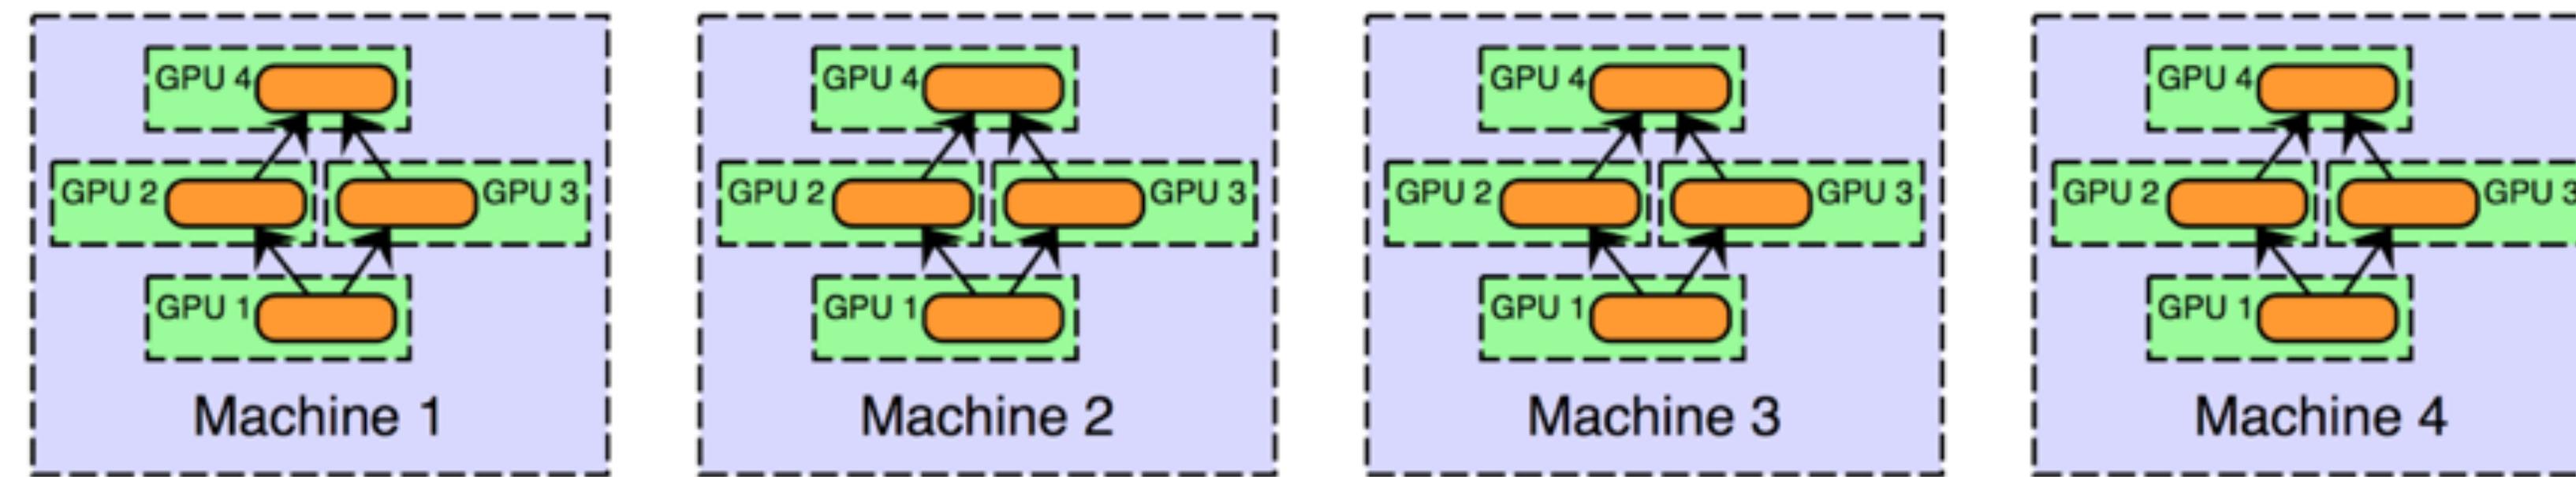
Data Parallelism



Model Parallelism

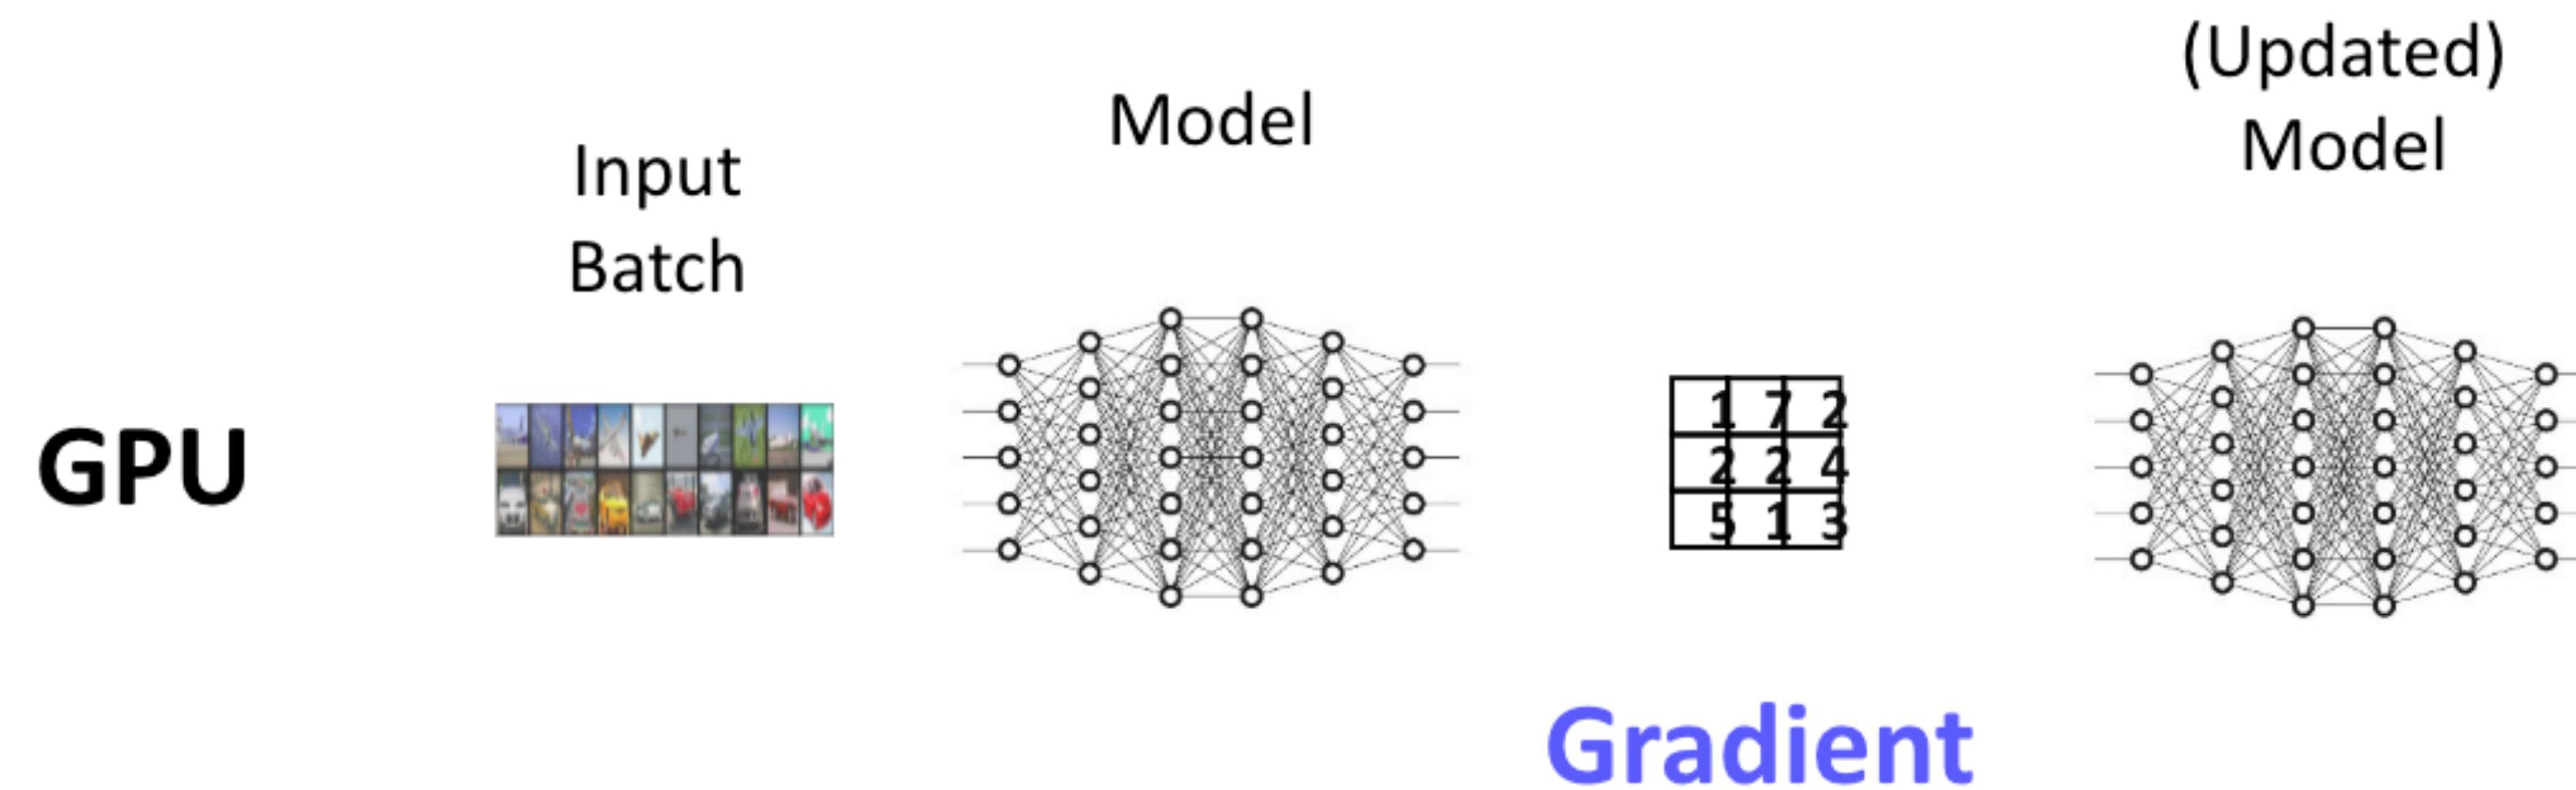


Hybrid (Model and Data) Parallelism

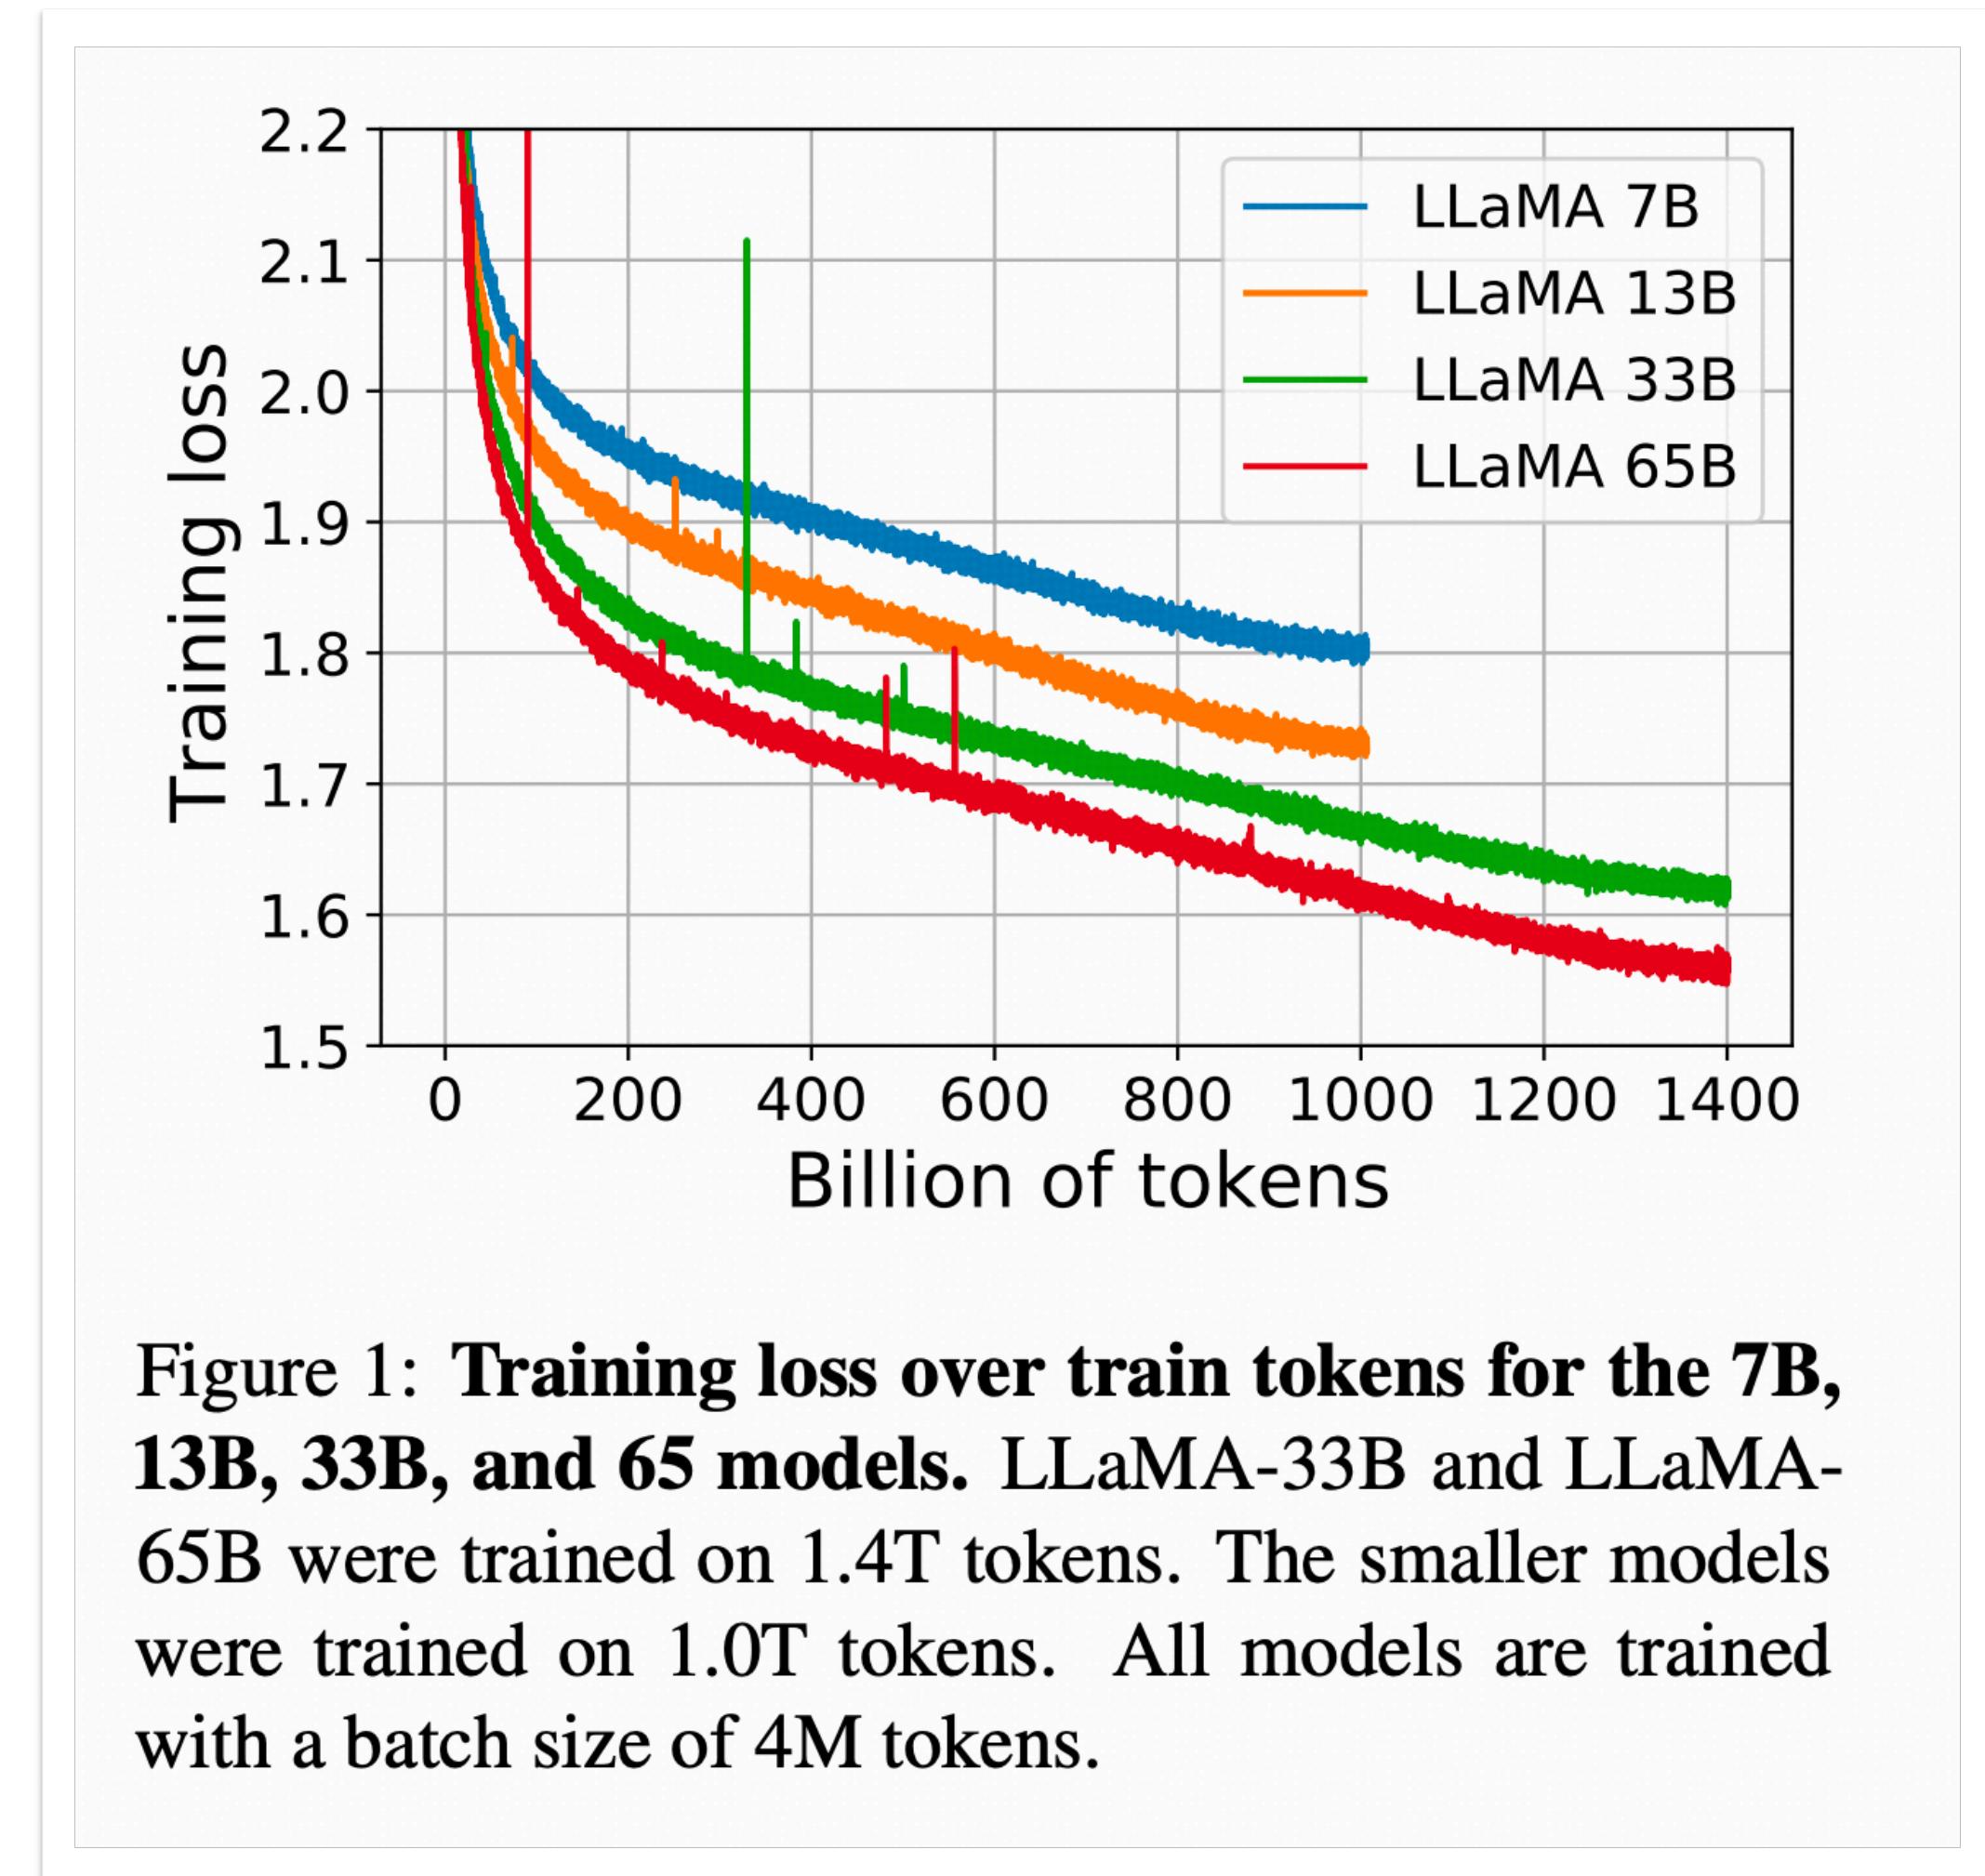


Sequential Training

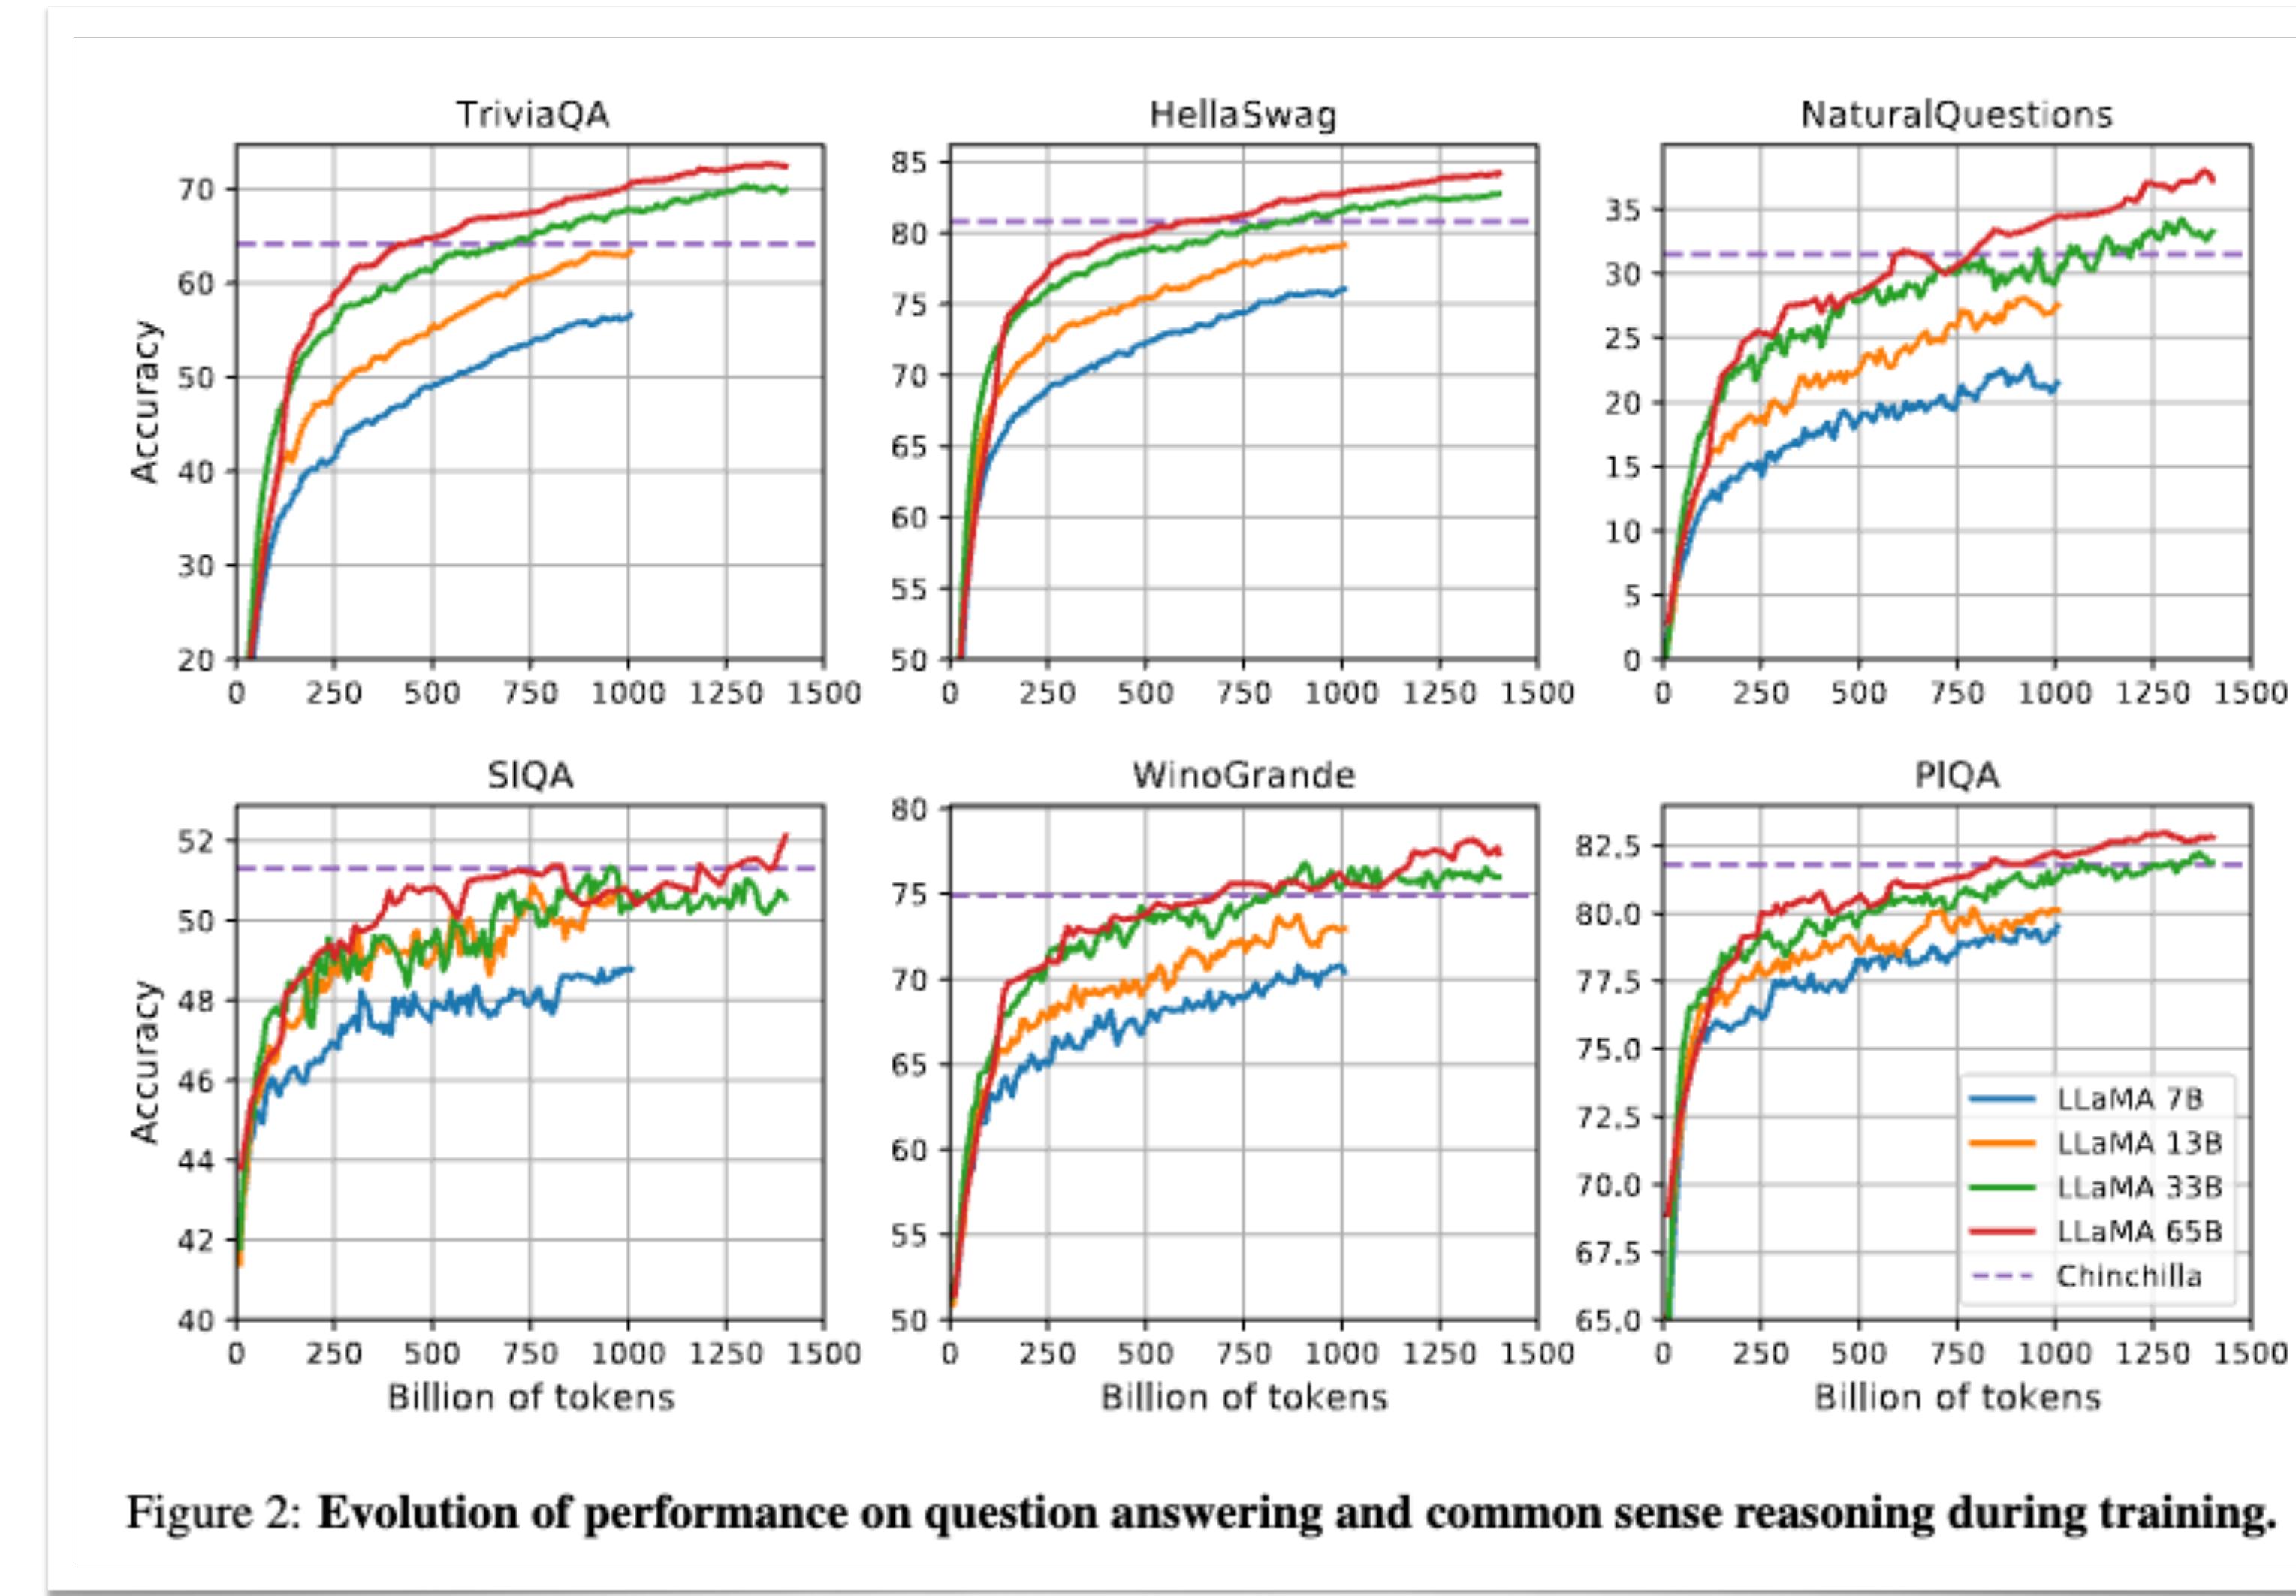
- At each training iteration:
 - Take batch size b training samples from dataset
 - Run a forward pass on each sample and compute each sample's loss
 - Run a backward pass and calculate the gradient
 - Update parameters via gradient



Loss Curve



LM Performance: More Compute, Better Results



Llama Family



# params	7B	13B	33B	65B	7B	13B	34B	70B	8B			70B
# training tokens	1T	1T	1.4T	1.4T	2T	2T	2T	2T	15T			15T
hidden embed dim	4096	5120	6656	8192	4096	5120		8192	4096			8192
# attn heads	32	40	52	64	32	40		64	32			64
# attn layers	32	40	60	80	32	40		80	32			80
attention	MHA	MHA	MHA	MHA	MHA	MHA	GQA	GQA	GQA			GQA
# kv heads	32	40	52	64	32	40		8	8			8
mlp intermediate size	11008	13824	17920	22016	11008	13824		28672	14336			28672
context	2048				4096				8192			
tokenizer	BPE sentencepiece				BPE sentencepiece				BPE tiktoken			
token vocabulary	32000				32000				128256			
fine-tuned models	-				Llama-2-Chat (Jul 2023) Code Llama (Aug 2023)				Llama-3-Instruct (Apr 2024)			

BPE: Byte Pair Encoding

■ Not released by Meta

MHA: Multi-Head Attention

GQA: Grouped-Query Attention

Try Training Llama By Yourself



github.com/facebookresearch/lingua

Meta Lingua

Mathurin Videau*, Badr Youbi Idrissi*, Daniel Haziza, Luca Wehrstedt, Jade Copet, Olivier Teytaud, David Lopez-Paz. *Equal and main contribution

Meta Lingua is a minimal and fast LLM training and inference library designed for research. Meta Lingua uses easy-to-modify PyTorch components in order to try new architectures, losses, data, etc. We aim for this code to enable end to end training, inference and evaluation as well as provide tools to better understand speed and stability. While Meta Lingua is currently under development, we provide you with multiple `apps` to showcase how to use this codebase.

The screenshot displays the GitHub repository for Meta Lingua. At the top, there's a graph showing training progress: MMLU 48% at 24 hours with 256 GPUs and a 7B model. Below the graph, several metrics are listed: MMLU 48%, ARC Chal 50%, CSQA 63%, TQA 51%, and an ellipsis (...). To the right of these metrics is the **∞ Meta Lingua** logo. Further down, there are three boxes: one for "Fully Sharded Tensor Parallelism" and another for "Fully Compilable Fully Reproducible". The third box shows "Automatic probing throughout model Layers" and "... and training Train step". On the far left, there are small snippets of code: `mississippi`, `tokens`, `layers`, and `step`. To the right of the "train step" box, a vertical list of files is shown: `train.py`, `eval.py`, `data.py`, `transformer.py`, `distributed.py`, `checkpoint.py`, and `optim.py`. A bracket labeled "Minimal" groups the `train.py` through `optim.py` files.

Outline

- Overview of LM Pretraining
- Pretraining Data
- Training Setups
- Scaling Law

Pretraining and Compute

- Goal: get a better pretrained model by “adding more compute”
- *“The biggest lesson that can be read from 70 years of AI research is that general methods that leverage computation are ultimately the most effective, and by a large margin.”*
 - *The Bitter Lesson*, Richard Sutton 2019

What Is Compute?

$$\alpha \times \text{Model size} \quad \text{robot icon} \quad \times \quad \text{Training data} \quad \text{books icon} \quad = \quad \text{Training compute} \quad \text{laptop icon}$$

	Model size (# parameters)	Training data (# tokens)	Training compute (FLOPs)	Resources

$$C \approx 6ND$$

N : number of model parameters

D : number of tokens

C : compute; floating point operations (FLOPs)

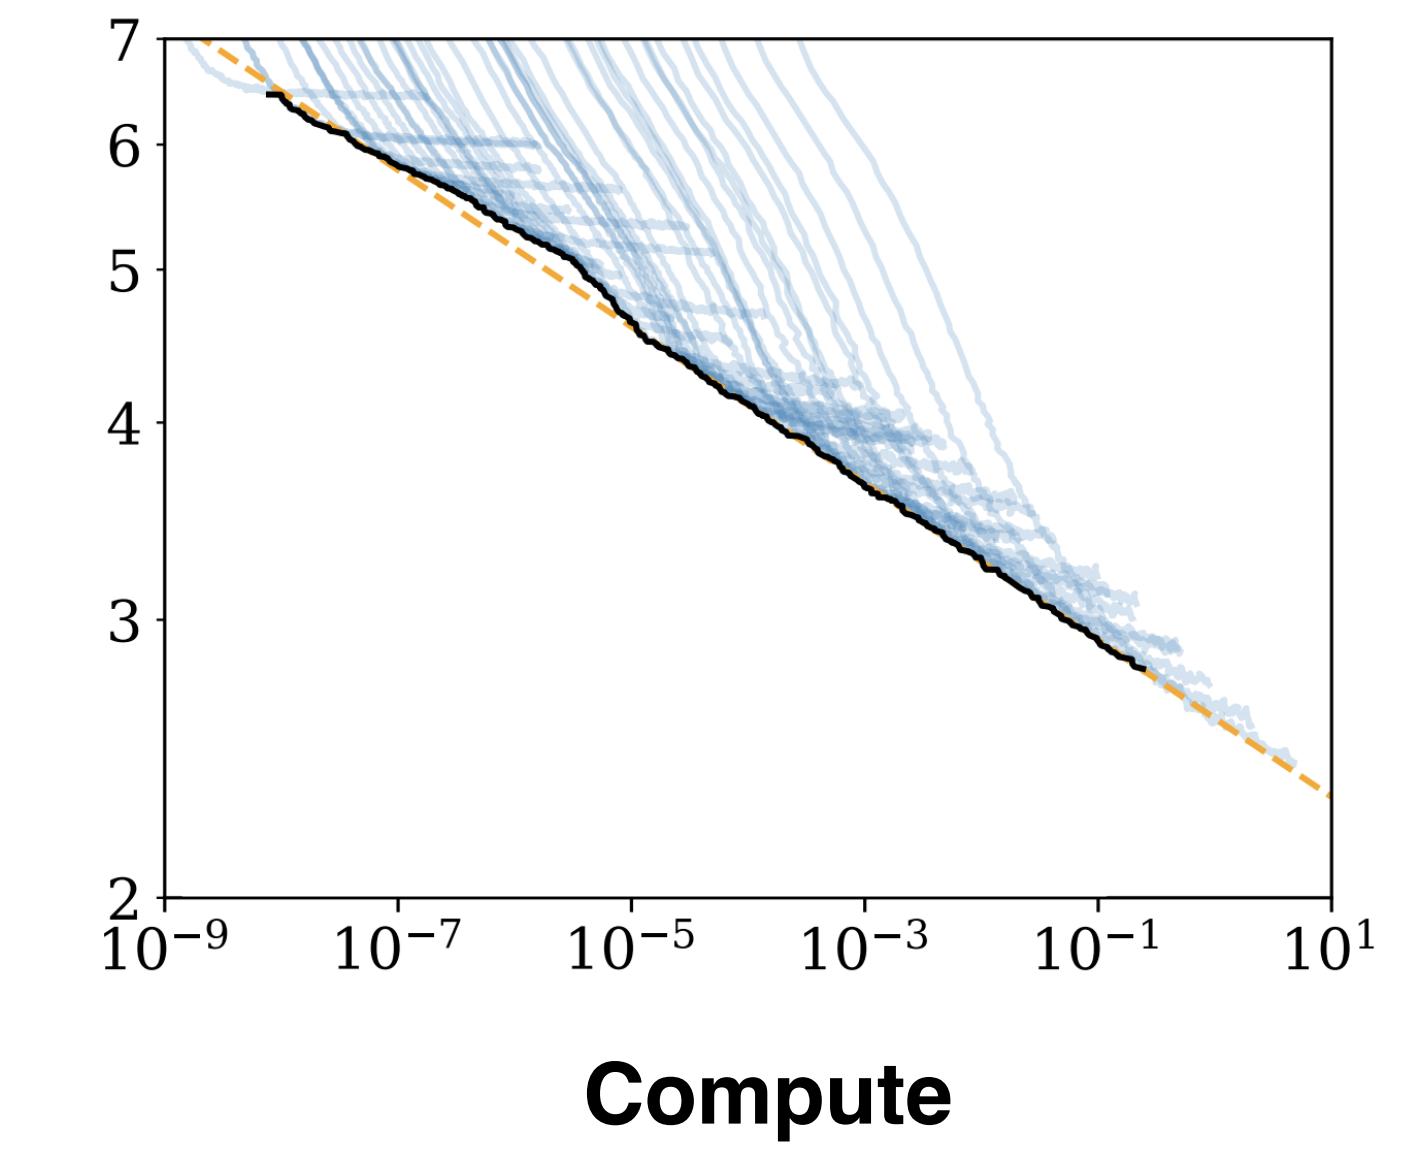
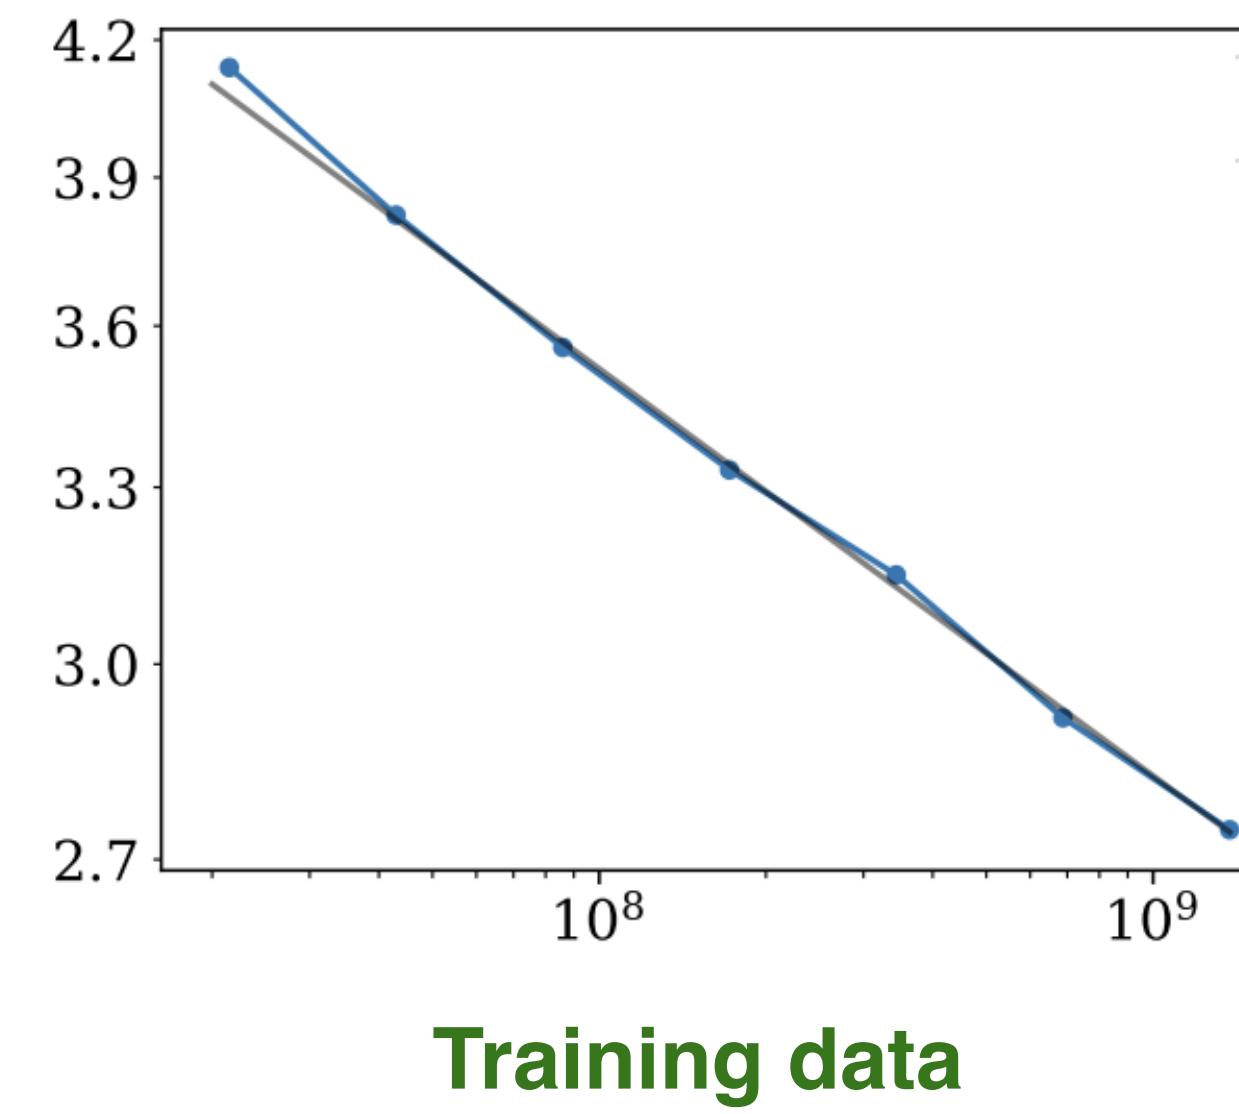
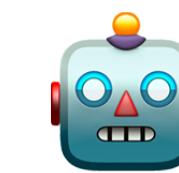
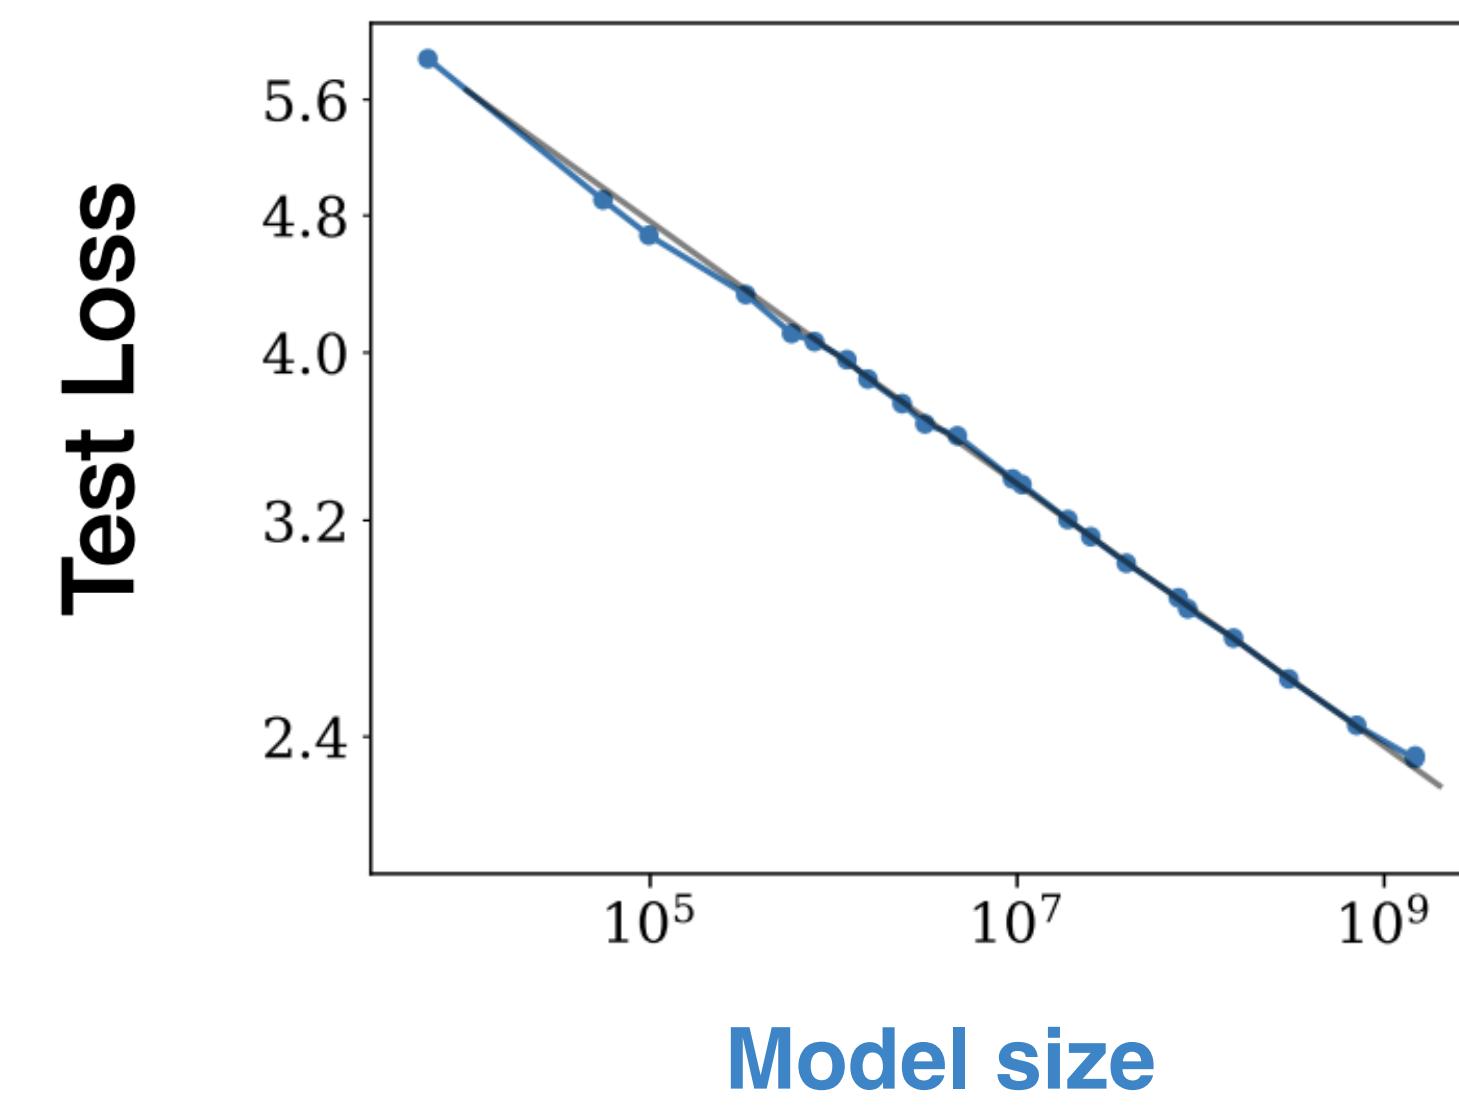
What Is Compute?

$$\alpha \times \text{Model size} \times \text{Training data} = \text{Training compute}$$

Icon descriptions: Model size (robot head), Training data (stack of books), Training compute (laptop).

	Model size (# parameters)	Training data (# tokens)	Training compute (FLOPs)	Resources
 BERT-base (2018)	109M	250B	1.6e20	64 TPU v2 for 4 days (16 V100 GPU for 33 hrs)
 GPT-3 (2020)	175B	300B	3.1e23	~1,000x BERT-base
 PaLM (2022)	540B	780B	2.5e24	6k TPU v4 for 2 months

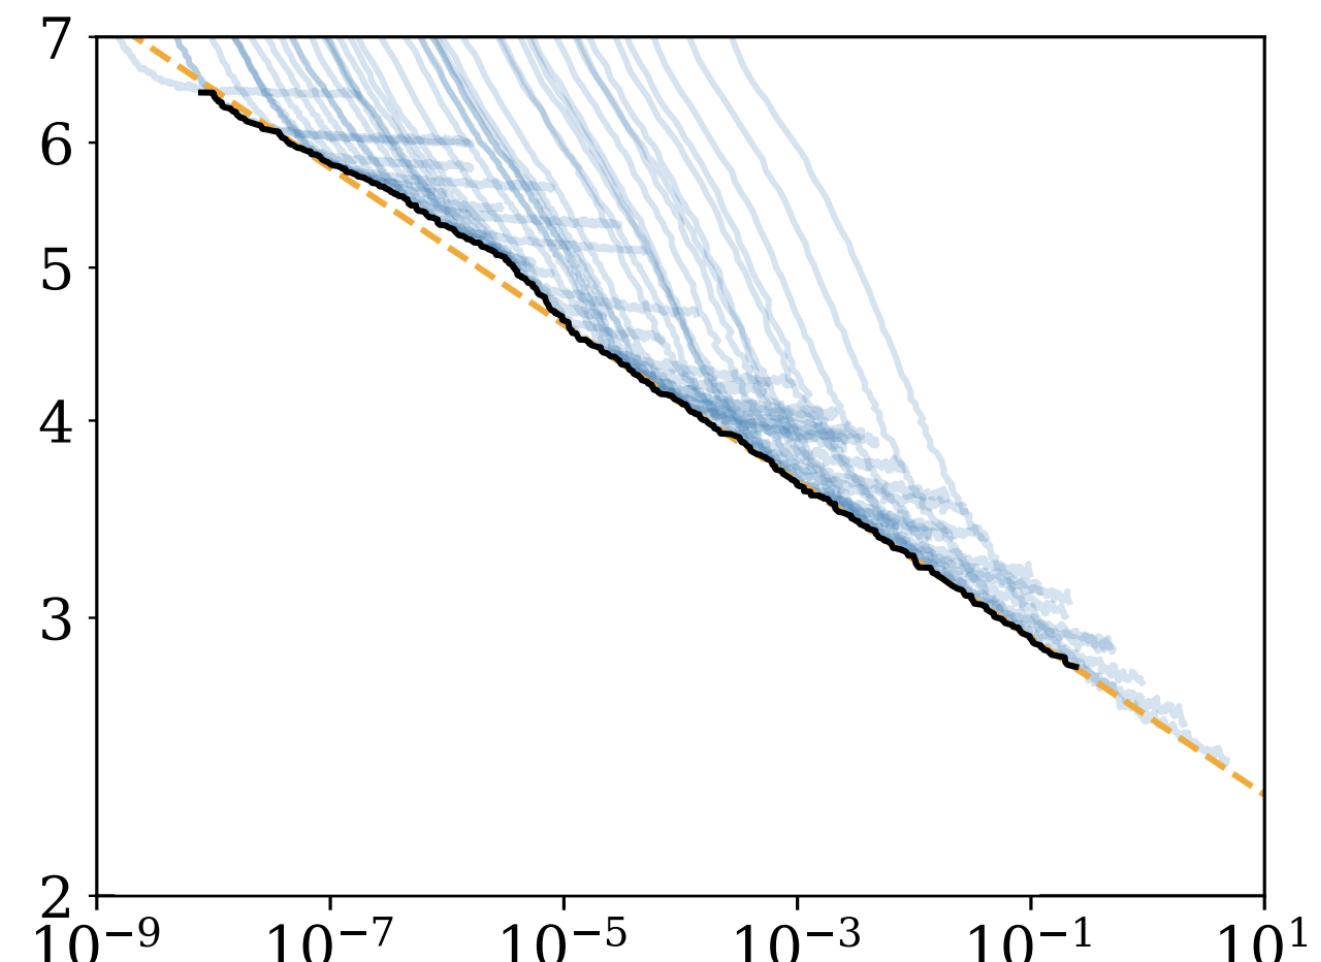
How important is scaling? (Return)



*Test loss predictably **improves** with more **compute***

Scaling Laws

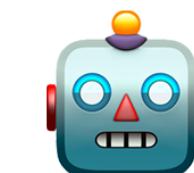
Given a fixed compute budget, can we predict the optimal test loss?



Train models of different sizes and numbers of tokens



Predictive formula



We can estimate loss (L) given **model size (N)**, **training data (D)**, and learned constants:

$$L(N, D) = \frac{A}{N^\alpha} + \frac{B}{D^\beta} + E$$

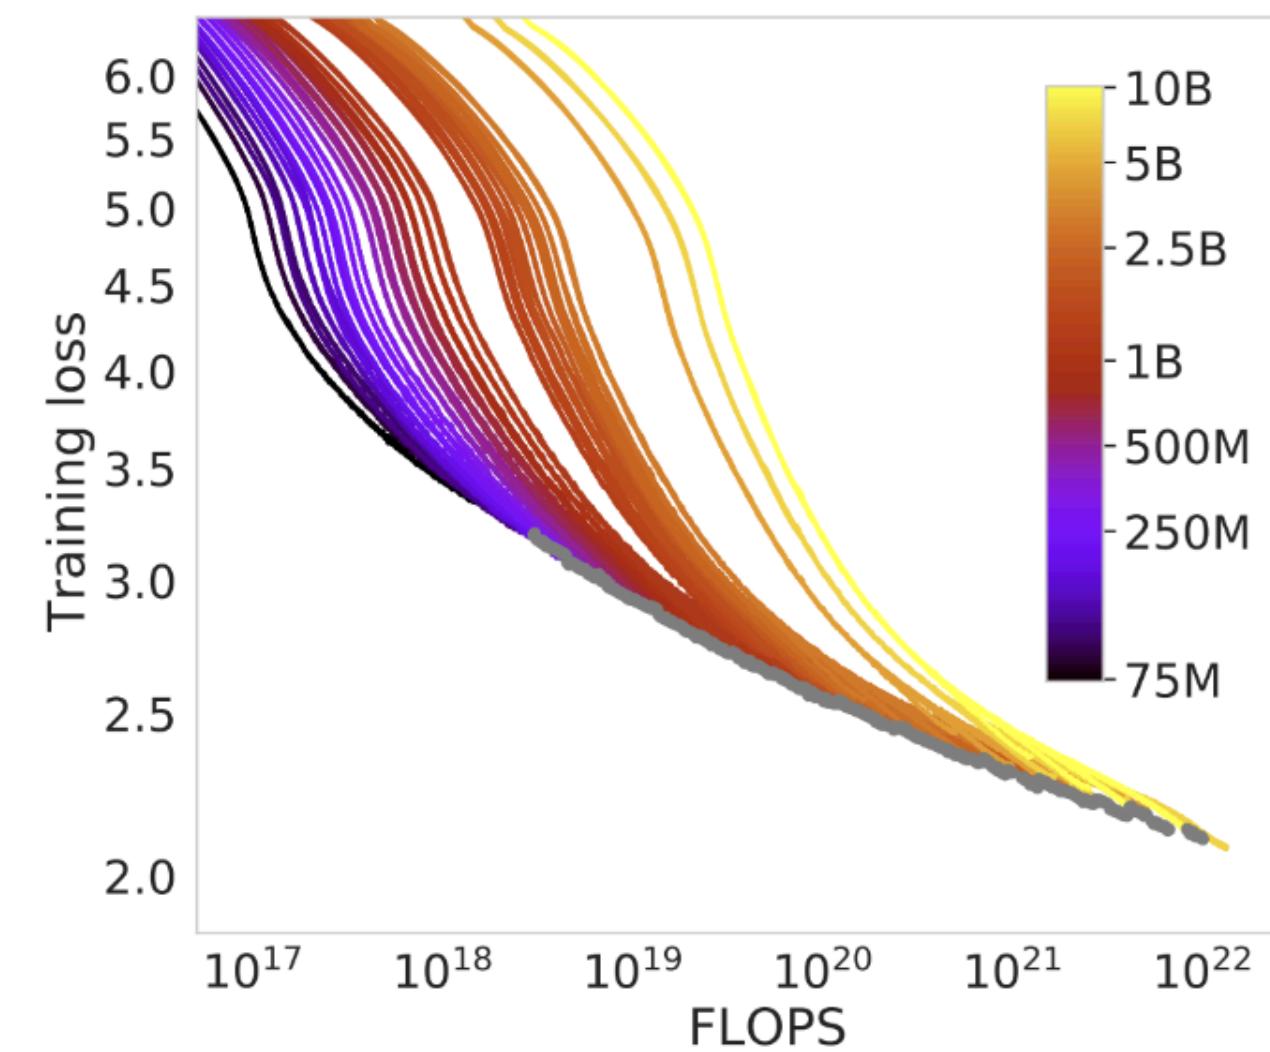
Fitting the constants, yields: $\alpha \approx \beta$

i.e. equal scaling of **N** and **D**.

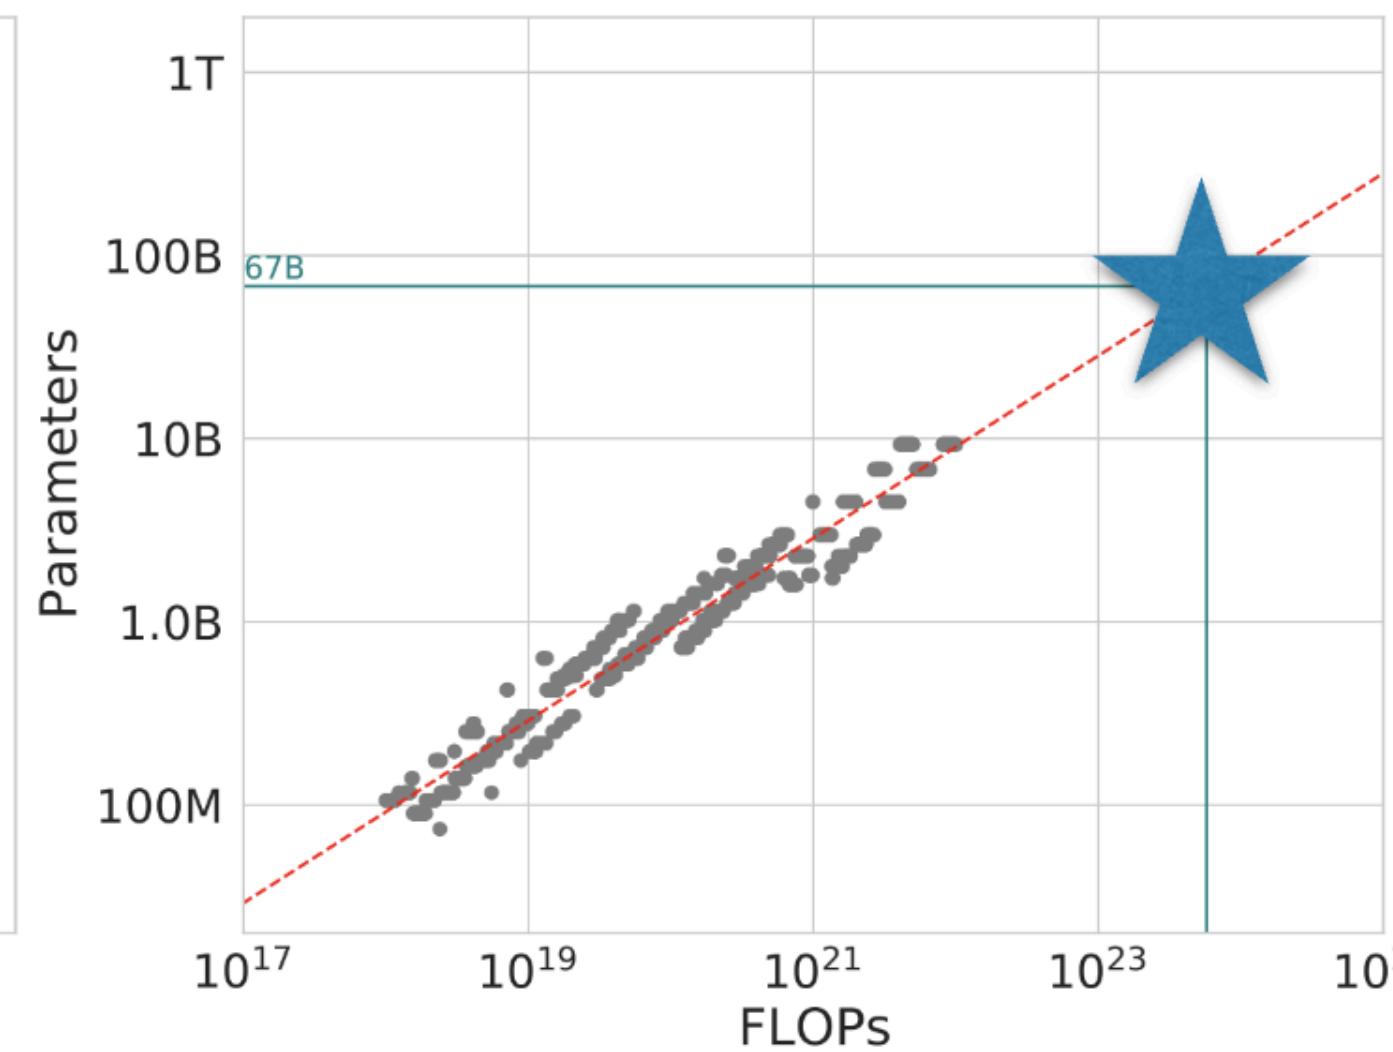
Using Scaling Laws

- Scaling laws are also used to choose hyper parameters
- Basic idea:
 - Run many experiments at a small scale
 - Use a scaling law to estimate the best hyper parameter for a large-scale model / training run

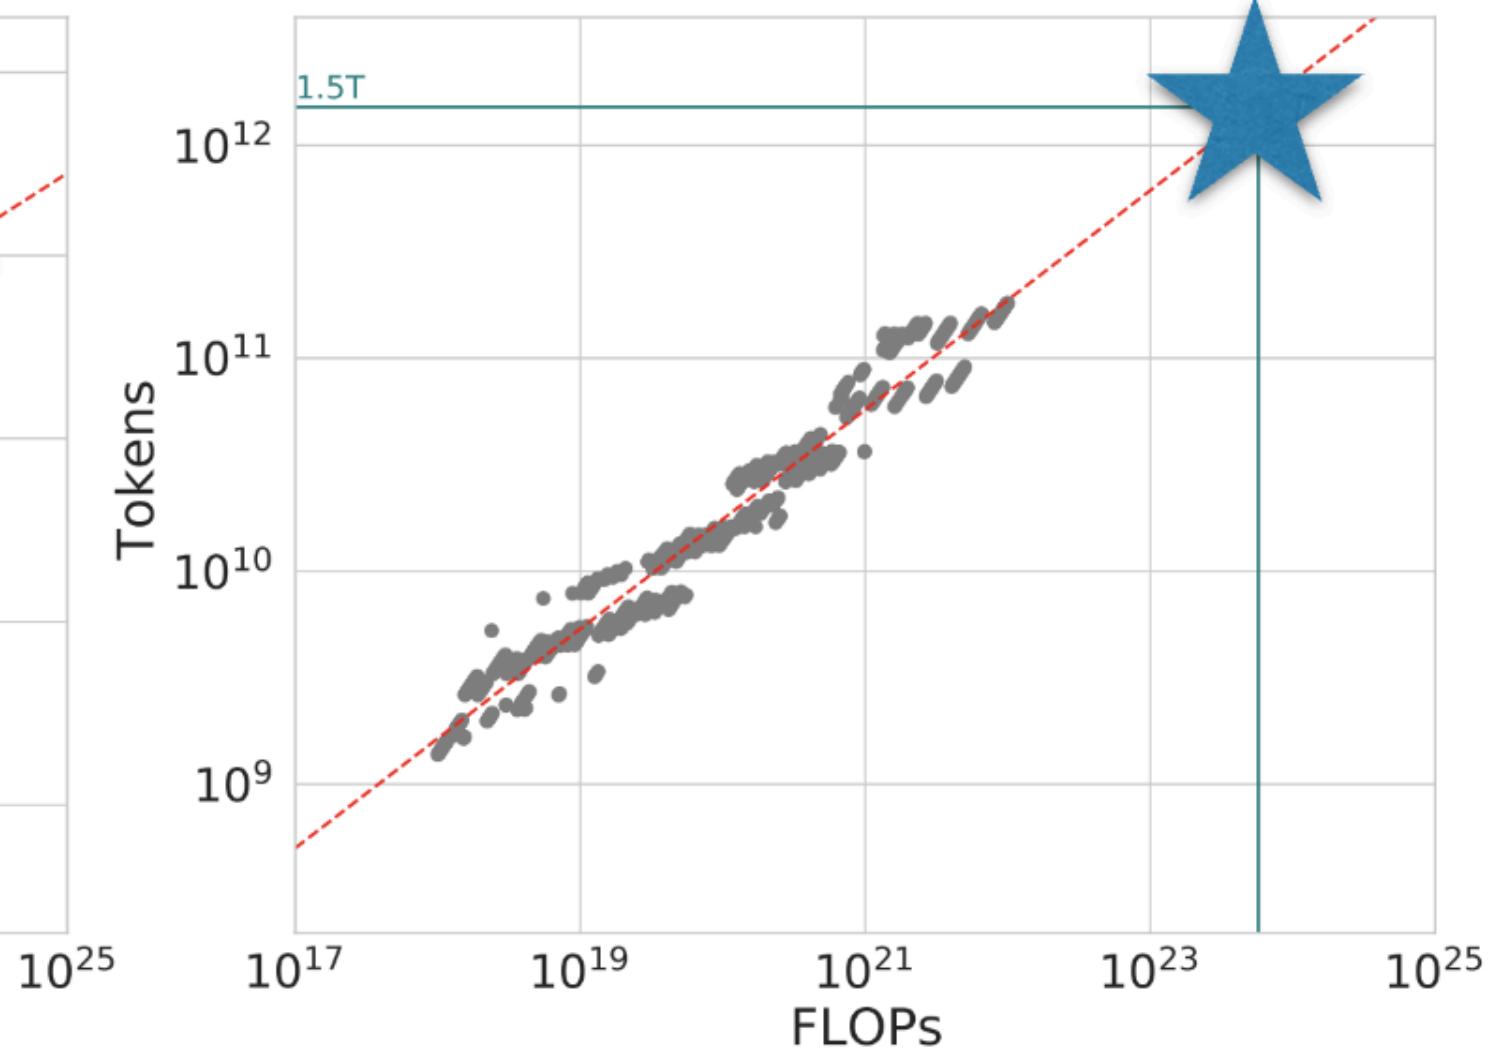
Example: choose model size and # of tokens



Run experiments



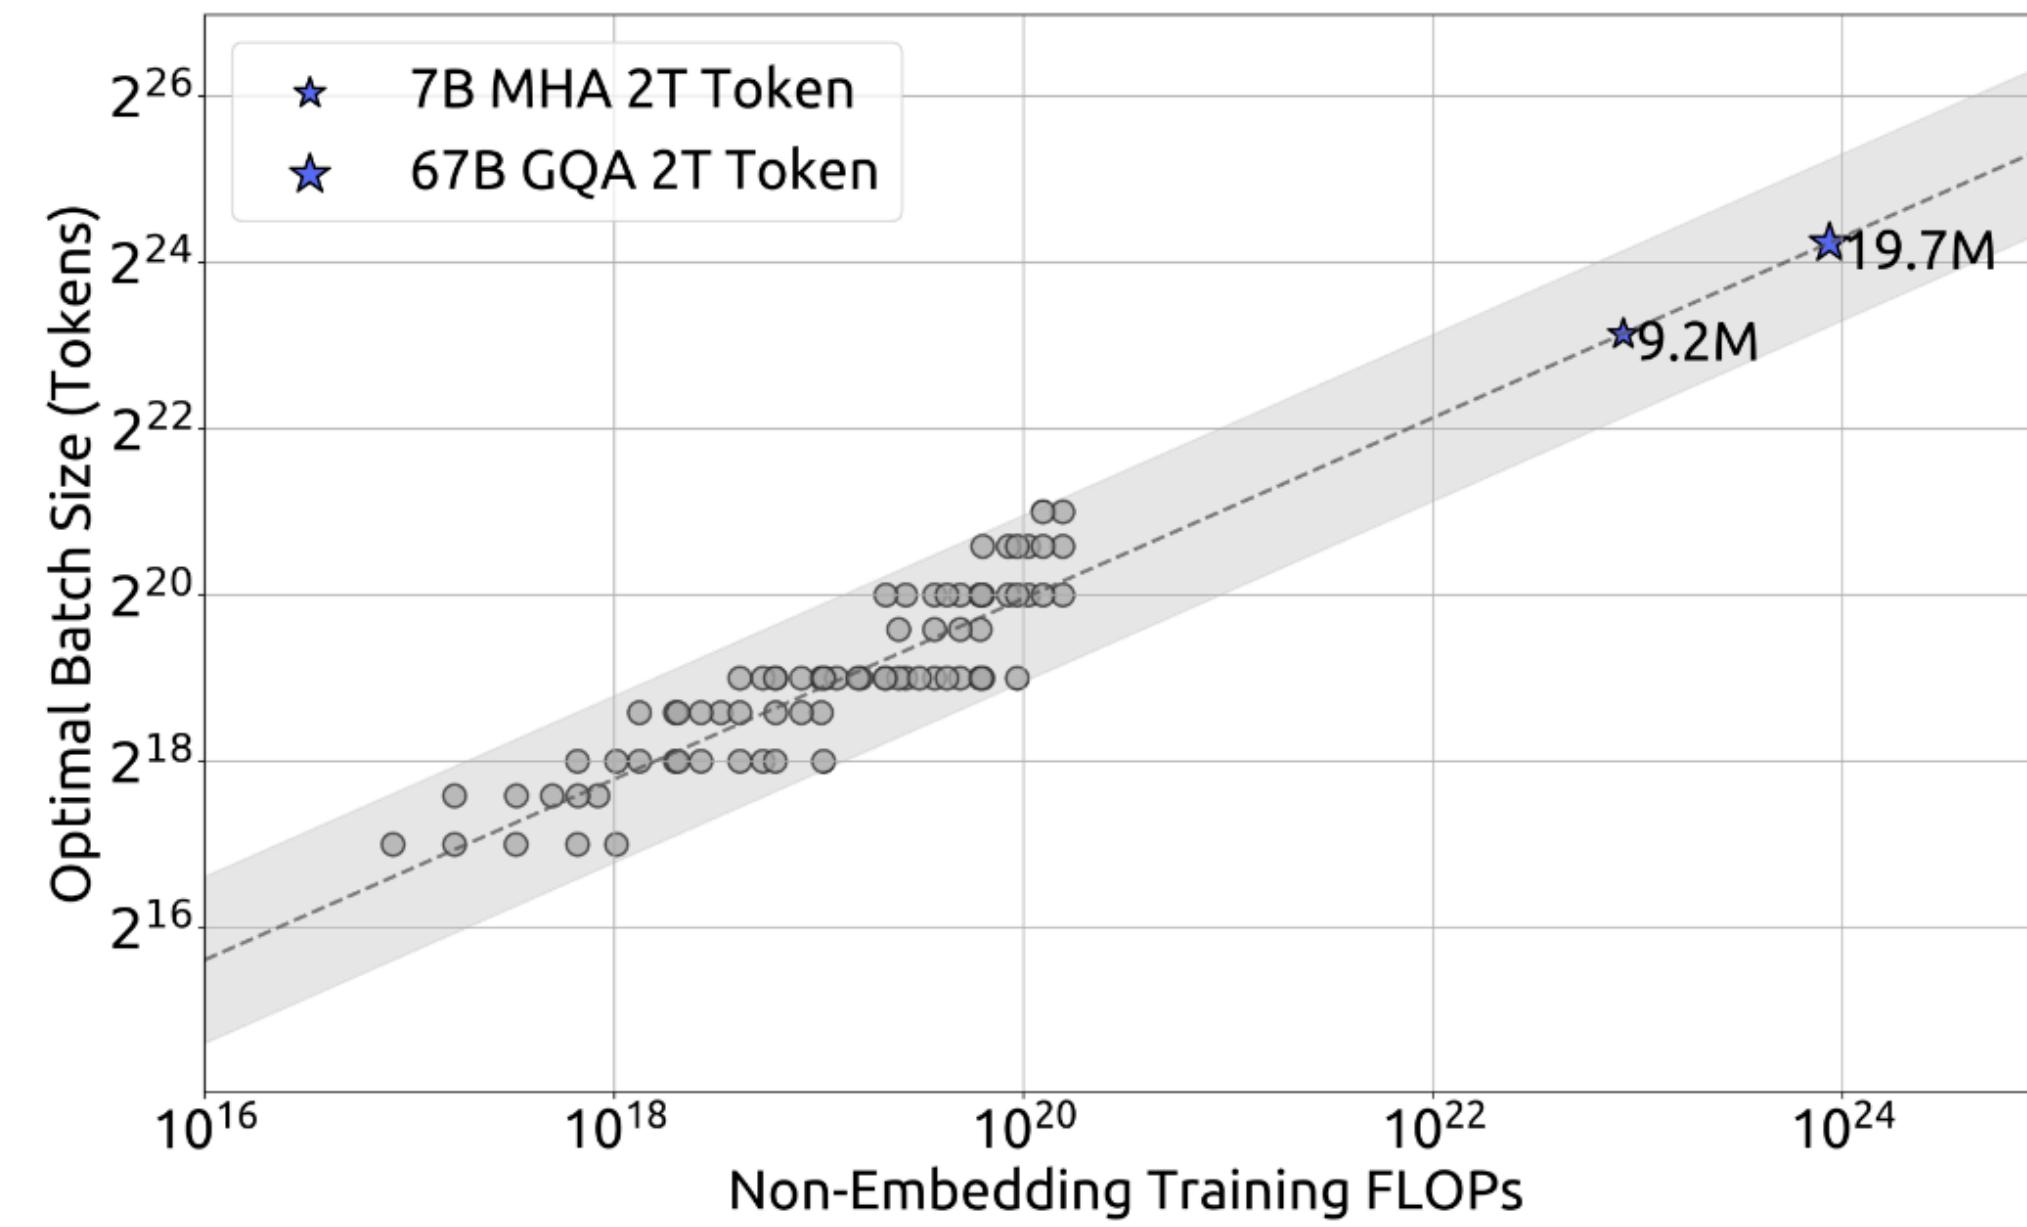
Fit a line and
predict optimal
model size



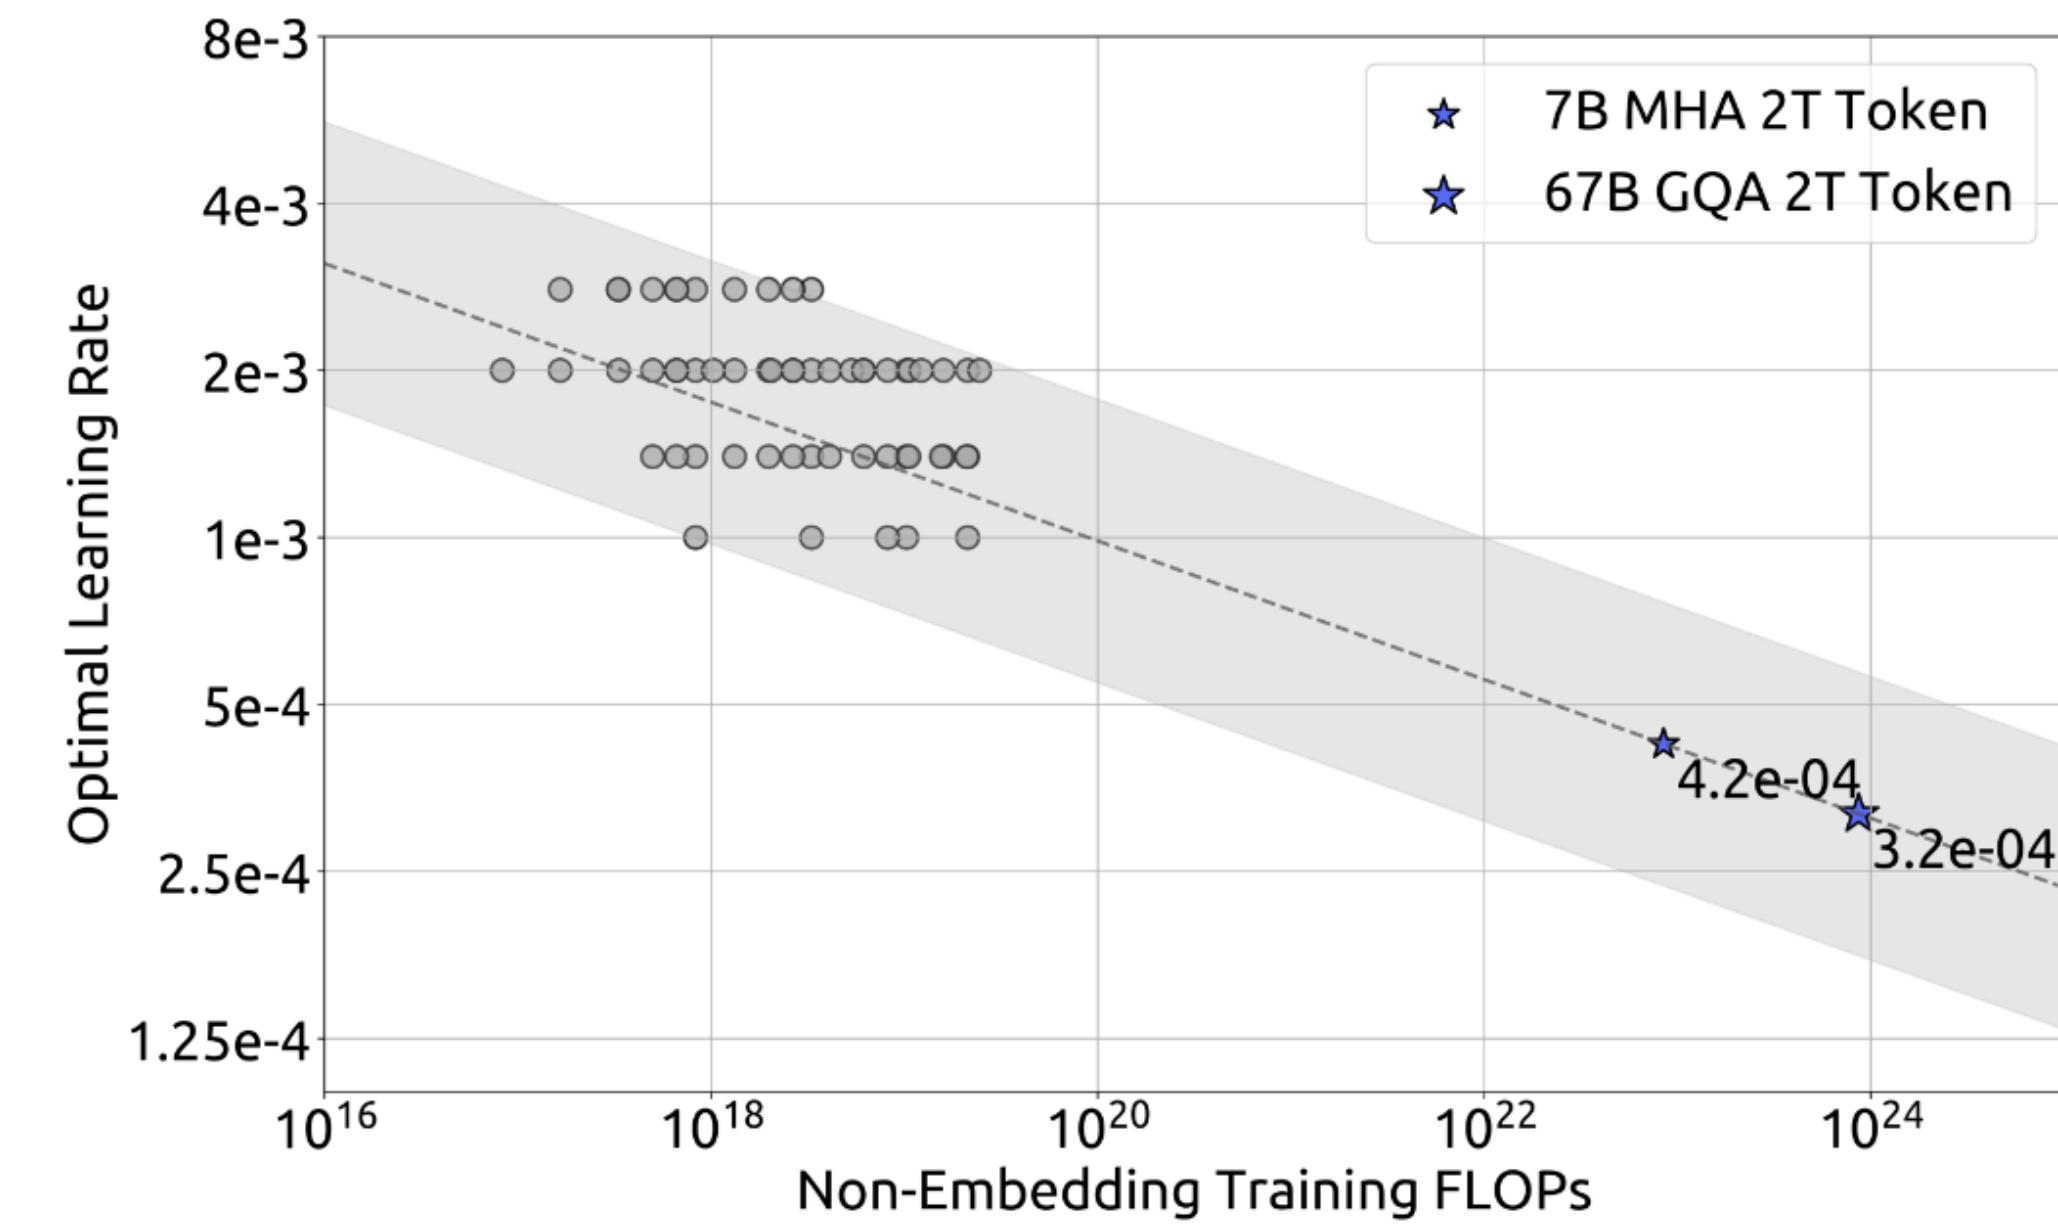
Fit a line and
predict optimal
of tokens

“Optimal”: best loss for a given compute budget (FLOPs)

Example: choose batch size, learning rate



Optimal batch size



Optimal learning rate

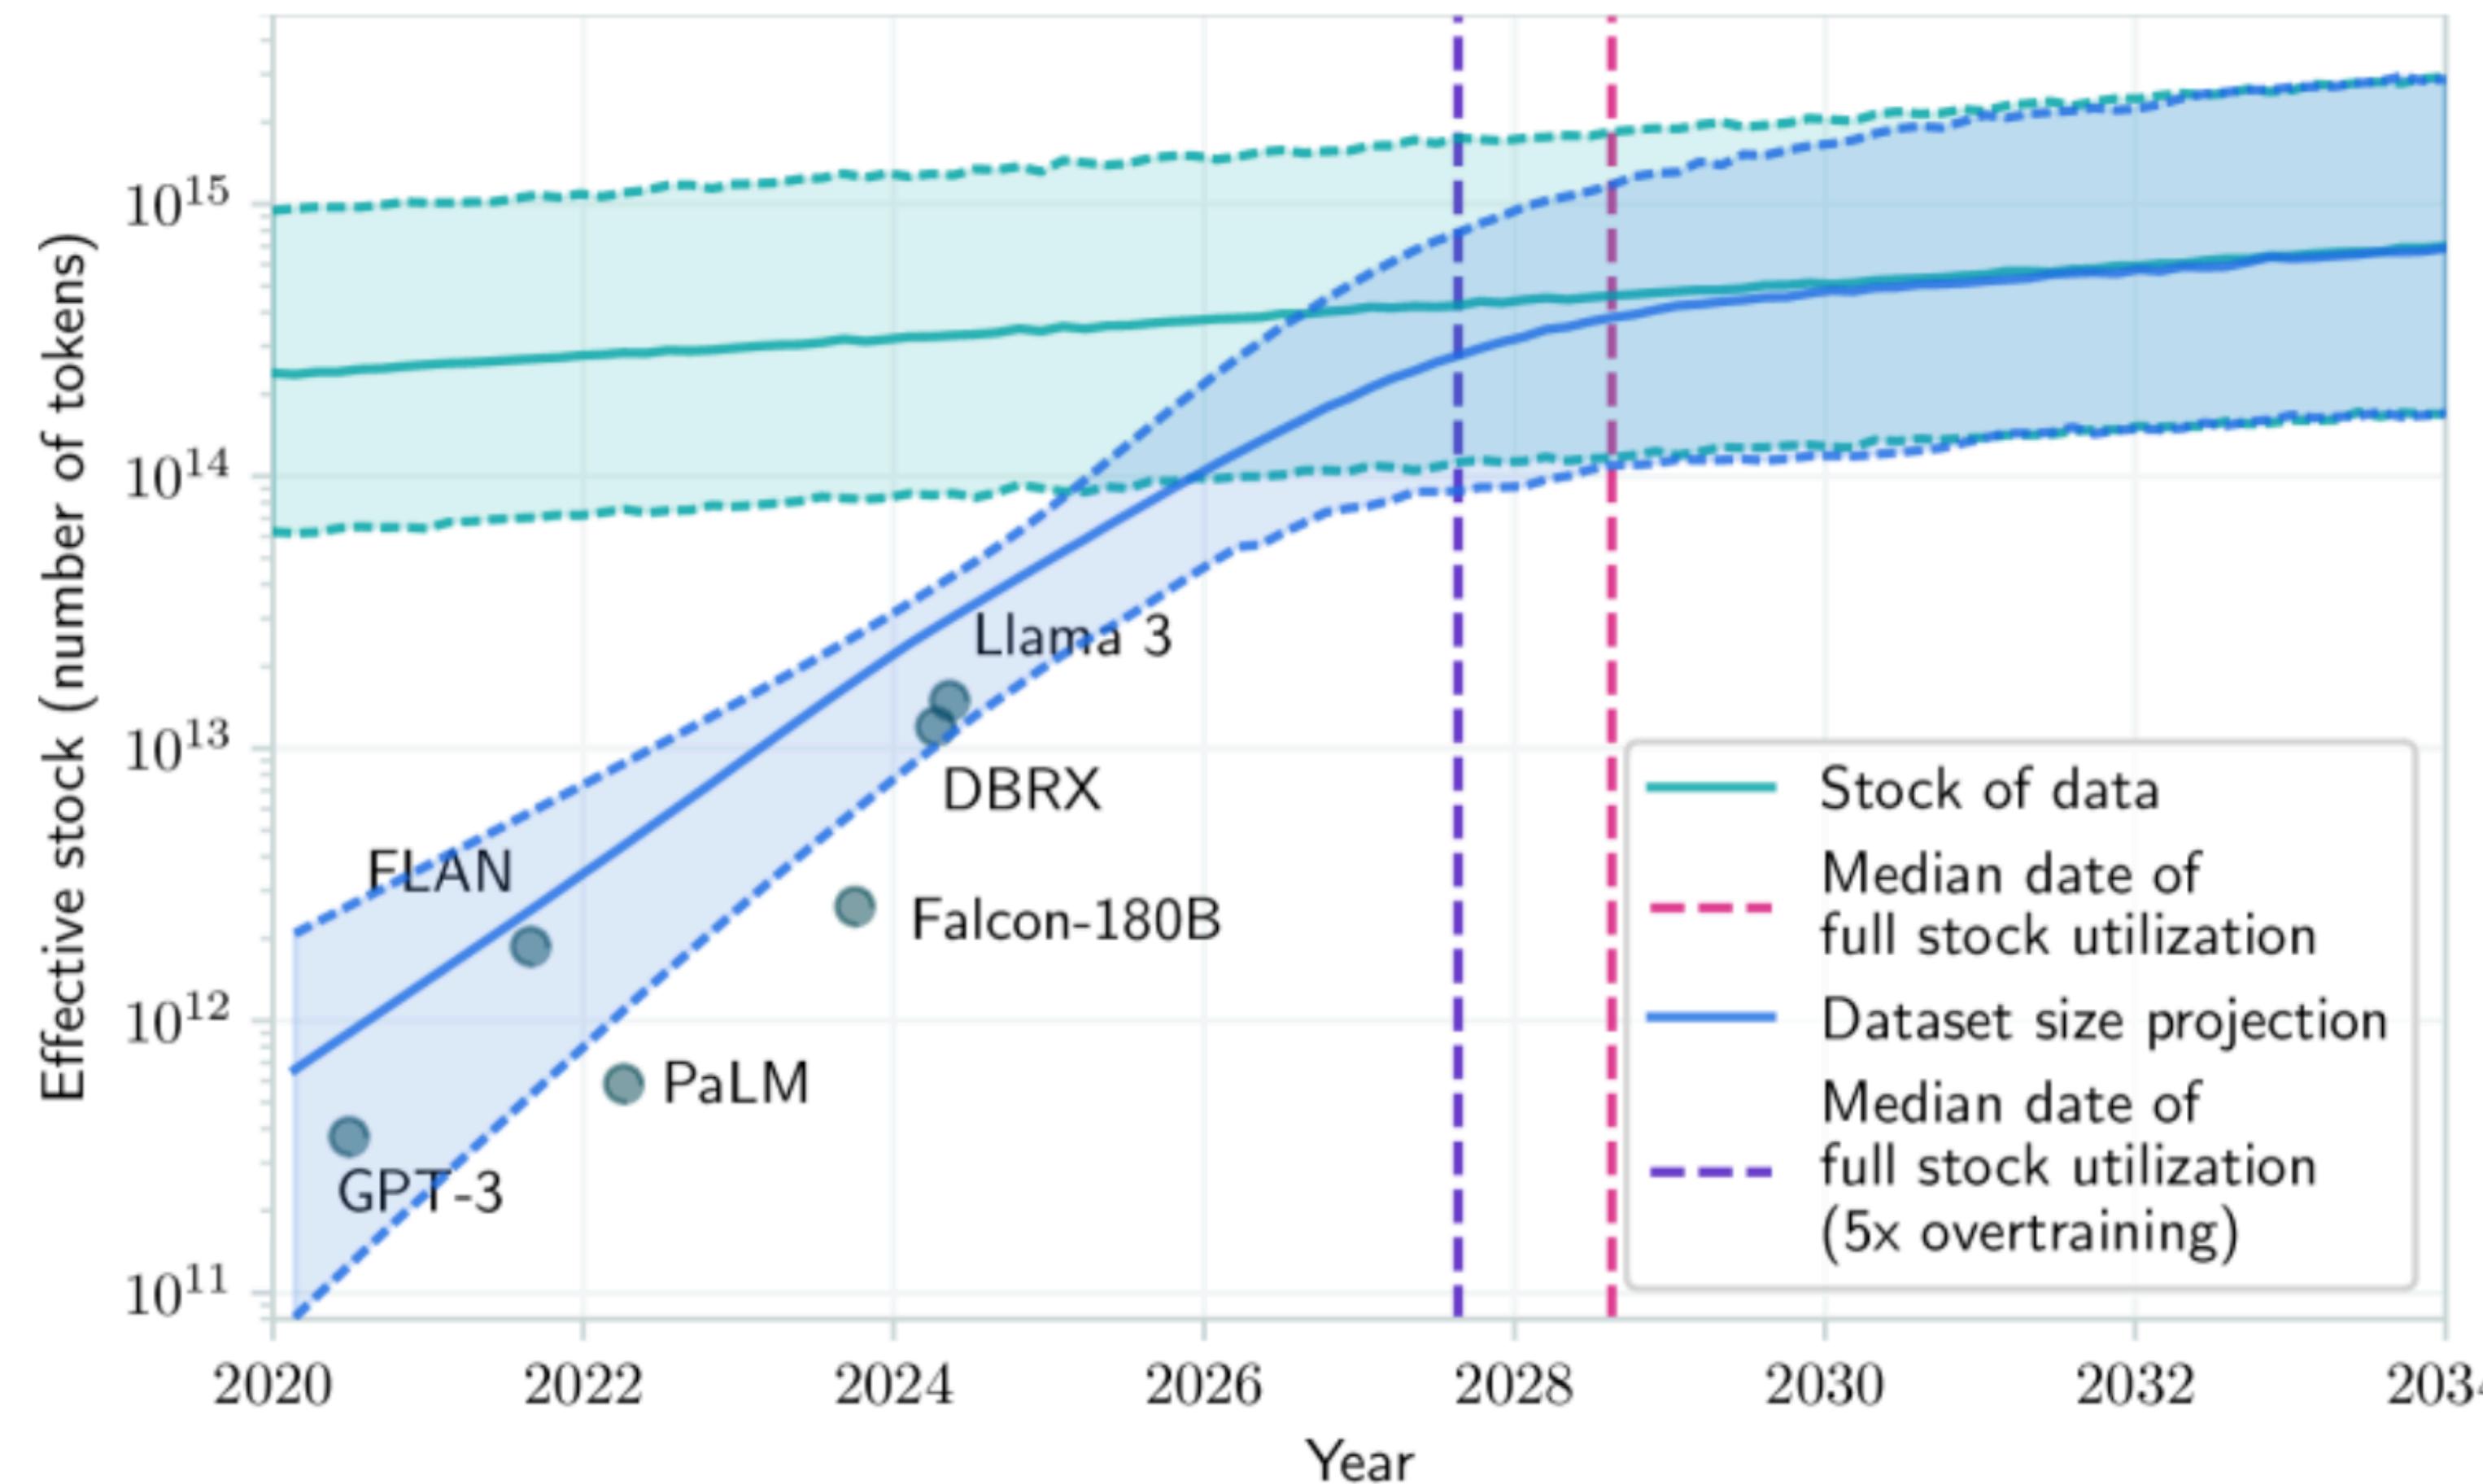
Limits of Scaling

Pretraining as we know it will end
—Ilya Sutskever

Limits of Scaling

- **Limits on data:** Modern LLMs are trained on basically the *entire internet* - we can't find 10 new internets out of nowhere
- **Limits on compute:** Big tech companies can't continue to 10x their model sizes for much longer

Limits on Data: Data Is Running Out

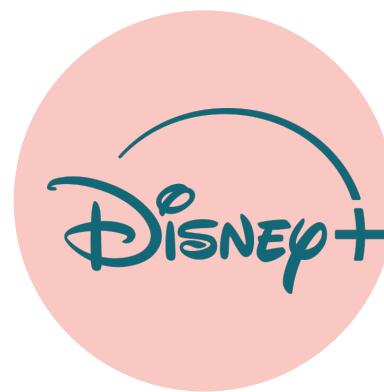
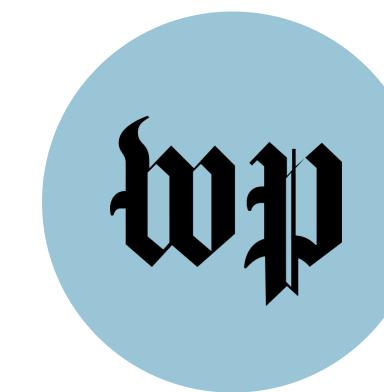


Limits on Data: Restrictions in Use

Public data is always usable, but proprietary/licensed data is not

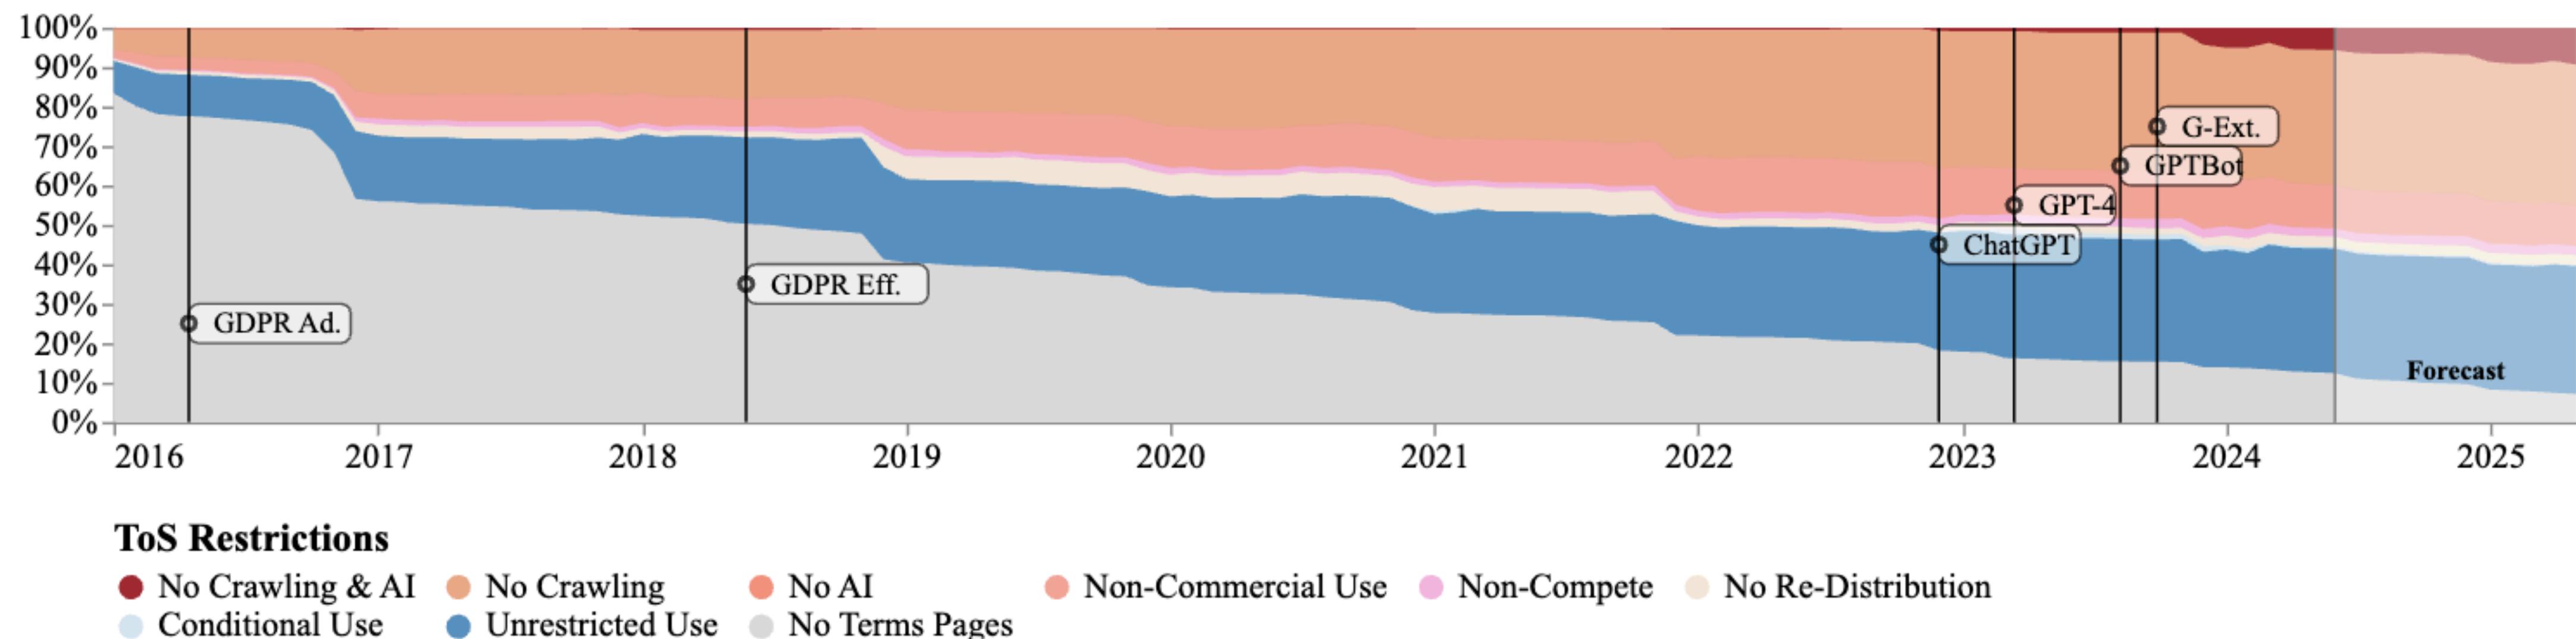


Public Data



Proprietary Data

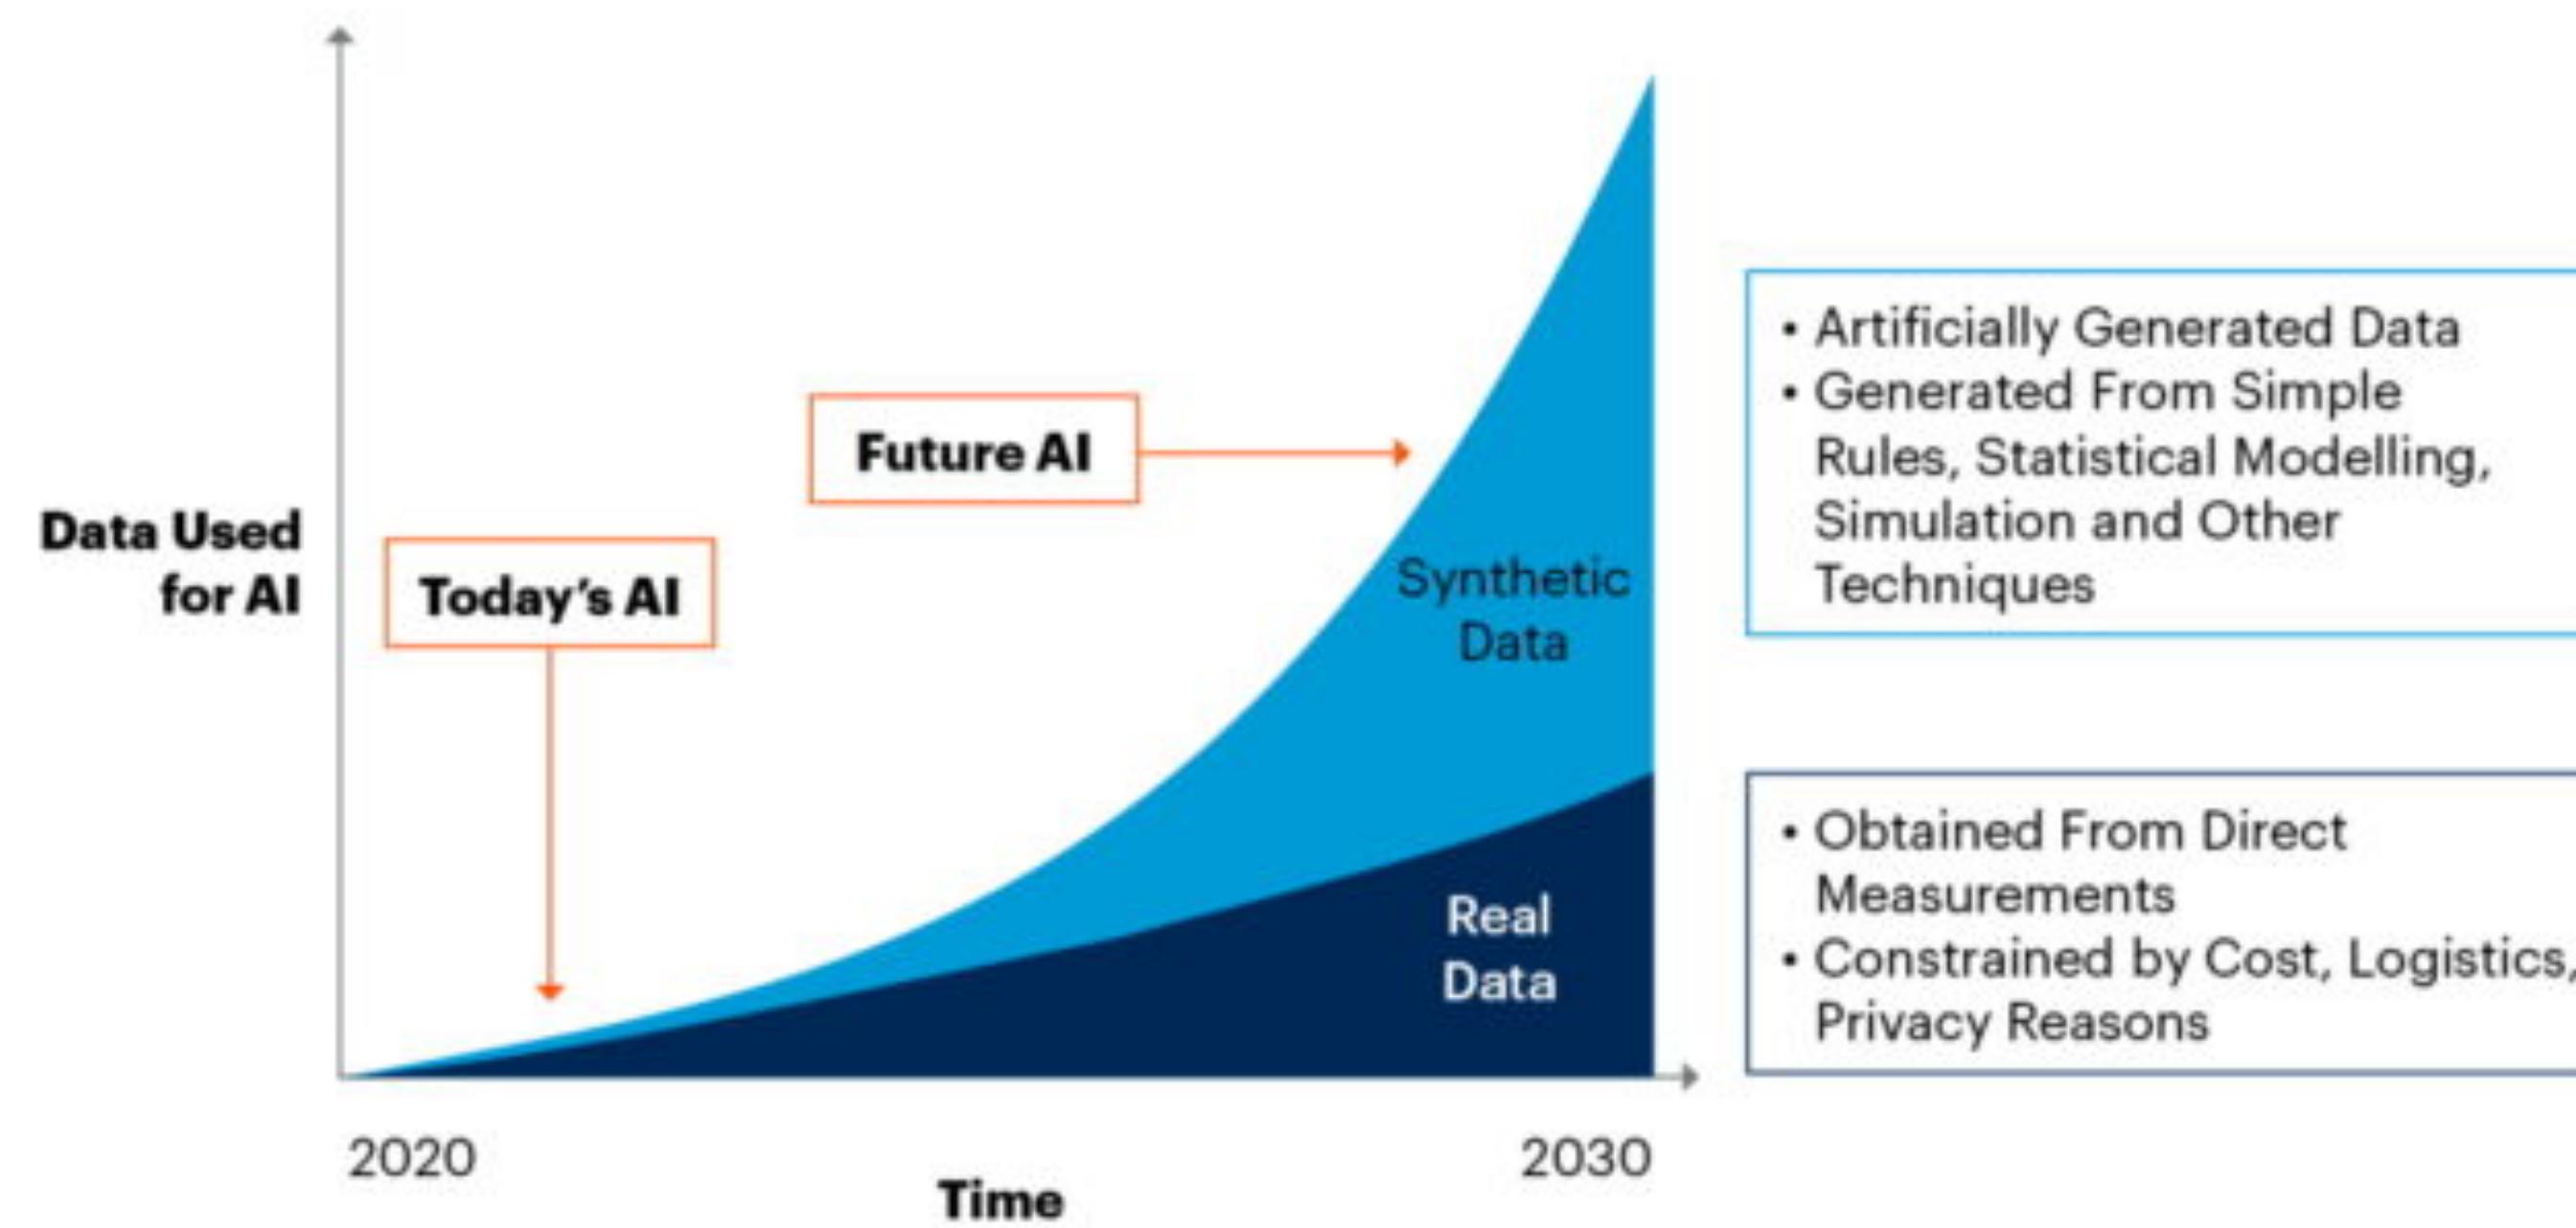
Limits on Data: Restrictions in Use



Terms of Service pages have imposed more **anti-crawling** and now **anti-AI** restrictions

Limits on Data: Synthetic Data

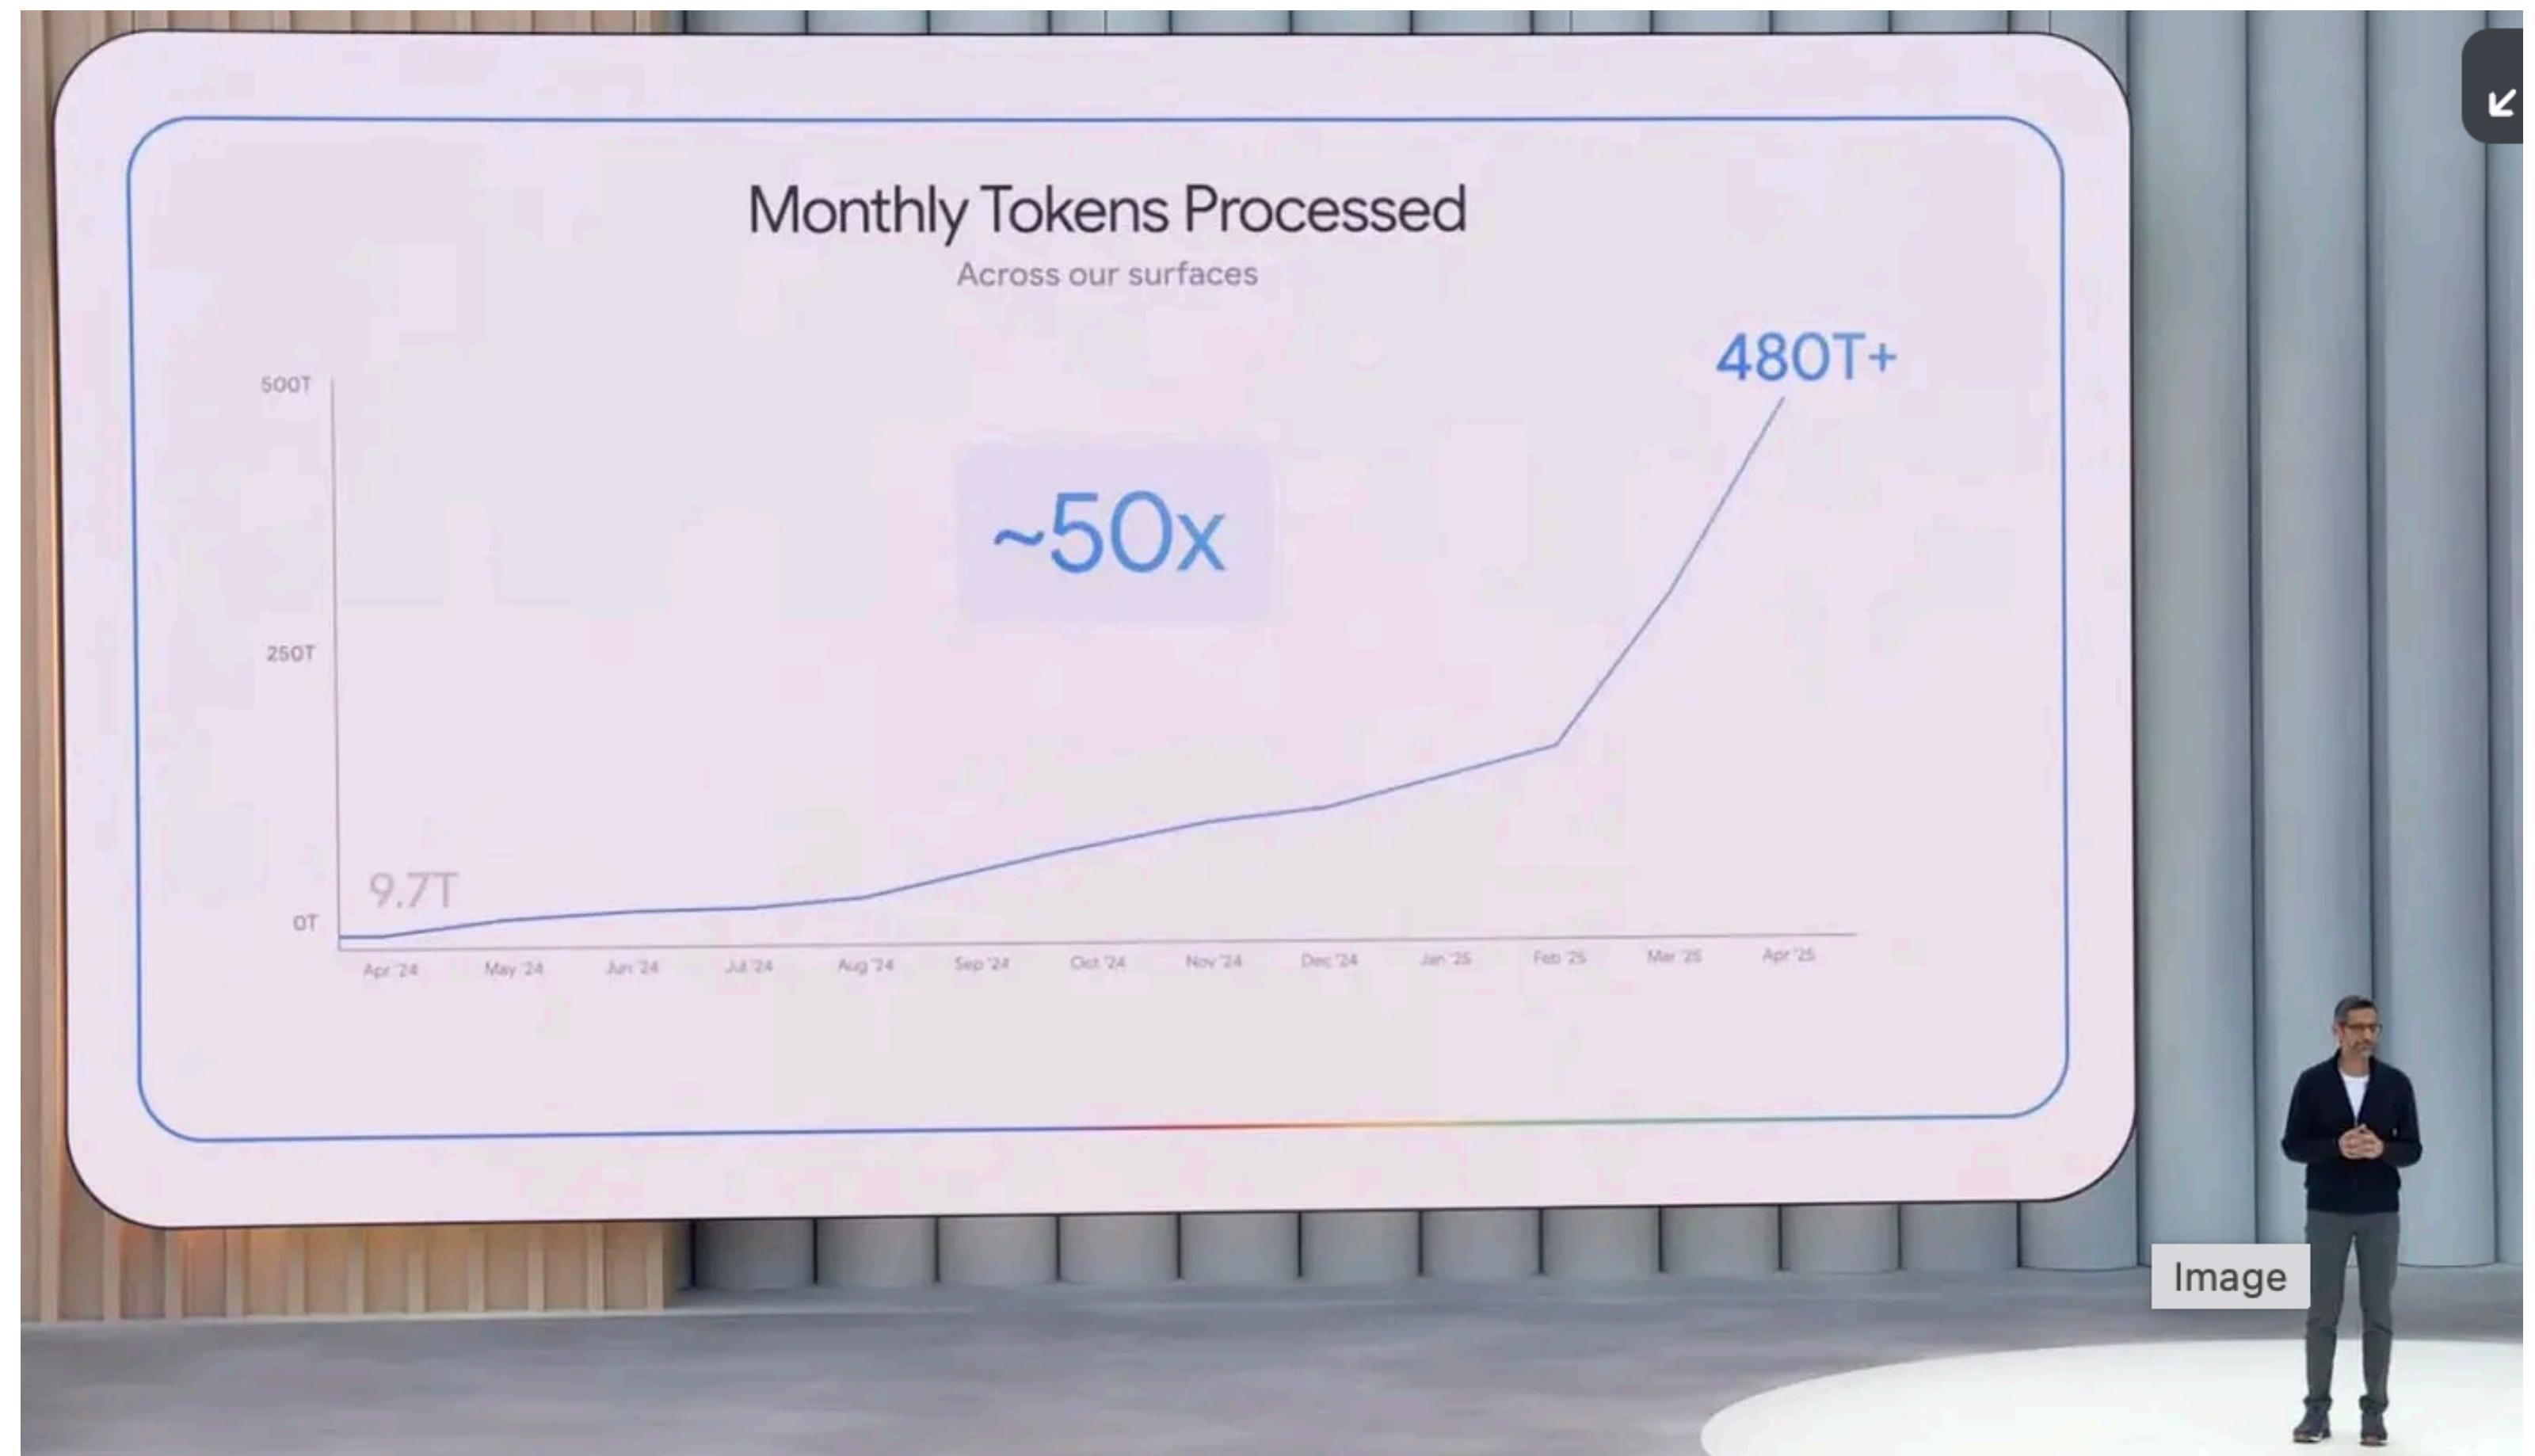
By 2030, Synthetic Data Will Completely Overshadow Real Data in AI Models



Gartner

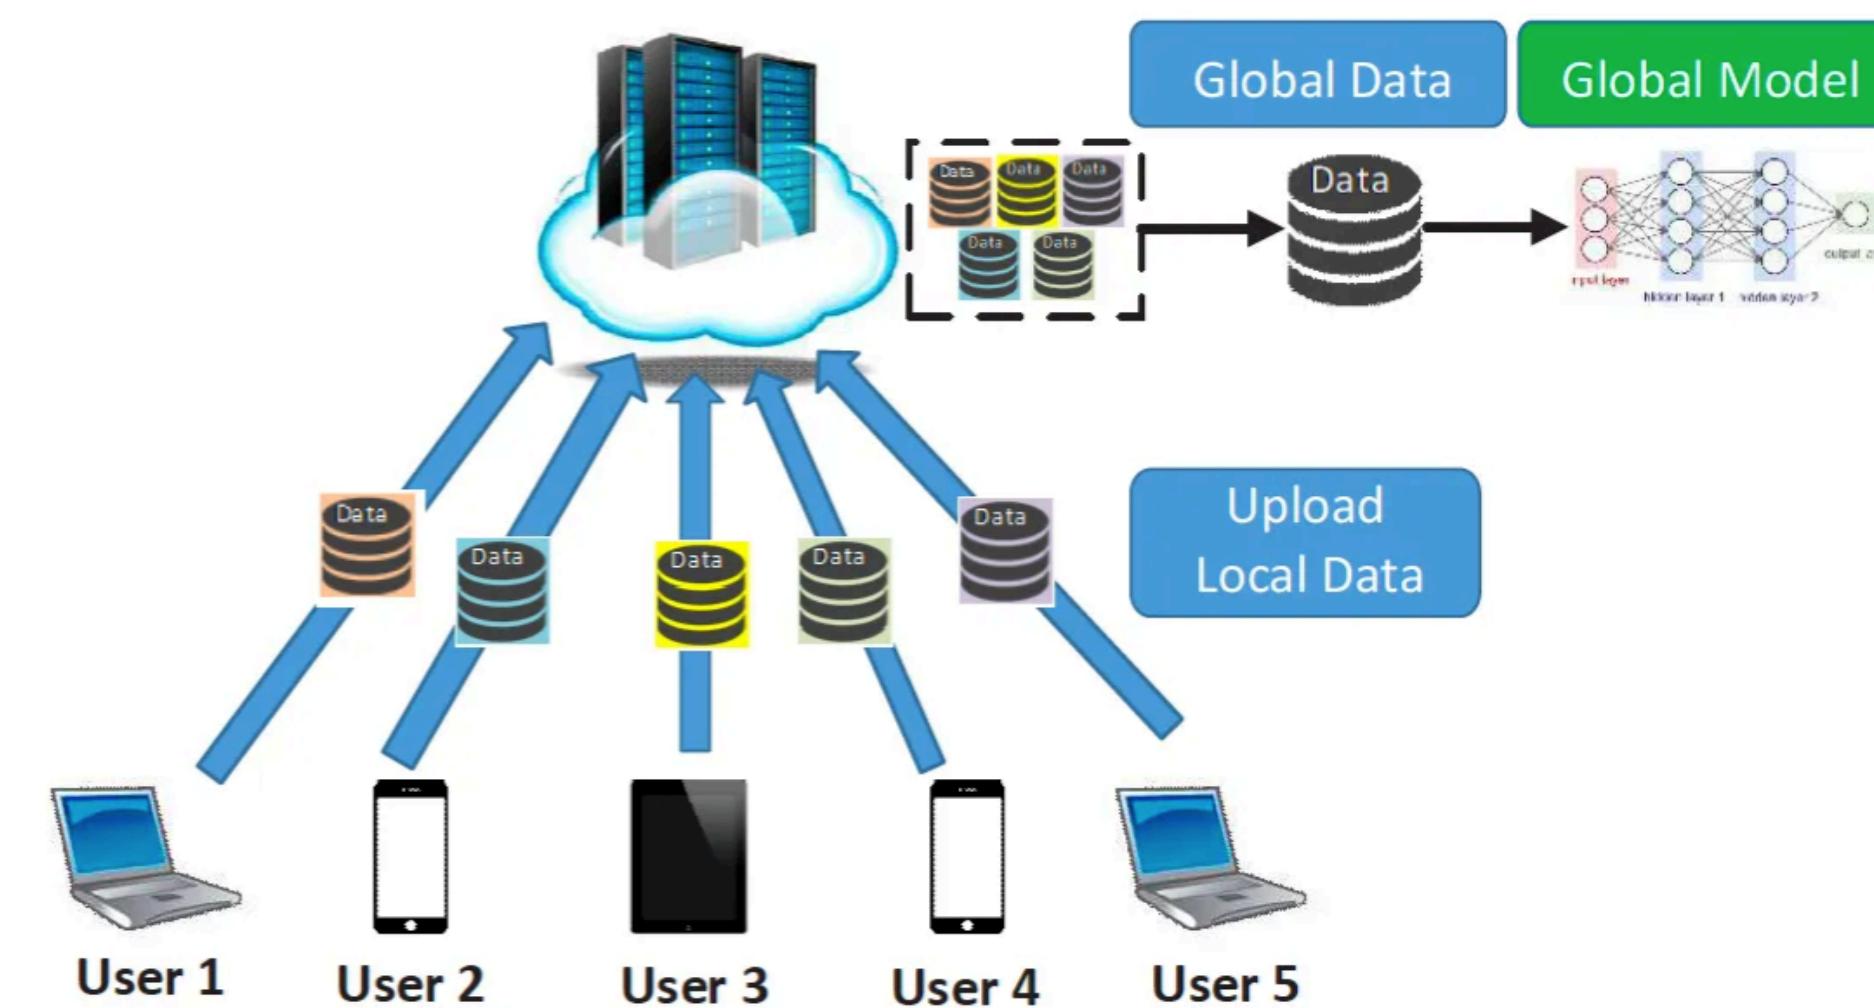
Limits on Data: Synthetic Data

- April 2024: 9.7T tokens
- December 2024: 90T tokens
- February 2025: 160T tokens
- March 2025: 300T tokens
- April 2025: 480T+ tokens



Limits on Compute: Pretraining is Centralized

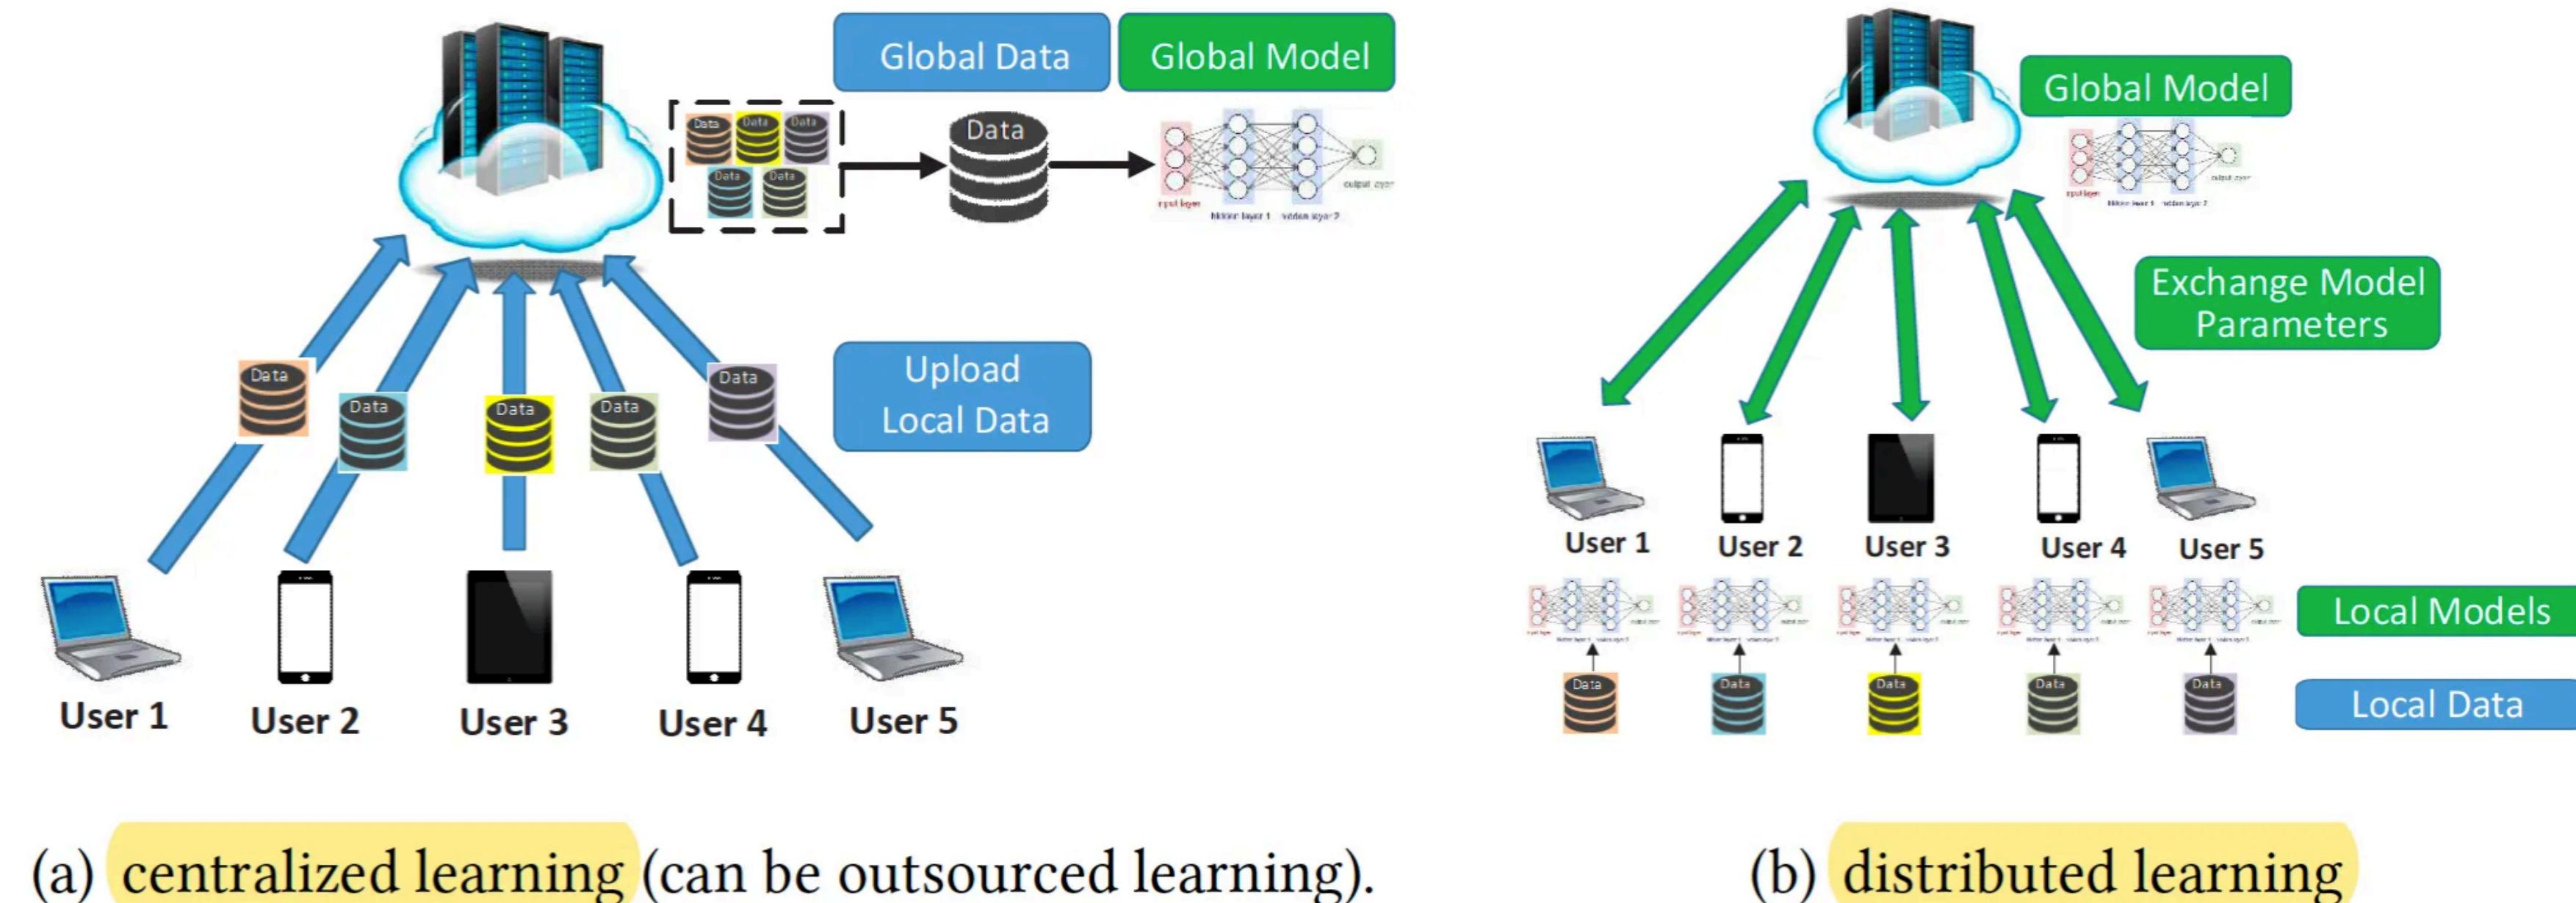
- Current pretraining requires GPUs' communications
- But one data center can hold fixed amount of GPUs



(a) centralized learning (can be outsourced learning).

Limits on Compute: Decentralized Training

- Current pretraining requires GPUs' communications
- But one data center can hold fixed amount of GPUs



What Comes Next?

What Comes Next?

- Agents
- Reasoning/Inference time compute
- Synthetic Data
- Self-improving AI
- Towards full embodiment: Systems that learn interactively from rich supervisions
- World Model

Questions?

Please share any feedback you might have:
<https://forms.gle/14M3DaNrJGDoBLoH8>